

Slides for the
Machine Learning exercise book

(published in Romanian)

by **Liviu Ciortuz, Alina Munteanu, Elena Bădăraș**

Faculty of Computer Science, University of Iași, Romania

Random Variables:

Some proofs

$$E[X + Y] = E[X] + E[Y]$$

where X and Y are random variables of the same type (i.e. either discrete or cont.)

The discrete case:

$$\begin{aligned} E[X + Y] &= \sum_{\omega \in \Omega} (X(\omega) + Y(\omega)) \cdot P(\omega) \\ &= \sum_{\omega} X(\omega) \cdot P(\omega) + \sum_{\omega} Y(\omega) \cdot P(\omega) = E[X] + E[Y] \end{aligned}$$

The continuous case:

$$\begin{aligned} E[X + Y] &= \int_x \int_y (x + y) p_{XY}(x, y) dy dx \\ &= \int_x \int_y x p_{XY}(x, y) dy dx + \int_x \int_y y p_{XY}(x, y) dy dx \\ &= \int_x x \int_y p_{XY}(x, y) dy dx + \int_y y \int_x p_{XY}(x, y) dx dy \\ &= \int_x x p_X(x) dx + \int_y y p_Y(y) dy = E[X] + E[Y] \end{aligned}$$

X and Y are independent $\Rightarrow E[XY] = E[X] \cdot E[Y]$,

X and Y being random variables of the same type (i.e. either discrete or continuous)

The discrete case:

$$\begin{aligned} E[XY] &= \sum_{x \in \text{Val}(X)} \sum_{y \in \text{Val}(Y)} xy P(X = x, Y = y) = \sum_{x \in \text{Val}(X)} \sum_{y \in \text{Val}(Y)} xy P(X = x) \cdot P(Y = y) \\ &= \sum_{x \in \text{Val}(X)} \left(x P(X = x) \sum_{y \in \text{Val}(Y)} y P(Y = y) \right) = \sum_{x \in \text{Val}(X)} x P(X = x) E[Y] = E[X] \cdot E[Y] \end{aligned}$$

The continuous case:

$$\begin{aligned} E[XY] &= \int_x \int_y xy p(X = x, Y = y) dy dx = \int_x \int_y xy p(X = x) \cdot p(Y = y) dy dx \\ &= \int_x x p(X = x) \left(\int_y y p(Y = y) dy \right) dx = \int_x x p(X = x) E[Y] dx \\ &= E[Y] \cdot \int_x x p(X = x) dx = E[X] \cdot E[Y] \end{aligned}$$

Binomial distribution: $b(r; n, p) \stackrel{\text{def.}}{=} C_n^r p^r (1 - p)^{n-r}$

Significance: $b(r; n, p)$ is the probability of drawing r *heads* in n independent flips of a coin having the head probability p .

$b(r; n, p)$ indeed represents a **probability distribution**:

- $b(r; n, p) = C_n^r p^r (1 - p)^{n-r} \geq 0$ for all $p \in [0, 1]$, $n \in \mathbb{N}$ and $r \in \{0, 1, \dots, n\}$,
- $\sum_{r=0}^n b(r; n, p) = 1$:

$$(1 - p)^n + C_n^1 p (1 - p)^{n-1} + \dots + C_n^{n-1} p^{n-1} (1 - p) + p^n = [p + (1 - p)]^n = 1$$

Binomial distribution: calculating the mean

$$\begin{aligned}
 E[b(r; n, p)] &\stackrel{\text{def.}}{=} \sum_{r=0}^n r \cdot b(r; n, p) = \\
 &= 1 \cdot C_n^1 p(1-p)^{n-1} + 2 \cdot C_n^2 p^2(1-p)^{n-2} + \dots + (n-1) \cdot C_n^{n-1} p^{n-1}(1-p) + n \cdot p^n \\
 &= p \left[C_n^1 (1-p)^{n-1} + 2 \cdot C_n^2 p(1-p)^{n-2} + \dots + (n-1) \cdot C_n^{n-1} p^{n-2}(1-p) + n \cdot p^{n-1} \right] \\
 &= np \left[(1-p)^{n-1} + C_{n-1}^1 p(1-p)^{n-2} + \dots + C_{n-1}^{n-2} p^{n-2}(1-p) + C_{n-1}^{n-1} p^{n-1} \right] \\
 &= np[p + (1-p)]^{n-1} = np
 \end{aligned}$$

Binomial distribution: calculating the variance

following www.proofwiki.org/wiki/Variance_of_Binomial_Distribution, which cites
 “Probability: An Introduction”, by Geoffrey Grimmett and Dominic Welsh,
 Oxford Science Publications, 1986

We will make use of the formula $Var[X] = E[X^2] - E^2[X]$.
 By denoting $q = 1 - p$, it follows:

$$\begin{aligned}
 E[b^2(r; n, p)] &\stackrel{\text{def.}}{=} \sum_{r=0}^n r^2 C_n^r p^r q^{n-r} = \sum_{r=0}^n r^2 \frac{n(n-1) \dots (n-r+1)}{r!} p^r q^{n-r} \\
 &= \sum_{r=1}^n r n \frac{(n-1) \dots (n-r+1)}{(r-1)!} p^r q^{n-r} = \sum_{r=1}^n r n C_{n-1}^{r-1} p^r q^{n-r} \\
 &= np \sum_{r=1}^n r C_{n-1}^{r-1} p^{r-1} q^{(n-1)-(r-1)}
 \end{aligned}$$

Binomial distribution: calculating the variance (cont'd)

By denoting $j = r - 1$ and $m = n - 1$, we'll get:

$$\begin{aligned}
 E[b^2(r; n, p)] &= np \sum_{j=0}^m (j+1) C_m^j p^j q^{m-j} \\
 &= np \left[\sum_{j=0}^m j C_m^j p^j q^{m-j} + \sum_{j=0}^m C_m^j p^j q^{m-j} \right] \\
 &= np \left[\sum_{j=0}^m j \frac{m \cdot \dots \cdot (m-j+1)}{j!} p^j q^{m-j} + \underbrace{(p+q)^m}_1 \right] \\
 &= np \left[\sum_{j=1}^m m C_{m-1}^{j-1} p^j q^{m-j} + 1 \right] = np \left[mp \sum_{j=1}^m C_{m-1}^{j-1} p^{j-1} q^{(m-1)-(j-1)} + 1 \right] \\
 &= np[(n-1)p \underbrace{(p+q)^{m-1}}_1 + 1] = np[(n-1)p + 1] = n^2 p^2 - np^2 + np
 \end{aligned}$$

Finally,

$$Var[X] = E[b^2(r; n, p)] - (E[b(r; n, p)])^2 = n^2 p^2 - np^2 + np - n^2 p^2 = np(1 - p)$$

Binomial distribution: calculating the variance

Another solution

- se demonstrează relativ ușor că orice variabilă aleatoare urmând distribuția binomială $b(r; n, p)$ poate fi văzută ca o sumă de n variabile independente care urmează distribuția Bernoulli de parametru p ; ^a
- știm (sau, se poate dovedi imediat) că varianța distribuției Bernoulli de parametru p este $p(1 - p)$;
- ținând cont de proprietatea de liniaritate a varianțelor — $Var[X_1 + X_2 + \dots + X_n] = Var[X_1] + Var[X_2] + \dots + Var[X_n]$, dacă X_1, X_2, \dots, X_n sunt variabile independente —, rezultă că $Var[X] = np(1 - p)$.

^aVezi www.proofwiki.org/wiki/Bernoulli_Process_as_Binomial_Distribution, care citează de asemenea ca sursă “Probability: An Introduction” de Geoffrey Grimmett și Dominic Welsh, Oxford Science Publications, 1986.

The Gaussian distribution: $p(X = x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x - \mu)^2}{2\sigma^2}}$

Calculating the mean

$$E[\mathcal{N}_{\mu,\sigma}(x)] \stackrel{\text{def.}}{=} \int_{-\infty}^{\infty} xp(x)dx = \frac{1}{\sqrt{2\pi}\sigma} \int_{-\infty}^{\infty} x \cdot e^{-\frac{(x - \mu)^2}{2\sigma^2}} dx$$

Using the variable transformation $v = \frac{x - \mu}{\sigma}$ will imply $x = \sigma v + \mu$ and $dx = \sigma dv$, so:

$$\begin{aligned} E[X] &= \frac{1}{\sqrt{2\pi}\sigma} \int_{-\infty}^{\infty} (\sigma v + \mu) e^{-\frac{v^2}{2}} (\sigma dv) = \frac{\sigma}{\sqrt{2\pi}\sigma} \left(\sigma \int_{-\infty}^{\infty} v e^{-\frac{v^2}{2}} dv + \mu \int_{-\infty}^{\infty} e^{-\frac{v^2}{2}} dv \right) \\ &= \frac{1}{\sqrt{2\pi}} \left(-\sigma \int_{-\infty}^{\infty} (-v) e^{-\frac{v^2}{2}} dv + \mu \int_{-\infty}^{\infty} e^{-\frac{v^2}{2}} dv \right) = \frac{1}{\sqrt{2\pi}} \left(\underbrace{-\sigma e^{-\frac{v^2}{2}} \Big|_{-\infty}^{\infty}}_{=0} + \mu \int_{-\infty}^{\infty} e^{-\frac{v^2}{2}} dv \right) \\ &= \frac{\mu}{\sqrt{2\pi}} \int_{-\infty}^{\infty} e^{-\frac{v^2}{2}} dv. \text{ The last integral is computed as shown on the next slide.} \end{aligned}$$

The Gaussian distribution: calculating the mean (Cont'd)

10.

$$\begin{aligned} \left(\int_{v=-\infty}^{\infty} e^{-\frac{v^2}{2}} dv \right)^2 &= \left(\int_{x=-\infty}^{\infty} e^{-\frac{x^2}{2}} dx \right) \cdot \left(\int_{y=-\infty}^{\infty} e^{-\frac{y^2}{2}} dy \right) = \int_{x=-\infty}^{\infty} \int_{y=-\infty}^{\infty} e^{-\frac{x^2+y^2}{2}} dy dx \\ &= \iint_{\mathbb{R}^2} e^{-\frac{x^2+y^2}{2}} dy dx \end{aligned}$$

By switching from x, y to polar coordinates r, θ (see the *Note* below), it follows:

$$\begin{aligned} \left(\int_{v=-\infty}^{\infty} e^{-\frac{v^2}{2}} dv \right)^2 &= \int_{r=0}^{\infty} \int_{\theta=0}^{2\pi} e^{-\frac{r^2}{2}} (r dr d\theta) = \int_{r=0}^{\infty} r e^{-\frac{r^2}{2}} \left(\int_{\theta=0}^{2\pi} d\theta \right) dr = \int_{r=0}^{\infty} r e^{-\frac{r^2}{2}} \theta \Big|_0^{2\pi} dr \\ &= 2\pi \int_{r=0}^{\infty} r e^{-\frac{r^2}{2}} dr = 2\pi \left(-e^{-\frac{r^2}{2}} \right) \Big|_0^{\infty} = 2\pi(0 - (-1)) = 2\pi \end{aligned}$$

Note: $x = r \cos \theta$ and $y = r \sin \theta$, with $r \geq 0$ and $\theta \in [0, 2\pi)$. Therefore, $x^2 + y^2 = r^2$, and the Jacobian matrix is

$$\frac{\partial(x, y)}{\partial(r, \theta)} = \begin{vmatrix} \frac{\partial x}{\partial r} & \frac{\partial x}{\partial \theta} \\ \frac{\partial y}{\partial r} & \frac{\partial y}{\partial \theta} \end{vmatrix} = \begin{vmatrix} \cos \theta & -r \sin \theta \\ \sin \theta & r \cos \theta \end{vmatrix} = r \cos^2 \theta + r \sin^2 \theta = r \geq 0. \text{ So, } dx dy = r dr d\theta.$$

The Gaussian distribution: calculating the variance

We will make use of the formula $\text{Var}[X] = E[X^2] - E^2[X]$.

$$E[X^2] = \int_{-\infty}^{\infty} x^2 p(x) dx = \frac{1}{\sqrt{2\pi}\sigma} \int_{-\infty}^{\infty} x^2 \cdot e^{-\frac{(x-\mu)^2}{2\sigma^2}} dx$$

Again, using the transformation $v = \frac{x-\mu}{\sigma}$ will imply $x = \sigma v + \mu$ and $dx = \sigma dv$. Therefore,

$$\begin{aligned} E[X^2] &= \frac{1}{\sqrt{2\pi}\sigma} \int_{-\infty}^{\infty} (\sigma v + \mu)^2 e^{-\frac{v^2}{2}} (\sigma dv) \\ &= \frac{\sigma}{\sqrt{2\pi}\sigma} \int_{-\infty}^{\infty} (\sigma^2 v^2 + 2\sigma\mu v + \mu^2) e^{-\frac{v^2}{2}} dv \\ &= \frac{1}{\sqrt{2\pi}} \left(\sigma^2 \int_{-\infty}^{\infty} v^2 e^{-\frac{v^2}{2}} dv + 2\sigma\mu \int_{-\infty}^{\infty} v e^{-\frac{v^2}{2}} dv + \mu^2 \int_{-\infty}^{\infty} e^{-\frac{v^2}{2}} dv \right) \end{aligned}$$

Note that we have already computed $\int_{-\infty}^{\infty} v e^{-\frac{v^2}{2}} dv = 0$ and $\int_{-\infty}^{\infty} e^{-\frac{v^2}{2}} dv = \sqrt{2\pi}$.

The Gaussian distribution: calculating the variance (Cont'd)

Therefore, we only need to compute

$$\begin{aligned}\int_{-\infty}^{\infty} v^2 e^{-\frac{v^2}{2}} dv &= \int_{-\infty}^{\infty} (-v) \left(-v e^{-\frac{v^2}{2}} \right) dv = \int_{-\infty}^{\infty} (-v) \left(e^{-\frac{v^2}{2}} \right)' dv \\ &= (-v) e^{-\frac{v^2}{2}} \Big|_{-\infty}^{\infty} - \int_{-\infty}^{\infty} (-1) e^{-\frac{v^2}{2}} dv = 0 + \int_{-\infty}^{\infty} e^{-\frac{v^2}{2}} dv = \sqrt{2\pi}.\end{aligned}$$

So,

$$E[X^2] = \frac{1}{\sqrt{2\pi}} \left(\sigma^2 \sqrt{2\pi} + 2\sigma\mu \cdot 0 + \mu^2 \sqrt{2\pi} \right) = \sigma^2 + \mu^2.$$

Finally,

$$\text{Var}[X] = E[X^2] - (E[X])^2 = (\sigma^2 + \mu^2) - \mu^2 = \sigma^2.$$

The covariance matrix Σ corresponding to a vector X made of n random variables is symmetric and positive semi-definite

a. $\text{Cov}(X)_{i,j} \stackrel{\text{def.}}{=} \text{Cov}(X_i, X_j)$, for all $i, j \in \{1, \dots, n\}$, and

$\text{Cov}(X_i, X_j) \stackrel{\text{def.}}{=} E[(X_i - E[X_i])(X_j - E[X_j])] = E[(X_j - E[X_j])(X_i - E[X_i])] = \text{Cov}(X_j, X_i)$, therefore $\text{Cov}(X)$ is a symmetric matrix.

b. We will show that $z^T \Sigma z \geq 0$ for any $z \in \mathbb{R}^n$ (seen as a column-vector):

$$\begin{aligned}
 z^T \Sigma z &= \sum_{i=1}^n z_i \left(\sum_{j=1}^n \Sigma_{ij} z_j \right) = \sum_{i=1}^n \sum_{j=1}^n (z_i \Sigma_{ij} z_j) = \sum_{i=1}^n \sum_{j=1}^n (z_i \text{Cov}[X_i, X_j] z_j) \\
 &= \sum_{i=1}^n \sum_{j=1}^n (z_i E[(X_i - E[X_i])(X_j - E[X_j])] z_j) = E \left[\sum_{i=1}^n \sum_{j=1}^n z_i (X_i - E[X_i])(X_j - E[X_j]) z_j \right] \\
 &= E \left[\left(\sum_{i=1}^n z_i (X_i - E[X_i]) \right) \left(\sum_{j=1}^n (X_j - E[X_j]) z_j \right) \right] \\
 &= E \left[\left(\sum_{i=1}^n (X_i - E[X_i]) z_i \right) \left(\sum_{j=1}^n (X_j - E[X_j]) z_j \right) \right] = E[(X - E[X])^T \cdot z]^2 \geq 0
 \end{aligned}$$

If the covariance matrix of a multi-variate Gaussian distribution is diagonal, then the density of this is equal to the product of independent univariate Gaussian densities

Let's consider $X = [X_1 \dots X_n]^T$, $\mu \in \mathbb{R}^n$ and $\Sigma \in \mathbb{S}_+^n$, where \mathbb{S}_+^n is the set of symmetric positive definite matrices (which implies $|\Sigma| \neq 0$ and $(x - \mu)^T \Sigma^{-1} (x - \mu) > 0$, therefore $-\frac{1}{2}(x - \mu)^T \Sigma^{-1} (x - \mu) < 0$, for any $x \in \mathbb{S}^n$, $x \neq \mu$).

The probability density function of a multi-variate Gaussian distribution of parameters μ and Σ is:

$$p(x; \mu, \Sigma) = \frac{1}{(2\pi)^{n/2} |\Sigma|^{1/2}} \exp \left(-\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right),$$

Notation: $X \sim \mathcal{N}(\mu, \Sigma)$.

We will make the **proof** for $n = 2$ (generalization to $n > 2$ will be easy):

$$x = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \quad \mu = \begin{bmatrix} \mu_1 \\ \mu_2 \end{bmatrix} \quad \Sigma = \begin{bmatrix} \sigma_1^2 & 0 \\ 0 & \sigma_2^2 \end{bmatrix}$$

Note: It is easy to show that if $\Sigma \in \mathbb{S}_+^n$ is diagonal, the elements on the principal diagonal Σ are indeed strictly positive. (It is enough to consider $z = (1, 0)$ and respectively $z = (0, 1)$ in formula for *positive-definiteness* of Σ .) This is why we wrote these elements of σ as σ_1^2 and σ_2^2 .

A property of multi-variate Gaussians whose covariance matrices are diagonal (Cont'd)

$$\begin{aligned}
 p(x; \mu, \Sigma) &= \frac{1}{2\pi \begin{vmatrix} \sigma_1^2 & 0 \\ 0 & \sigma_2^2 \end{vmatrix}^{\frac{1}{2}}} \exp \left(-\frac{1}{2} \begin{bmatrix} x_1 - \mu_1 \\ x_2 - \mu_2 \end{bmatrix}^T \begin{bmatrix} \sigma_1^2 & 0 \\ 0 & \sigma_2^2 \end{bmatrix}^{-1} \begin{bmatrix} x_1 - \mu_1 \\ x_2 - \mu_2 \end{bmatrix} \right) \\
 &= \frac{1}{2\pi \sigma_1 \sigma_2} \exp \left(-\frac{1}{2} \begin{bmatrix} x_1 - \mu_1 \\ x_2 - \mu_2 \end{bmatrix}^T \begin{bmatrix} \frac{1}{\sigma_1^2} & 0 \\ 0 & \frac{1}{\sigma_2^2} \end{bmatrix} \begin{bmatrix} x_1 - \mu_1 \\ x_2 - \mu_2 \end{bmatrix} \right) \\
 &= \frac{1}{2\pi \sigma_1 \sigma_2} \exp \left(-\frac{1}{2} \begin{bmatrix} x_1 - \mu_1 \\ x_2 - \mu_2 \end{bmatrix}^T \begin{bmatrix} \frac{1}{\sigma_1^2}(x_1 - \mu_1) \\ \frac{1}{\sigma_2^2}(x_2 - \mu_2) \end{bmatrix} \right) \\
 &= \frac{1}{2\pi \sigma_1 \sigma_2} \exp \left(-\frac{1}{2\sigma_1^2}(x_1 - \mu_1)^2 - \frac{1}{2\sigma_2^2}(x_2 - \mu_2)^2 \right) \\
 &= p(x_1; \mu_1, \sigma_1^2) p(x_2; \mu_2, \sigma_2^2).
 \end{aligned}$$

Random Variables:


Some exercises

Exemplifying
the computation of *expected values* for random variables
and the [use of] *sensitivity* and *specificity* of a test
in a real-world application

CMU, 2009 fall, Geoff Gordon, HW1, pr. 2

There is a disease which affects 1 in 500 people. A 100.00 dollar blood **test** can help reveal whether a person has the disease. A positive outcome indicates that the person may have the disease.

The test has perfect **sensitivity** (*true positive rate*), i.e., a person who has the disease tests positive 100% of the time. However, the test has 99% **specificity** (*true negative rate*), i.e., a healthy person tests positive 1% of the time.



$$\text{sensitivity (or: recall): } \frac{tp}{tp + fn}$$

$$\text{specificity: } \frac{tn}{tn + fp}$$

a. A randomly selected individual is tested and the result is positive.

What is the *probability* of the individual having the disease?

b. There is a second more expensive test which costs 10, 000.00 dollars but is exact with 100% *sensitivity* and *specificity*.

If we require all people who test positive with the less expensive test to be tested with the more expensive test, what is the *expected cost* to check whether an individual has the disease?

c. A pharmaceutical company is attempting to decrease the cost of the second (perfect) test.

How much would it have to make the second test cost, so that the first test is no longer needed? That is, at what cost is it cheaper simply to use the perfect test alone, instead of screening with the cheaper test as described in part *b*?

Answer:

Let's define the following *random variables*:

B : $\begin{cases} 1/\text{true} & \text{for persons affected by that disease,} \\ 0/\text{false} & \text{otherwise;} \end{cases}$

T_1 : the result of the first test: + (in case of disease) or – (otherwise);

T_2 : the result of the second test: again + or –.

Known facts:

$$P(B) = \frac{1}{500}$$

$$P(T_1 = + \mid B) = 1, \quad P(T_1 = + \mid \bar{B}) = \frac{1}{100},$$

$$P(T_2 = + \mid B) = 1, \quad P(T_2 = + \mid \bar{B}) = 0$$

a.

$$\begin{aligned} P(B \mid T_1 = +) &\stackrel{TBayes}{=} \frac{P(T_1 = + \mid B) \cdot P(B)}{P(T_1 = + \mid B) \cdot P(B) + P(T_1 = + \mid \bar{B}) \cdot P(\bar{B})} \\ &= \frac{1 \cdot \frac{1}{500}}{1 \cdot \frac{1}{500} + \frac{1}{100} \cdot \frac{499}{500}} = \frac{100}{599} \approx 0.1669 \end{aligned}$$

b.

Let's consider a new random variable:

$$C = \begin{cases} c_1 & \text{if the person only takes the first test} \\ c_1 + c_2 & \text{if the person takes the two tests} \end{cases}$$

$$\Rightarrow P(C = c_1) = P(T_1 = -) \text{ and } P(C = c_1 + c_2) = P(T_1 = +)$$

$$\begin{aligned} \Rightarrow E[C] &= c_1 \cdot (1 - P(T_1 = +)) + (c_1 + c_2) \cdot P(T_1 = +) \\ &= c_1 - c_1 \cdot P(T_1 = +) + c_1 \cdot P(T_1 = +) + c_2 \cdot P(T_1 = +) \\ &= c_1 + c_2 \cdot P(T_1 = +) \\ &= 100 + 10000 \cdot \frac{599}{50000} = 219.8 \approx 220\$ \end{aligned}$$

Note: Here above we used

$$\begin{aligned} P(T_1 = +) &\stackrel{\text{total probability form.}}{=} P(T_1 = + | B) \cdot P(B) + P(T_1 = + | \bar{B}) \cdot P(\bar{B}) \\ &= 1 \cdot \frac{1}{500} + \frac{1}{100} \cdot \frac{499}{500} = \frac{599}{50000} = 0.01198 \end{aligned}$$

c.

c_n ^{not.} = the new price for the second test (T'_2)

$$\begin{aligned} c_n \leq E[C'] &= c_1 \cdot P(C = c_1) + (c_1 + c_n) \cdot P(C = c_1 + c_n) \\ &= c_1 + c_n \cdot P(T_1 = +) = 100 + c_n \cdot \frac{599}{50000} \end{aligned}$$

$$c_n = 100 + c_n \cdot 0.01198 \Rightarrow c_n \approx 101.2125.$$

Using the Central Limit Theorem (the i.i.d. version)
to compute the *real error* of a classifier
CMU, 2008 fall, Eric Xing, HW3, pr. 3.3

Chris recently adopts a new (binary) classifier to filter email spams. He wants to quantitatively evaluate how good the classifier is.

He has a small dataset of 100 emails on hand which, you can assume, are randomly drawn from all emails.

He tests the classifier on the 100 emails and gets 83 classified correctly, so the error rate on the small dataset is 17%.

However, the number on 100 samples could be either higher or lower than the real error rate just by chance.

With a confidence level of 95%, what is likely to be the range of the real error rate? Please write down all important steps.

(Hint: You need some approximation in this problem.)

Notations:

Let X_i , $i = 1, \dots, n = 100$ be defined as:

$X_i = 1$ if the email i was incorrectly classified, and 0 otherwise;

$$E[X_i] \stackrel{\text{not.}}{=} \mu \stackrel{\text{not.}}{=} e_{\text{real}}; \quad \text{Var}(X_i) \stackrel{\text{not.}}{=} \sigma^2$$

$$e_{\text{sample}} \stackrel{\text{not.}}{=} \frac{X_1 + \dots + X_n}{n} = 0.17$$

$$Z_n = \frac{X_1 + \dots + X_n - n\mu}{\sqrt{n} \sigma} \quad (\text{the standardized form of } X_1 + \dots + X_n)$$

Key insight:

Calculating the real error of the classifier (more exactly, a symmetric interval around the real error $p \stackrel{\text{not.}}{=} \mu$) with a “confidence” of 95% amounts to finding $a > 0$ such that $P(|Z_n| \leq a) \geq 0.95$.

Calculus:

$$\begin{aligned}
 |Z_n| \leq a &\Leftrightarrow \left| \frac{X_1 + \dots + X_n - n\mu}{\sqrt{n} \sigma} \right| \leq a \Leftrightarrow \left| \frac{X_1 + \dots + X_n - n\mu}{n\sigma} \right| \leq \frac{a}{\sqrt{n}} \\
 &\Leftrightarrow \left| \frac{X_1 + \dots + X_n - n\mu}{n} \right| \leq \frac{a\sigma}{\sqrt{n}} \Leftrightarrow \left| \frac{X_1 + \dots + X_n}{n} - \mu \right| \leq \frac{a\sigma}{\sqrt{n}} \\
 &\Leftrightarrow |e_{sample} - e_{real}| \leq \frac{a\sigma}{\sqrt{n}} \Leftrightarrow |e_{real} - e_{sample}| \leq \frac{a\sigma}{\sqrt{n}} \\
 &\Leftrightarrow -\frac{a\sigma}{\sqrt{n}} \leq e_{real} - e_{sample} \leq \frac{a\sigma}{\sqrt{n}} \\
 &\Leftrightarrow e_{sample} - \frac{a\sigma}{\sqrt{n}} \leq e_{real} \leq e_{sample} + \frac{a\sigma}{\sqrt{n}} \\
 &\Leftrightarrow e_{real} \in \left[e_{sample} - \frac{a\sigma}{\sqrt{n}}, e_{sample} + \frac{a\sigma}{\sqrt{n}} \right]
 \end{aligned}$$

Important facts:

The Central Limit Theorem: $Z_n \rightarrow N(0; 1)$

Therefore, $P(|Z_n| \leq a) \approx P(|X| \leq a) = \Phi(a) - \Phi(-a)$, where $X \sim N(0; 1)$ and Φ is the cumulative function distribution of $N(0; 1)$.

Calculus:

$$\Phi(-a) + \Phi(a) = 1 \Rightarrow P(|Z_n| \leq a) = \Phi(a) - \Phi(-a) = 2\Phi(a) - 1$$

$$P(|Z_n| \leq a) = 0.95 \Leftrightarrow 2\Phi(a) - 1 = 0.95 \Leftrightarrow \Phi(a) = 0.975 \Leftrightarrow a \cong 1.97 \text{ (see } \Phi \text{ table)}$$

Finally:

$\sigma^2 \stackrel{\text{not.}}{=} \text{Var}_{real} \approx \text{Var}_{sample}$ due to the above theorem, and

$\text{Var}_{sample} = e_{sample}(1 - e_{sample})$ because X_i are Bernoulli variables.

$$\Rightarrow \frac{a\sigma}{\sqrt{n}} = 1.97 \cdot \frac{\sqrt{0.17(1 - 0.17)}}{\sqrt{100}} \cong 0.07$$

$$\begin{aligned} |e_{real} - e_{sample}| \leq 0.07 &\Leftrightarrow |e_{real} - 0.17| \leq 0.07 \Leftrightarrow -0.07 \leq e_{real} - 0.17 \leq 0.07 \\ &\Leftrightarrow e_{real} \in [0.10, 0.24] \end{aligned}$$

Estimating the parameters of some probability distributions: Exemplifications

Estimating the parameter of the Bernoulli
distribution,
in sense MLE and MAP

CMU, 2015 spring, Tom Mitchell, Nina Balcan, HW2, pr. 2

Suppose we observe the values of n i.i.d. (independent, identically distributed) random variables X_1, \dots, X_n drawn from a single Bernoulli distribution with parameter θ . In other words, for each X_i , we know that

$$P(X_i = 1) = \theta \quad \text{and} \quad P(X_i = 0) = 1 - \theta.$$

Our *goal* is to estimate the value of θ from the observed values of X_1, \dots, X_n .

Maximum Likelihood Estimation

For any hypothetical value $\hat{\theta}$, we can compute the probability of observing the outcome X_1, \dots, X_n if the true parameter value θ were equal to $\hat{\theta}$.

This probability of the observed data is often called the *data likelihood*, and the function $L(\hat{\theta})$ that maps each $\hat{\theta}$ to the corresponding likelihood is called the *likelihood function*.

A natural way to estimate the unknown parameter θ is to choose the $\hat{\theta}$ that maximizes the likelihood function. Formally,

$$\hat{\theta}_{MLE} = \operatorname{argmax}_{\hat{\theta}} L(\hat{\theta}).$$

a. Write a formula for the likelihood function, $L(\hat{\theta})$.

Your function should depend on the random variables X_1, \dots, X_n and the hypothetical parameter $\hat{\theta}$.

Does the likelihood function depend on the order of the random variables?

Solution:

Since the X_i are independent, we have

$$\begin{aligned} L(\hat{\theta}) &= P_{\hat{\theta}}(X_1, \dots, X_n) = \prod_{i=1}^n P_{\hat{\theta}}(X_i) = \prod_{i=1}^n (\hat{\theta}^{X_i} \cdot (1 - \hat{\theta})^{1-X_i}) \\ &= \hat{\theta}^{\#\{X_i=1\}} \cdot (1 - \hat{\theta})^{\#\{X_i=0\}}, \end{aligned}$$

where $\#\{\cdot\}$ counts the number of X_i for which the condition in braces holds true. In the third equality we used the trick $X_i = \mathbb{I}\{X_i = 1\}$.

The likelihood function does not depend on the order of the data.

b. Suppose that $n = 10$ and the data set contains six 1s and four 0s.


Write a short computer program that plots the likelihood function of this data.

For the plot, the x -axis should be $\hat{\theta}$ and the y -axis $L(\hat{\theta})$. Scale your y -axis so that you can see some variation in its value.

Estimate $\hat{\theta}_{MLE}$ by marking on the x -axis the value of $\hat{\theta}$ that maximizes the likelihood.

Solution:

MLE; $n = 10$, six 1s, four 0s



c. Find a closed-form formula for $\hat{\theta}_{MLE}$, the MLE estimate of $\hat{\theta}$. Does the closed form agree with the plot?

Solution:

Let's consider $l(\theta) = \ln(L(\theta))$. Since the \ln function is increasing, the $\hat{\theta}$ that maximizes the log-likelihood is the same as the θ that maximizes the likelihood. Using the properties of the \ln function, we can rewrite $l(\hat{\theta})$ as follows:

$$l(\hat{\theta}) = \ln(\hat{\theta}^{n_1} \cdot (1 - \hat{\theta})^{n_0}) = n_1 \ln(\hat{\theta}) + n_0 \ln(1 - \hat{\theta}).$$

Assuming that $\hat{\theta} \neq 0$ and $\hat{\theta} \neq 1$, the first and second derivatives of l are given by


$$l'(\hat{\theta}) = \frac{n_1}{\hat{\theta}} - \frac{n_0}{1 - \hat{\theta}} \quad \text{and} \quad l''(\hat{\theta}) = -\frac{n_1}{\hat{\theta}^2} - \frac{n_0}{(1 - \hat{\theta})^2}$$

Since $l''(\hat{\theta})$ is always negative, the l function is concave, and we can find its maximizer by solving the equation $l'(\theta) = 0$. The solution to this equation is given by $\hat{\theta}_{MLE} = \frac{n_1}{n_1 + n_0}$.


d. Create three more likelihood plots: one where $n = 5$ and the data set contains three 1s and two 0s; one where $n = 100$ and the data set contains sixty 1s and forty 0s; and one where $n = 10$ and there are five 1s and five 0s.

Solution:

MLE; $n = 5$, three 1s, two 0s




MLE; $n = 100$, sixty 1s, forty 0s



Solution (to part d.):

MLE; $n = 10$, five 1s, five 0s



e. Describe how the likelihood functions and maximum likelihood estimates compare for the different data sets.

Solution (to part e.):

The MLE is equal to the proportion of 1s observed in the data, so for the first three plots the MLE is always at 0.6, while for the last plot it is at 0.5.

As the number of samples n increases, the likelihood function gets more peaked at its maximum value, and the values it takes on decrease.

Maximum a Posteriori Probability Estimation

In the maximum likelihood estimate, we treated the true parameter value θ as a fixed (non-random) number. In cases where we have some prior knowledge about θ , it is useful to treat θ itself as a random variable, and express our prior knowledge in the form of a prior probability distribution over θ .

For *example*, suppose that the X_1, \dots, X_n are generated in the following way:

- First, the value of θ is drawn from a given prior probability distribution
- Second, X_1, \dots, X_n are drawn independently from a Bernoulli distribution using this value for θ .

Since both θ and the sequence X_1, \dots, X_n are random, they have a joint probability distribution. In this setting, a natural way to estimate the value of θ is to simply choose its most probable value given its prior distribution plus the observed data X_1, \dots, X_n .

$$\hat{\theta}_{MAP} = \operatorname{argmax}_{\hat{\theta}} P(\theta = \hat{\theta} | X_1, \dots, X_n).$$

This is called the maximum a posteriori probability (MAP) estimate of θ .

Using Bayes rule, we can rewrite the posterior probability as follows:

$$P(\theta = \hat{\theta} | X_1, \dots, X_n) = \frac{P(X_1, \dots, X_n | \theta = \hat{\theta}) P(\theta = \hat{\theta})}{P(X_1, \dots, X_n)}.$$

Since the probability in the denominator does not depend on $\hat{\theta}$, the MAP estimate is given by

$$\begin{aligned} \hat{\theta}_{MAP} &= \operatorname{argmax}_{\hat{\theta}} P(X_1, \dots, X_n | \theta = \hat{\theta}) P(\theta = \hat{\theta}) \\ &= \operatorname{argmax}_{\hat{\theta}} L(\hat{\theta}) P(\theta = \hat{\theta}). \end{aligned}$$

In words, the MAP estimate for θ is the value $\hat{\theta}$ that maximizes the likelihood function multiplied by the prior distribution on θ . The MAP estimate for θ is given by

$$\hat{\theta}_{MAP} = \operatorname{argmax}_{\hat{\theta}} L(\hat{\theta}) p(\hat{\theta}).$$

We will consider a $Beta(3,3)$ prior distribution for θ , which has the density function given by $p(\hat{\theta}) = \frac{\hat{\theta}^2(1-\hat{\theta})^2}{B(3,3)}$, where $B(\alpha, \beta)$ is the beta function and $B(3,3) \approx 0.0333$.


f. Suppose, as in part c, that $n = 10$ and we observed six 1s and four 0s.

Write a short computer program that plots the function $\hat{\theta} \mapsto L(\hat{\theta})p(\hat{\theta})$ for the same values of $\hat{\theta}$ as in part c.

Estimate $\hat{\theta}_{MAP}$ by marking on the x -axis the value of $\hat{\theta}$ that maximizes the function.

Solution:

MAP; $n = 10$, six 1s, four 0s; $Beta(3,3)$



g. Find a closed form formula for $\hat{\theta}_{MAP}$, the MAP estimate of $\hat{\theta}$. Does the closed form agree with the plot?

Solution:

As in the case of the MLE, we will apply the \ln function before finding the maximizer. We want to maximize the function

$$l(\hat{\theta}) = \ln(L(\hat{\theta}) \cdot p(\hat{\theta})) = \ln(\hat{\theta}^{n_1+2} \cdot (1 - \hat{\theta})^{n_0+2}) - \ln(B(3, 3)).$$

The normalizing constant for the prior appears as an additive constant and therefore the first and second derivatives are identical to those in the case of the MLE (except with $n_1 + 2$ and $n_0 + 2$ instead of n_1 and n_0 , respectively).

It follows that the closed form formula for the MAP estimate is given by

$$\hat{\theta}_{MAP} = \frac{n_1 + 2}{n_1 + n_0 + 4}$$

h. Compare the MAP estimate to the MLE computed from the same data in part c. Briefly explain any significant difference.

Solution:

The MAP estimate is equal to the MLE with four additional virtual random variables, two that are equal to 1, and two that are equal to 0. This pulls the value of the MAP estimate closer to the value 0.5, which is why $\hat{\theta}_{MAP}$ is smaller than $\hat{\theta}_{MLE}$.

i. Comment on the relationship between the MAP and MLE estimates as n goes to infinity, while the ratio $\#\{X_i = 1\}/\#\{X_i = 0\}$ remains constant.

Solution:

As n goes to infinity, the influence of the 4 virtual random variables diminishes, and the two estimators become equal.

**The Gaussian [uni-variate] distribution:
estimating μ when σ^2 is known**

CMU, 2011 fall, Tom Mitchell, Aarti Singh, HW2, pr. 1

CMU, 2010 fall, Ziv Bar-Joseph, HW1, pr. 1.2-3

Assume we have n samples, x_1, \dots, x_n , independently drawn from a normal distribution with *known* variance σ^2 and *unknown* mean μ .

a. Derive the MLE estimator for the mean μ .

Solution:

$$P(x_1, \dots, x_n | \mu) = \prod_{i=1}^n P(x_i | \mu) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x_i - \mu)^2}{2\sigma^2}}$$

$$\Rightarrow \ln P(x_1, \dots, x_n | \mu) = \sum_{i=1}^n \left(\ln \frac{1}{\sqrt{2\pi}\sigma} - \frac{(x_i - \mu)^2}{2\sigma^2} \right)$$

$$\Rightarrow \frac{\partial}{\partial \mu} \ln P(x_1, \dots, x_n | \mu) = \sum_{i=1}^n \frac{x_i - \mu}{\sigma^2}$$

$$\frac{\partial}{\partial \mu} \ln P(x_1, \dots, x_n | \mu) = 0 \Leftrightarrow \sum_{i=1}^n \frac{x_i - \mu}{\sigma^2} = 0 \Leftrightarrow \sum_{i=1}^n (x_i - \mu) = 0 \Leftrightarrow \sum_{i=1}^n x_i = n\mu$$

$$\Rightarrow \mu_{MLE} = \frac{\sum_{i=1}^n x_i}{n}$$

b. Show that $E[\mu_{MLE}] = \mu$.

Solution:

The sample x_1, \dots, x_n can be seen as the realization of n independent random variables X_1, \dots, X_n of Gaussian distribution of mean μ and variance σ^2 . Then, due to the property of linearity for the expectation of random variables, we get:

$$E[\mu_{MLE}] = E\left[\frac{X_1 + \dots + X_n}{n}\right] = \frac{E[X_1] + \dots + E[X_n]}{n} = \frac{n\mu}{n} = \mu$$

Therefore, the μ_{MLE} estimator is unbiased.

c. What is $Var[\mu_{MLE}]$?

Solution:

$$Var[\mu_{MLE}] = Var\left[\frac{1}{n} \sum_{i=1}^n X_i\right] \stackrel{i.i.d.}{=} \frac{1}{n^2} \sum_{i=1}^n Var[X_i] = n \frac{1}{n^2} Var[X_1] = \frac{\sigma^2}{n}$$

Therefore, $Var[\mu_{MLE}] \rightarrow 0$ as $n \rightarrow \infty$.

d. Now derive the MAP estimator for the mean μ . Assume that the prior distribution for the mean is itself a normal distribution with mean ν and variance β^2 .

Solution 1:

$$P(\mu|x_1, \dots, x_n) \stackrel{T. Bayes}{=} \frac{P(x_1, \dots, x_n|\mu) P(\mu)}{P(x_1, \dots, x_n)} \quad (1)$$

$$= \frac{\left(\prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} \right) e^{-\frac{(x_i - \mu)^2}{2\sigma^2}} \frac{1}{\sqrt{2\pi}\beta} e^{-\frac{(\mu - \nu)^2}{2\beta^2}}}{C} \quad (2)$$

where $C \stackrel{not.}{=} P(x_1, \dots, x_n)$.

$$\Rightarrow \ln P(\mu|x_1, \dots, x_n) = - \sum_{i=1}^n \left(\ln \sqrt{2\pi}\sigma + \frac{(x_i - \mu)^2}{2\sigma^2} \right) - \ln \sqrt{2\pi}\beta - \frac{(\mu - \nu)^2}{2\beta^2} - \ln C$$

$$\Rightarrow \frac{\partial}{\partial \mu} \ln P(\mu|x_1, \dots, x_n) = \sum_{i=1}^n \frac{x_i - \mu}{\sigma^2} - \frac{\mu - \nu}{\beta^2}$$

$$\frac{\partial}{\partial \mu} \ln P(\mu|x_1, \dots, x_n) = 0 \Leftrightarrow \sum_{i=1}^n \frac{x_i - \mu}{\sigma^2} = \frac{\mu - \nu}{\beta^2} \Leftrightarrow \mu \left(\frac{1}{\beta^2} + \frac{n}{\sigma^2} \right) = \frac{\sum_{i=1}^n x_i}{\sigma^2} + \frac{\nu}{\beta^2}$$

$$\Rightarrow \mu_{MAP} = \frac{\sigma^2 \nu + \beta^2 \sum_{i=1}^n x_i}{\sigma^2 + n\beta^2}$$

Solution 2:

Instead of computing the derivative of the posterior distribution $P(\mu|x_1, \dots, x_n)$, we will first show that the right hand side of (2) is itself a Gaussian, and then we will use the fact that the mean of a Gaussian is where it achieves its maximum value.

$$\begin{aligned}
 P(\mu|x_1, \dots, x_n) &= \frac{1}{C} \left(\prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} \right) e^{-\frac{(x_i - \mu)^2}{2\sigma^2}} \frac{1}{\sqrt{2\pi}\beta} e^{-\frac{(\mu - \nu)^2}{2\beta^2}} \\
 &= \text{const} \cdot e^{-\frac{\sum_{i=1}^n (x_i - \mu)^2}{2\sigma^2} - \frac{(\mu - \nu)^2}{2\beta^2}} \\
 &= \text{const} \cdot e^{-\frac{\beta^2 \sum_{i=1}^n (x_i - \mu)^2 + \sigma^2 (\mu - \nu)^2}{2\sigma^2 \beta^2}} \\
 &= \text{const} \cdot e^{-\frac{n\beta^2 + \sigma^2}{2\sigma^2 \beta^2} \mu^2 + \frac{\beta^2 \sum_{i=1}^n x_i + \nu \sigma^2}{\sigma^2 \beta^2} \mu - \frac{\beta^2 \sum_{i=1}^n x_i + \nu^2 \sigma^2}{2\sigma^2 \beta^2}}
 \end{aligned}$$

$$\begin{aligned}
P(\mu|x_1, \dots, x_n) &= \\
&= \text{const} \cdot \exp \left(- \frac{\mu^2 - 2\mu \frac{\beta^2 \sum_{i=1}^n x_i + \nu \sigma^2}{n\beta^2 + \sigma^2} + \frac{\beta^2 \sum_{i=1}^n x_i + \nu^2 \sigma^2}{n\beta^2 + \sigma^2}}{\frac{2\sigma^2 \beta^2}{n\beta^2 + \sigma^2}} \right) \\
&= \text{const} \cdot \exp \left(- \frac{(\mu - \frac{\beta^2 \sum_{i=1}^n x_i + \nu \sigma^2}{n\beta^2 + \sigma^2})^2 - \left(\frac{\beta^2 \sum_{i=1}^n x_i + \nu \sigma^2}{n\beta^2 + \sigma^2} \right)^2 + \frac{\beta^2 \sum_{i=1}^n x_i + \nu^2 \sigma^2}{n\beta^2 + \sigma^2}}{2 \frac{\sigma^2 \beta^2}{n\beta^2 + \sigma^2}} \right) \\
&= \text{const} \cdot \exp \left(- \frac{\left(\mu - \frac{\beta^2 \sum_{i=1}^n x_i + \nu \sigma^2}{n\beta^2 + \sigma^2} \right)^2}{2 \frac{\sigma^2 \beta^2}{n\beta^2 + \sigma^2}} \right) \cdot \exp \left(\frac{\left(\frac{\beta^2 \sum_{i=1}^n x_i + \nu \sigma^2}{n\beta^2 + \sigma^2} \right)^2 - \frac{\beta^2 \sum_{i=1}^n x_i + \nu^2 \sigma^2}{n\beta^2 + \sigma^2}}{2 \frac{\sigma^2 \beta^2}{n\beta^2 + \sigma^2}} \right) \\
&= \text{const}' \cdot \exp \left(- \frac{\left(\mu - \frac{\beta^2 \sum_{i=1}^n x_i + \nu \sigma^2}{n\beta^2 + \sigma^2} \right)^2}{2 \frac{\sigma^2 \beta^2}{n\beta^2 + \sigma^2}} \right)
\end{aligned}$$

The exp term in the last equality being a Gaussian of mean $\frac{\beta^2 \sum_{i=1}^n x_i + \nu \sigma^2}{n\beta^2 + \sigma^2}$ and variance $\frac{\sigma^2 \beta^2}{n\beta^2 + \sigma^2}$, it follows that its maximum is obtained for $\mu = \frac{\beta^2 \sum_{i=1}^n x_i + \nu \sigma^2}{n\beta^2 + \sigma^2} = \mu_{MAP}$.

e. Please comment on what happens to the MLE and MAP estimators for the mean μ as the number of samples n goes to infinity.

Solution:

$$\begin{aligned}\mu_{MLE} &= \frac{\sum_{i=1}^n x_i}{n} \\ \mu_{MAP} &= \frac{\sigma^2 \nu + \beta^2 \sum_{i=1}^n x_i}{\sigma^2 + n\beta^2} = \frac{\sigma^2 \nu}{\sigma^2 + n\beta^2} + \frac{\beta^2 \sum_{i=1}^n x_i}{\sigma^2 + n\beta^2} \\ &= \frac{\sigma^2 \nu}{\sigma^2 + n\beta^2} + \frac{\frac{1}{n} \sum_{i=1}^n x_i}{1 + \frac{\sigma^2}{n\beta^2}} = \frac{\sigma^2 \nu}{\sigma^2 + n\beta^2} + \frac{\mu_{MLE}}{1 + \frac{\sigma^2}{n\beta^2}}\end{aligned}$$

$$n \rightarrow \infty \Rightarrow \frac{\sigma^2 \nu}{\sigma^2 + n\beta^2} \rightarrow 0 \text{ and } \frac{\sigma^2}{n\beta^2} \rightarrow 0 \Rightarrow \mu_{MAP} \rightarrow \mu_{MLE}$$

**The Gaussian [uni-variate] distribution:
estimating σ^2 when $\mu = 0$**

CMU, 2009 spring, Ziv Bar-Joseph, HW1, pr. 2.1

Let X be a random variable distributed according to a Normal distribution with 0 mean, and σ^2 variance, i.e. $X \sim N(0, \sigma^2)$.

a. Find the maximum likelihood estimate for σ^2 , i.e. σ_{MLE}^2 .

Solution:

Let X_1, X_2, \dots, X_n be drawn i.i.d. $\sim N(0, \sigma^2)$. Let f be the density function corresponding to X . Then we can write the likelihood function as:

$$\begin{aligned} L(\sigma^2 | X_1, X_2, \dots, X_n) &= \prod_{i=1}^n f(X_i; \mu = 0, \sigma^2) \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \right)^n \prod_{i=1}^n \exp \left(-\frac{(X_i - 0)^2}{2\sigma^2} \right) = \left(\frac{1}{\sqrt{2\pi}\sigma} \right)^n \exp \left(-\frac{\sum_{i=1}^n X_i^2}{2\sigma^2} \right) \\ \Rightarrow \ln L &= \text{constant} - \frac{n}{2} \ln \sigma^2 - \frac{1}{2\sigma^2} \sum_{i=1}^n X_i^2 \\ \Rightarrow \frac{\partial \ln L}{\partial \sigma^2} &= -\frac{n}{2\sigma^2} + \frac{1}{2\sigma^4} \sum_{i=1}^n X_i^2. \text{ Therefore, } \frac{\partial \ln L}{\partial \sigma^2} = 0 \Leftrightarrow \sigma_{MLE}^2 = \frac{1}{n} \sum_{i=1}^n X_i^2 \end{aligned}$$

Note: It can be easily shown that $L(\sigma^2 | X_1, X_2, \dots, X_n)$ indeed reaches its maximum for $\sigma^2 = \frac{1}{n} \sum_{i=1}^n X_i^2$.

b. Is the estimator you obtained biased?

Solution:

It is unbiased, since:

$$\begin{aligned} E\left[\frac{1}{n} \sum_{i=1}^n X_i^2\right] &= \frac{n}{n} E[X^2] && \text{since i.i.d.} \\ &= \text{Var}[X] + (E[X])^2 \\ &= \text{Var}[X] = \sigma^2 && \text{since } E[X] = 0 \end{aligned}$$

**The Gaussian [uni-variate] distribution:
estimating σ^2 (without restrictions on μ)**

CMU, 2010 fall, Ziv Bar-Joseph, HW1, pr. 2.1.1-2

Let $\mathbf{x} = (x_1, \dots, x_n)$ be observed i.i.d. samples from a Gaussian distribution $N(x|\mu, \sigma^2)$.

a. Derive σ_{MLE}^2 , the MLE for σ^2 .

Solution:

The p.d.f. for $N(x|\mu, \sigma^2)$ has the form $f(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$.

The log likelihood function of the data \mathbf{x} is:

$$\begin{aligned} \ln \mathcal{L}(\mathbf{x} \mid \mu, \sigma^2) &= \ln \prod_{i=1}^n f(x_i) = \sum_{i=1}^n \left(-\frac{1}{2} \ln(2\pi\sigma^2) - \frac{(x_i - \mu)^2}{2\sigma^2} \right) \\ &= -\frac{n}{2} \ln(2\pi\sigma^2) - \frac{1}{\sigma^2} \sum_{i=1}^n (x_i - \mu)^2 \end{aligned}$$

The partial derivative of $\ln \mathcal{L}$ w.r.t. σ^2 : $\frac{\partial \ln \mathcal{L}(\mathbf{x} \mid \mu, \sigma^2)}{\partial \sigma^2} = -\frac{n}{2\sigma^2} + \frac{1}{\sigma^4} \sum_{i=1}^n (x_i - \mu)^2$.

Solving the equation $\frac{\partial \ln \mathcal{L}(\mathbf{x} \mid \mu, \sigma^2)}{\partial \sigma^2} = 0$, we get: $\sigma_{MLE}^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \mu_{MLE})^2$.

Note that we had took into account the optimal value of μ (see problem CMU, 2011 fall, T. Mitchell, A. Singh, HW2, pr. 1)

b. Show that $E[\sigma_{MLE}] = \frac{n-1}{n}\sigma^2$.

Solution:

$$\begin{aligned}
 E[\sigma_{MLE}] &= E\left[\frac{1}{n}\sum_{i=1}^n(x_i - \mu_{MLE})^2\right] = E[(x_1 - \mu_{MLE})^2] = E\left[\left(x_1 - \frac{1}{n}\sum_{i=1}^n x_i\right)^2\right] \\
 &= E\left[x_1^2 - \frac{2}{n}x_1\sum_{i=1}^n x_i + \frac{1}{n^2}\left(\sum_{i=1}^n x_i\right)^2\right] \\
 &= E\left[x_1^2 - \frac{2}{n}x_1\sum_{i=1}^n x_i + \frac{1}{n^2}\sum_{i=1}^n x_i^2 + \frac{2}{n^2}\sum_{i<j} x_i x_j\right] \\
 &= E[x_1^2] + \frac{1}{n^2}\sum_{i=1}^n E[x_i^2] - \frac{2}{n}\sum_{i=1}^n E[x_1 x_i] + \frac{2}{n^2}\sum_{i<j} E[x_i x_j] \\
 &= E[x_1^2] + \frac{1}{n^2}nE[x_1^2] - \frac{2}{n}E[x_1^2] - \frac{2}{n}(n-1)E[x_1 x_2] + \frac{2}{n^2}\frac{n(n-1)}{2}E[x_1 x_2] \\
 &= \frac{n-1}{n}E[x_1^2] - \frac{n-1}{n}E[x_1 x_2]
 \end{aligned}$$

$$\sigma^2 = \text{Var}(x_1) = E[x_1^2] - (E[x_1])^2 = E[x_1^2] - \mu^2 \Rightarrow E[x_1^2] = \sigma^2 + \mu^2$$

**Because x_1 and x_2 are independent, it follows that $\text{Cov}(x_1, x_2) = 0$.
Therefore,**

$$\begin{aligned} 0 &= \text{Cov}(x_1, x_2) = E[(x_1 - E[x_1])(x_2 - E[x_2])] = E[(x_1 - \mu)(x_2 - \mu)] \\ &= E[x_1 x_2] - \mu E[x_1 + x_2] + \mu^2 = E[x_1 x_2] - \mu(E[x_1] + E[x_2]) + \mu^2 \\ &= E[x_1 x_2] - \mu(2\mu) + \mu^2 = E[x_1 x_2] - \mu^2 \end{aligned}$$

So, $E[x_1 x_2] = \mu^2$.

By substituting $E[x_1^2] = \sigma^2 + \mu^2$ and $E[x_1 x_2] = \mu^2$ into the previously obtained equality ($E[\sigma_{MLE}] = \frac{n-1}{n}E[x_1^2] - \frac{n-1}{n}E[x_1 x_2]$), we get:

$$E[\sigma_{MLE}] = \frac{n-1}{n}(\sigma^2 + \mu^2) - \frac{n-1}{n}\mu^2 = \frac{n-1}{n}\sigma^2$$

c. Find an unbiased estimator for σ^2 .

Solution:

It can be immediately proven that $\frac{1}{n-1} \sum_{i=1}^n (x_i - \mu_{MLE})^2$ is an unbiased estimator of σ^2 .

Linear Regression
with only one parameter;
MLE and MAP estimation

CMU, 2012 fall, Tom Mitchell, Ziv Bar-Joseph, midterm, pr. 3

Consider real-valued variables X and Y . The Y variable is generated, conditional on X , from the following process:

$$\begin{aligned}\varepsilon &\sim N(0, \sigma^2) \\ Y &= aX + \varepsilon,\end{aligned}$$

where every ε is an independent variable, called a *noise* term, which is drawn from a Gaussian distribution with mean 0, and standard deviation σ .

This is a one-feature *linear regression* model, where a is the only weight parameter.

The conditional probability of Y has the distribution $p(Y|X, a) \sim N(aX, \sigma^2)$, so it can be written as

$$p(Y|X, a) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{1}{2\sigma^2}(Y - aX)^2\right)$$

MLE estimation

a. Assume we have a training dataset of n pairs (X_i, Y_i) for $i = 1, \dots, n$, and σ is known. Which ones of the following equations correctly represent the maximum likelihood problem for estimating a ? Say *yes* or *no* to each one. More than one of them should have the answer *yes*.

i. $\arg \max_a \sum_i \frac{1}{\sqrt{2\pi}\sigma} \exp \left(-\frac{1}{2\sigma^2} (Y_i - aX_i)^2 \right)$

ii. $\arg \max_a \prod_i \frac{1}{\sqrt{2\pi}\sigma} \exp \left(-\frac{1}{2\sigma^2} (Y_i - aX_i)^2 \right)$

iii. $\arg \max_a \sum_i \exp \left(-\frac{1}{2\sigma^2} (Y_i - aX_i)^2 \right)$

iv. $\arg \max_a \prod_i \exp \left(-\frac{1}{2\sigma^2} (Y_i - aX_i)^2 \right)$

v. $\arg \max_a \sum_i (Y_i - aX_i)^2$

vi. $\operatorname{argmin}_a \sum_i (Y_i - aX_i)^2$

Answer:

$$L_D(a) \stackrel{\text{def.}}{=} p(Y_1, \dots, Y_n | a) = p(Y_1, \dots, Y_n | X_1, \dots, X_n, a)$$

$$\stackrel{i.i.d.}{=} \prod_{i=1}^n p(Y_i | X_i, a) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{1}{2\sigma^2}(Y_i - aX_i)^2\right)$$

Therefore

$$a_{MLE} \stackrel{\text{def.}}{=} \arg \max_a L_D(a) = \arg \max_a \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{1}{2\sigma^2}(Y_i - aX_i)^2\right) \quad (ii.)$$

$$= \arg \max_a \left(\frac{1}{\sqrt{2\pi}\sigma}\right)^n \prod_{i=1}^n \exp\left(-\frac{1}{2\sigma^2}(Y_i - aX_i)^2\right) = \arg \max_a \frac{1}{(\sqrt{2\pi}\sigma)^n} \exp\left(-\sum_{i=1}^n \frac{1}{2\sigma^2}(Y_i - aX_i)^2\right)$$

$$= \arg \max_a \prod_{i=1}^n \exp\left(-\frac{1}{2\sigma^2}(Y_i - aX_i)^2\right) \quad (iv.)$$

$$= \arg \max_a \ln \prod_{i=1}^n \exp\left(-\frac{1}{2\sigma^2}(Y_i - aX_i)^2\right) = \arg \max_a \sum_{i=1}^n -\frac{1}{2\sigma^2}(Y_i - aX_i)^2$$

$$= \arg \max_a -\frac{1}{2\sigma^2} \sum_{i=1}^n (Y_i - aX_i)^2 = \arg \min_a \sum_{i=1}^n (Y_i - aX_i)^2 \quad (vi.)$$

b. Derive the maximum likelihood estimate of the parameter a in terms of the training example X_i 's and Y_i 's. We recommend you start with the simplest form of the problem you found above.

Answer:

$$\begin{aligned} a_{MLE} &= \arg \min_a \sum_{i=1}^n (Y_i - aX_i)^2 = \arg \min_a \left(a^2 \sum_{i=1}^n X_i^2 - 2a \sum_{i=1}^n X_i Y_i + \sum_{i=1}^n Y_i^2 \right) \\ &= -\frac{-2 \sum_{i=1}^n X_i Y_i}{2 \sum_{i=1}^n X_i^2} = \frac{\sum_{i=1}^n X_i Y_i}{\sum_{i=1}^n X_i^2} \end{aligned}$$

MAP estimation

Let's put a prior on a . Assume $a \sim N(0, \lambda^2)$, so

$$p(a|\lambda) = \frac{1}{\sqrt{2\pi}\lambda} \exp\left(-\frac{1}{2\lambda^2}a^2\right)$$

The posterior probability of a is

$$p(a|Y_1, \dots, Y_n, X_1, \dots, X_n, \lambda) = \frac{p(Y_1, \dots, Y_n|X_1, \dots, X_n, a) p(a|\lambda)}{\int_{a'} p(Y_1, \dots, Y_n|X_1, \dots, X_n, a') p(a'|\lambda) da'}$$

We can ignore the denominator when doing MAP estimation.

c. Assume $\sigma = 1$, and a fixed prior parameter λ . Solve for the MAP estimate of a ,

$$\operatorname{argmax}_a [\ln p(Y_1, \dots, Y_n|X_1, \dots, X_n, a) + \ln p(a|\lambda)]$$

Your solution should be in terms of X_i 's, Y_i 's, and λ .

Answer:

$$\begin{aligned}
 & p(Y_1, \dots, Y_n | X_1, \dots, X_n, a) \cdot p(a | \lambda) \\
 &= \left(\prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} \exp \left(-\frac{1}{2\sigma^2} (Y_i - aX_i)^2 \right) \right) \cdot \frac{1}{\sqrt{2\pi}\lambda} \exp \left(-\frac{a^2}{2\lambda^2} \right) \\
 &\stackrel{\sigma=1}{=} \left(\prod_{i=1}^n \frac{1}{\sqrt{2\pi}} \exp \left(-\frac{1}{2} (Y_i - aX_i)^2 \right) \right) \cdot \frac{1}{\sqrt{2\pi}\lambda} \exp \left(-\frac{a^2}{2\lambda^2} \right)
 \end{aligned}$$

Therefore the MAP optimization problem is

$$\begin{aligned}
 & \arg \max_a \left(n \ln \frac{1}{\sqrt{2\pi}} - \frac{1}{2} \sum_{i=1}^n (Y_i - aX_i)^2 + \ln \frac{1}{\sqrt{2\pi}\lambda} - \frac{1}{2\lambda^2} a^2 \right) \\
 &= \arg \max_a \left(-\frac{1}{2} \sum_{i=1}^n (Y_i - aX_i)^2 - \frac{1}{2\lambda^2} a^2 \right) \\
 &= \arg \min_a \left(\sum_{i=1}^n (Y_i - aX_i)^2 + \frac{a^2}{\lambda^2} \right) = \arg \min_a \left(a^2 \left(\sum_{i=1}^n X_i^2 + \frac{1}{\lambda^2} \right) - 2a \sum_{i=1}^n X_i Y_i + \sum_{i=1}^n Y_i^2 \right) \\
 &\Rightarrow a_{MAP} = \frac{\sum_{i=1}^n X_i Y_i}{\sum_{i=1}^n X_i^2 + \frac{1}{\lambda^2}}
 \end{aligned}$$

d. Under the following conditions, how do the prior and conditional likelihood curves change? Do a^{MLE} and a^{MAP} become closer together, or further apart?

	$p(a \lambda)$ prior probability: wider, narrower, or same?	$p(Y_1, \dots, Y_n X_1, \dots, X_n, a)$ conditional likelihood: wider, narrower, or same?	$ a^{MLE} - a^{MAP} $ increase or decrease?
As $\lambda \rightarrow \infty$			
As $\lambda \rightarrow 0$			
More data: as $n \rightarrow \infty$ (fixed λ)			

Answer:

	$p(a \lambda)$ prior probability: wider, narrower, or same?	$p(Y_1, \dots, Y_n X_1, \dots, X_n, a)$ conditional likelihood: wider, narrower, or same?	$ a^{MLE} - a^{MAP} $ increase or decrease?
As $\lambda \rightarrow \infty$	wider	same	decrease
As $\lambda \rightarrow 0$	narrower	same	increase
More data: as $n \rightarrow \infty$ (fixed λ)	same	narrower	decrease


Linear Regression in \mathbb{R}^2

[without “intercept” term]

with either Gaussian or Laplace noise

CMU, 2009 fall, Carlos Guestrin, HW3, pr. 1.5.2

CMU, 2012 fall, Eric Xing, Aarti Singh, HW1, pr. 2



This figure shows a system S which takes two inputs x_1, x_2 and outputs a linear combination of those two inputs, $c_1 x_1 + c_2 x_2$, where c_1 and c_2 are two unknown real numbers.

The device you use to measure the output of S , i.e., $c_1 x_1 + c_2 x_2$, introduces an additive error ε , which is a random variable following some distribution. Thus, the output y that you observe is given by equation (3):

$$y = c_1 x_1 + c_2 x_2 + \varepsilon \quad (3)$$

Assume that you have $n > 2$ instances $\langle x_{j1}, x_{j2}, y_j \rangle_{j=1, \dots, n}$ or equivalently $\langle x_j, y_j \rangle_{j=1, \dots, n}$, where $x_j \stackrel{not.}{=} [x_{j1}, x_{j2}]$. In other words, having n measurements in your hands is equivalent to having n equations of the following form: $y_j = c_1 x_{j1} + c_2 x_{j2} + \varepsilon_j$, $j = 1, \dots, n$.

The *goal* is to estimate c_1 and c_2 from those measurements using the maximum likelihood.

a. Assume that the ε_i for $i = 1, \dots, n$ are i.i.d. Gaussian random variables with zero mean and variance σ^2 .

Compute the loglikelihood function and use it to prove that the maximum likelihood estimate $c^* = [c_1^*, c_2^*]$ is the solution of a least squares approximation problem. Find the solution of the least squares problem.

Answer:

$\varepsilon_i = y_i - (c_1 x_{i1} + c_2 x_{i2}) \sim \mathcal{N}(0, \sigma^2)$. Therefore $y_i \sim \mathcal{N}(c_1 x_{i1} + c_2 x_{i2}, \sigma^2)$. Since the noise are i.i.d., the likelihood function is given by

$$L(c_1, c_2) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y_i - c_1 x_{i1} - c_2 x_{i2})^2}{2\sigma^2}\right).$$

Taking the logarithm, we get the loglikelihood function:

$$l(c_1, c_2) = -\frac{n}{2} \log(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (y_i - c_1 x_{i1} - c_2 x_{i2})^2.$$

Let $y \in \mathbb{R}^n$ be the vector containing the measurements, X the $n \times 2$ matrix with $X_{ij} = x_{ij}$ and $c = [c_1, c_2]^\top$, then we are trying to minimize $\|y - Xc\|_2^2$ resulting in a *solution* $c = (X^\top X)^{-1} X^\top y$.

b. Assume that the ε_i for $i = 1, \dots, n$ are independent Gaussian random variables with zero mean and variance $\text{Var}(\varepsilon_i) = \sigma_i^2$.

Compute the loglikelihood function and find $c^* = [c_1^*, c_2^*]$ which maximizes it, i.e., the MLE.

Answer:

$$\varepsilon_i = y_i - (c_1 x_{i1} + c_2 x_{i2}) \sim \mathcal{N}(0, \sigma_i^2).$$

Similar as before,

$$l(c_1, c_2) = -\frac{n}{2} \log(2\pi) - \sum_{i=1}^n \frac{(y_i - c_1 x_{i1} - c_2 x_{i2})^2}{2\sigma_i^2}.$$


Now we are trying to minimize $\|W(y - Xc)\|_2^2$, where W is a diagonal matrix, with $w_{ii} = \frac{1}{\sigma_i}$, resulting the solution $c = (X^\top W^\top W X)^{-1} X^\top W^\top W y$.

c. Assume that ε_i for $i = 1, \dots, n$ has density $f_{\varepsilon_i}(x) = f(x) = \frac{1}{2b} \exp(-\frac{|x|}{b})$. In other words, our noise is i.i.d. following a Laplace distribution with location parameter $\mu = 0$ and scale parameter b . Compute the loglikelihood function under this noise model and explain why this model leads to more robust solutions.

Answer:

$$l(c_1, c_2) = -n \log(2b) - \sum_{i=1}^n \|y - Xc\|_1^2.$$

It is prepared to see higher values of residuals because it has a larger tail [LC: than the Gaussian]. Thus it is more robust to noise and outliers.



**Derivation of entropy definition,
starting from a set of desirable properties**
CMU, 2005 fall, T. Mitchell, A. Moore, HW1, pr. 2.2

Remark:

The definition $H_n(X) = -\sum_i p_i \log p_i$ is not very intuitive.

Theorem:

If $\psi_n(p_1, \dots, p_n)$ satisfies the following axioms


A1. H_n should be continuous in p_i and symmetric in its arguments;

A2. if $p_i = 1/n$ then H_n should be a monotonically increasing function of n ;
(If all events are equally likely, then having more events means being more uncertain.)


A3. if a choice among N events is broken down into successive choices,
then entropy should be the weighted sum of the entropy at each stage;

then $\psi_n(p_1, \dots, p_n) = -K \sum_i p_i \log p_i$ where K is a positive constant.

Example for the axiom A3:



Encoding 1



Encoding 2

$$H\left(\frac{1}{2}, \frac{1}{3}, \frac{1}{6}\right) = \frac{1}{2} \log 2 + \frac{1}{3} \log 3 + \frac{1}{6} \log 6 = \left(\frac{1}{2} + \frac{1}{6}\right) \log 2 + \left(\frac{1}{3} + \frac{1}{6}\right) \log 3 = \frac{2}{3} + \frac{1}{2} \log 3$$

$$H\left(\frac{1}{2}, \frac{1}{2}\right) + \frac{1}{2} H\left(\frac{2}{3}, \frac{1}{3}\right) = 1 + \frac{1}{2} \left(\frac{2}{3} \log \frac{3}{2} + \frac{1}{3} \log 3\right) = 1 + \frac{1}{2} \left(\log 3 - \frac{2}{3}\right) = \frac{2}{3} + \frac{1}{2} \log 3$$

The next 3 slides:

Case 1: $p_i = 1/n$ for $i = 1, \dots, n$; proof steps

a. $A(n) \stackrel{not.}{=} \psi(1/n, 1/n, \dots, 1/n)$ implies

$$A(s^m) = m A(s) \text{ for any } s, m \in \mathbb{N}^*. \quad (1)$$

b. If $s, m \in \mathbb{N}^*$ (fixed), $s \neq 1$, and $t, n \in \mathbb{N}^*$ such that $s^m \leq t^n \leq s^{m+1}$, then

$$\left| \frac{m}{n} - \frac{\log t}{\log s} \right| \leq \frac{1}{n}. \quad (2)$$

c. For $s^m \leq t^n \leq s^{m+1}$ as above, it follows (immediately)

$$\psi_{s^m} \left(\frac{1}{s^m}, \dots, \frac{1}{s^m} \right) \leq \psi_{t^n} \left(\frac{1}{t^n}, \dots, \frac{1}{t^n} \right) \leq \psi_{s^{m+1}} \left(\frac{1}{s^{m+1}}, \dots, \frac{1}{s^{m+1}} \right)$$

i.e. $A(s^m) \leq A(t^n) \leq A(s^{m+1})$

Show that

$$\left| \frac{m}{n} - \frac{A(t)}{A(s)} \right| \leq \frac{1}{n} \text{ for } s \neq 1. \quad (3)$$

d. Combining (2) + (3) gives immediately



$$\left| \frac{A(t)}{A(s)} - \frac{\log t}{\log s} \right| \leq \frac{2}{n} \text{ pentru } s \neq 1 \quad (4)$$

Show that this inequation implies

$$A(t) = K \log t \text{ with } K > 0 \text{ (due to A2)}. \quad (5)$$

Proof

a.



Applying the axion A3 on the right encoding from above gives:

$$\begin{aligned}
 A(s^m) &= A(s) + s \cdot \frac{1}{s} A(s) + s^2 \cdot \frac{1}{s^2} A(s) + \dots + s^{m-1} \cdot \frac{1}{s^{m-1}} A(s) \\
 &= \underbrace{A(s) + A(s) + A(s) + \dots + A(s)}_{m \text{ times}} = mA(s)
 \end{aligned}$$

Proof (cont'd)

b.

$$s^m \leq t^n \leq s^{m+1} \Rightarrow m \log s \leq n \log t \leq (m+1) \log s \Rightarrow$$

$$\frac{m}{n} \leq \frac{\log t}{\log s} \leq \frac{m}{n} + \frac{1}{n} \Rightarrow 0 \leq \frac{\log t}{\log s} - \frac{m}{n} \leq \frac{1}{n} \Rightarrow \left| \frac{\log t}{\log s} - \frac{m}{n} \right| \leq \frac{1}{n}$$

c.

$$A(s^m) \leq A(t^n) \leq A(s^{m+1}) \xrightarrow{1} m A(s) \leq n A(t) \leq (m+1) A(s) \xrightarrow{s \neq 1}$$

$$\frac{m}{n} \leq \frac{A(t)}{A(s)} \leq \frac{m}{n} + \frac{1}{n} \Rightarrow 0 \leq \frac{A(t)}{A(s)} - \frac{m}{n} \leq \frac{1}{n} \Rightarrow \left| \frac{A(t)}{A(s)} - \frac{m}{n} \right| \leq \frac{1}{n}$$

d. Consider again $s^m \leq t^n \leq s^{m+1}$ with s, t fixed. If $m \rightarrow \infty$ then $n \rightarrow \infty$ and from $\left| \frac{A(t)}{A(s)} - \frac{\log t}{\log s} \right| \leq \frac{1}{n}$ it follows that $\left| \frac{A(t)}{A(s)} - \frac{\log t}{\log s} \right| \rightarrow 0$.

Therefore $\left| \frac{A(t)}{A(s)} - \frac{\log t}{\log s} \right| = 0$ and so $\frac{A(t)}{A(s)} = \frac{\log t}{\log s}$.

Finally, $A(t) = \frac{A(s)}{\log s} \log t = K \log t$, where $K = \frac{A(s)}{\log s} > 0$ (if $s \neq 1$).

Case 2: $p_i \in \mathbb{Q}$ for $i = 1, \dots, n$

Let's consider a set of N equiprobable random events, and $\mathcal{P} = (S_1, S_2, \dots, S_k)$ a partition of this set. Let's denote $p_i = |S_i| / N$.


A “natural” two-step encoding (as shown in the nearby figure) leads to $A(N) = \psi_k(p_1, \dots, p_k) + \sum_i p_i A(|S_i|)$, based on the axiom A3.

Finally, using the result $A(t) = K \log t$, gives:

$$K \log N = \psi_k(p_1, \dots, p_k) + K \sum_i p_i \log |S_i|$$

$$\Rightarrow \psi_k(p_1, \dots, p_k) = K \left[\log N - \sum_i p_i \log |S_i| \right]$$

$$= K \left[\log N \sum_i p_i - \sum_i p_i \log |S_i| \right] = -K \sum_i p_i \log \frac{|S_i|}{N} = -K \sum_i p_i \log p_i$$



**Entropie, entropie corelată,
entropie condițională, câștig de informație:
definiții și proprietăți imediate**

CMU, 2005 fall, T. Mitchell, A. Moore, HW1, pr. 2

Definiții

- **Entropia variabilei X :**

$$H(X) \stackrel{\text{def.}}{=} - \sum_i P(X = x_i) \log P(X = x_i) \stackrel{\text{not.}}{=} E_X[-\log P(X)].$$

- **Entropia condițională specifică a variabilei Y în raport cu valoarea x_k a variabilei X :**

$$H(Y | X = x_k) \stackrel{\text{def.}}{=} - \sum_j P(Y = y_j | X = x_k) \log P(Y = y_j | X = x_k) \\ \stackrel{\text{not.}}{=} E_{Y|X=x_k}[-\log P(Y | X = x_k)].$$

- **Entropia condițională medie a variabilei Y în raport cu variabila X :**

$$H(Y | X) \stackrel{\text{def.}}{=} \sum_k P(X = x_k) H(Y | X = x_k) \stackrel{\text{not.}}{=} E_X[H(Y | X)].$$

- **Entropia corelată a variabilelor X și Y :**

$$H(X, Y) \stackrel{\text{def.}}{=} - \sum_i \sum_j P(X = x_i, Y = y_j) \log P(X = x_i, Y = y_j) \\ \stackrel{\text{not.}}{=} E_{X,Y}[-\log P(X, Y)].$$

- **Informația mutuală a variabilelor X și Y , numită de asemenea *câștigul de informație* al variabilei X în raport cu variabila Y (sau invers):**

$$MI(X, Y) \stackrel{\text{not.}}{=} IG(X, Y) \stackrel{\text{def.}}{=} H(X) - H(X | Y) = H(Y) - H(Y | X)$$

(Observație: ultima egalitate de mai sus are loc datorită rezultatului de la punctul c de mai jos.)

a.

$$H(X) \geq 0.$$

$$H(X) = - \sum_i P(X = x_i) \log P(X = x_i) = \sum_i \underbrace{P(X = x_i)}_{\geq 0} \underbrace{\log \frac{1}{P(X = x_i)}}_{\geq 0} \geq 0$$

Mai mult, $H(X) = 0$ dacă și numai dacă variabila X este constantă:

„ \Rightarrow “ Presupunem că $H(X) = 0$, adică $\sum_i P(X = x_i) \log \frac{1}{P(X = x_i)} = 0$. Datorită faptului că fiecare termen din această sumă este mai mare sau egal cu 0, rezultă că $H(X) = 0$ doar dacă pentru $\forall i$, $P(X = x_i) = 0$ sau $\log \frac{1}{P(X = x_i)} = 0$, adică dacă pentru $\forall i$, $P(X = x_i) = 0$ sau $P(X = x_i) = 1$. Cum însă $\sum_i P(X = x_i) = 1$ rezultă că există o singură valoare x_1 pentru X astfel încât $P(X = x_1) = 1$, iar $P(X = x) = 0$ pentru orice $x \neq x_1$. Altfel spus, variabila aleatoare discretă X este constantă.

„ \Leftarrow “ Presupunem că variabila X este constantă, ceea ce înseamnă că X ia o singură valoare x_1 , cu probabilitatea $P(X = x_1) = 1$. Prin urmare, $H(X) = -1 \cdot \log 1 = 0$.

b.

$$H(Y | X) = - \sum_i \sum_j P(X = x_i, Y = y_j) \log P(Y = y_j | X = x_i)$$

$$\begin{aligned}
 H(Y | X) &= \sum_i P(X = x_i) H(Y | X = x_i) \\
 &= \sum_i P(X = x_i) \left[- \sum_j P(Y = y_j | X = x_i) \log P(Y = y_j | X = x_i) \right] \\
 &= - \sum_i \sum_j \underbrace{P(X = x_i) P(Y = y_j | X = x_i)}_{=P(X=x_i, Y=y_j)} \log P(Y = y_j | X = x_i) \\
 &= - \sum_i \sum_j P(X = x_i, Y = y_j) \log P(Y = y_j | X = x_i)
 \end{aligned}$$

c.

$$H(X, Y) = H(X) + H(Y | X) = H(Y) + H(X | Y)$$

$$\begin{aligned}
 H(X, Y) &= - \sum_i \sum_j p(x_i, y_j) \log p(x_i, y_j) \\
 &= - \sum_i \sum_j p(x_i) \cdot p(y_j | x_i) \log[p(x_i) \cdot p(y_j | x_i)] \\
 &= - \sum_i \sum_j p(x_i) \cdot p(y_j | x_i) [\log p(x_i) + \log p(y_j | x_i)] \\
 &= - \sum_i \sum_j p(x_i) \cdot p(y_j | x_i) \log p(x_i) - \sum_i \sum_j p(x_i) \cdot p(y_j | x_i) \log p(y_j | x_i) \\
 &= - \sum_i p(x_i) \log p(x_i) \cdot \underbrace{\sum_j p(y_j | x_i)}_{=1} - \sum_i p(x_i) \sum_j p(y_j | x_i) \log p(y_j | x_i) \\
 &= H(X) + \sum_i p(x_i) H(Y | X = x_i) = H(X) + H(Y | X)
 \end{aligned}$$

Mai general (regula de înlănțuire):

$$H(X_1, \dots, X_n) = H(X_1) + H(X_2 \mid X_1) + \dots + H(X_n \mid X_1, \dots, X_{n-1})$$

$$\begin{aligned} H(X_1, \dots, X_n) &= E \left[\log \frac{1}{p(x_1, \dots, x_n)} \right] \\ &= - E_{p(x_1, \dots, x_n)} [\log p(x_1, \dots, x_n)] \\ &= - E_{p(x_1, \dots, x_n)} [\log p(x_1) + \log p(x_2 \mid x_1) + \dots + \log p(x_n \mid x_1, \dots, x_{n-1})] \\ &= - E_{p(x_1)} [\log p(x_1)] - E_{p(x_1, x_2)} [\log p(x_2 \mid x_1)] - \dots \\ &\quad - E_{p(x_1, \dots, x_n)} [\log p(x_n \mid x_1, \dots, x_{n-1})] \\ &= H(X_1) + H(X_2 \mid X_1) + \dots + H(X_n \mid X_1, \dots, X_{n-1}) \end{aligned}$$

Relative entropy a.k.a. the Kulback-Leibler divergence,
and the [relationship to] information gain;
some basic properties

CMU, 2007 fall, C. Guestrin, HW1, pr. 1.2

[adapted by Liviu Ciortuz]

The *relative entropy* — also known as the *Kullback-Leibler (KL) divergence* — from a distribution p to a distribution q is defined as

$$KL(p||q) \stackrel{def.}{=} - \sum_{x \in X} p(x) \log \frac{q(x)}{p(x)}$$

From an information theory perspective, the KL-divergence specifies the number of additional bits required on average to transmit values of X if the values are distributed with respect to p but we encode them assuming the distribution q .

Notes

1. KL is not a *distance measure*, since it is not symmetric (i.e., in general $KL(p||q) \neq KL(q||p)$).

Another measure, which is defined as $JSD(p||q) = \frac{1}{2}(KL(p||q) + KL(q||p))$, and is called the **Jensen-Shannon divergence** is symmetric.

2. The quantity

$$\begin{aligned} d(X, Y) &\stackrel{def}{=} H(X, Y) - IG(X; Y) = H(X) + H(Y) - 2IG(X; Y) \\ &= H(X | Y) + H(Y | X) \end{aligned}$$

known as **variation of information**, is a distance metric, i.e., it is non-negative, symmetric, implies indiscernability, and satisfies the triangle inequality.

a. Show that $KL(p||q) = 0$ iff $p(x) = q(x)$ for all x .
 (More genrally, the smaller the KL-divergence, the more similar the two distributions.)

Indicație:

Pentru a demonstra punctul acesta puteți folosi **inegalitatea lui Jensen**:

Dacă $\varphi : \mathbb{R} \rightarrow \mathbb{R}$ este o funcție convexă, atunci pentru orice $t \in [0, 1]$ și orice $x_1, x_2 \in \mathbb{R}$ urmează $\varphi(tx_1 + (1 - t)x_2) \leq t\varphi(x_1) + (1 - t)\varphi(x_2)$.

Dacă φ este funcție strict convexă, atunci egalitatea are loc doar dacă $x_1 = x_2$.

Mai general, pentru orice $a_i \geq 0$, $i = 1, \dots, n$ cu $\sum_i a_i \neq 0$ și orice $x_i \in \mathbb{R}$, $i = 1, \dots, n$, avem

$$\varphi\left(\frac{\sum_i a_i x_i}{\sum_j a_j}\right) \leq \frac{\sum_i a_i \varphi(x_i)}{\sum_j a_j}.$$

Dacă φ este strict convexă, atunci egalitatea are loc doar dacă $x_1 = \dots = x_n$.

Evident, rezultate similare pot fi formulate și pentru funcții concave.

Answer

Vom dovedi inegalitatea $KL(p||q) \geq 0$ folosind inegalitatea lui Jensen, în expresia căreia vom înlocui φ cu funcția convexă $-\log_2$, pe a_i cu $p(x_i)$ și pe x_i cu $\frac{q(x_i)}{p(x_i)}$.

(Pentru conveniență, în cele ce urmează vor renunța la indicele variabilei x .)

Vom avea:

$$\begin{aligned}
 KL(p \parallel q) &\stackrel{\text{def.}}{=} - \sum_x p(x) \log \frac{p(x)}{q(x)} \\
 &\stackrel{\text{Jensen}}{\geq} - \log \left(\sum_x p(x) \frac{q(x)}{p(x)} \right) = - \log \left(\underbrace{\sum_x q(x)}_1 \right) = - \log 1 = 0
 \end{aligned}$$

Vom demonstra acum că $KL(p||q) \geq 0 \Leftrightarrow p = q$.

\Leftarrow

Egalitatea $p(x) = q(x)$ implică $\frac{q(x)}{p(x)} = 1$, deci $\log \frac{q(x)}{p(x)} = 0$ pentru orice x , de unde rezultă imediat $KL(p||q) = 0$.

\Rightarrow

Știm că în inegalitatea lui Jensen are loc egalitatea doar în cazul în care $x_i = x_j$ pentru orice i și j .

În cazul de față, această condiție se traduce prin faptul că raportul $\frac{q(x)}{p(x)}$ este același pentru orice valoare a lui x .

Ținând cont că $\sum_x p(x) = 1$ și $\sum_x p(x) \frac{q(x)}{p(x)} = \sum_x q(x) = 1$, rezultă că $\frac{q(x)}{p(x)} = 1$ sau, altfel spus, $p(x) = q(x)$ pentru orice x , ceea ce înseamnă că distribuțiile p și q sunt identice.

b. We can define the *information gain* as the KL-divergence from the observed joint distribution of X and Y to the product of their observed marginals:

$$IG(X, Y) \stackrel{\text{def.}}{=} KL(p_{X,Y} \parallel (p_X p_Y)) = - \sum_x \sum_y p(x, y) \log \left(\frac{p(x)p(y)}{p(x, y)} \right).$$

Prove that this definition of information gain is equivalent to the one given in problem CMU, 2005 fall, T. Mitchell, A. Moore, HW1, pr. 2. That is, show that $IG(X, Y) = H[X] - H[X|Y] = H[Y] - H[Y|X]$, starting from the definition in terms of KL-divergence.

Remark:

It follows that

$$\begin{aligned} IG(X, Y) &= \sum_y p(y) \sum_x p(x | y) \log \frac{p(x | y)}{p(x)} = \sum_y p(y) KL(p_{X|Y} \parallel p_X) \\ &= E_Y[KL(p_{X|Y} \parallel p_X)] \end{aligned}$$

Answer

By making use of the multiplication rule, namely $p(x, y) = p(x | y)p(y)$, we will have:

$$\begin{aligned}
 & KL(p_{XY} || (p_X p_Y)) \\
 & \stackrel{\text{def. } KL}{=} - \sum_x \sum_y p(x, y) \log \left(\frac{p(x)p(y)}{p(x, y)} \right) \\
 & = - \sum_x \sum_y p(x, y) \log \left(\frac{p(x)p(y)}{p(x | y)p(y)} \right) = - \sum_x \sum_y p(x, y) [\log p(x) - \log p(x | y)] \\
 & = - \sum_x \sum_y p(x, y) \log p(x) - \left(- \sum_x \sum_y p(x, y) \log p(x | y) \right) \\
 & = - \sum_x \log p(x) \underbrace{\sum_y p(x, y)}_{=p(x)} - H[X | Y] \\
 & = H[X] - H[X | Y] = IG(X, Y)
 \end{aligned}$$

c.

Show that $IG(X, Y) \geq 0$ for any discrete random variables X and Y .

Moreover, $IG(X, Y) = 0$ iff X and Y are independent.

Answer:

The two statements are immediate consequences of parts a. and b. already proven.

Remark

95.

Putem demonstra inegalitatea $IG(X, Y) \geq 0$ și în manieră directă, folosind rezultatul de la punctul b. și aplicând inegalitatea lui Jensen în forma generalizată, cu următoarele „amendamente“:

- în locul unui singur indice, se vor considera doi indici (așadar în loc de a_i și x_i vom avea a_{ij} și respectiv x_{ij});
- vom lua $\varphi = -\log_2$ iar $a_{ij} \leftarrow p(x_i, y_j)$ și $x_{ij} \leftarrow \frac{p(x_i)p(y_j)}{p(x_i, y_j)}$;
- în fine, vom ține cont că $\sum_i \sum_j p(x_i, y_j) = 1$.

Prin urmare,

$$\begin{aligned} IG(X, Y) &= \sum_i \sum_j p(x_i, y_j) \log \frac{p(x_i, y_j)}{p(x_i) \cdot p(y_j)} = \sum_i \sum_j p(x_i, y_j) \left[-\log \frac{p(x_i) \cdot p(y_j)}{p(x_i, y_j)} \right] \\ &\geq -\log \left(\sum_i \sum_j p(x_i, y_j) \frac{p(x_i) \cdot p(y_j)}{p(x_i, y_j)} \right) = -\log \left(\sum_i \sum_j p(x_i) \cdot p(y_j) \right) \\ &= -\log \left(\underbrace{\sum_i p(x_i)}_1 \cdot \underbrace{\sum_j p(y_j)}_1 \right) = -\log 1 = 0 \end{aligned}$$

În concluzie, $IG(X, Y) \geq 0$.

Remark (cont'd)

Dacă X și Y sunt variabilele independente,
atunci $p(x_i, y_j) = p(x_i)p(y_j)$ pentru orice i și j .

În consecință, toți logaritmi din partea dreaptă a primei egalități din calculul de mai sus sunt 0 și rezultă $IG(X, Y) = 0$.


Invers, presupunând că $IG(X, Y) = 0$, vom ține cont de faptul că putem exprima câștigul de informație cu ajutorul divergenței KL și vom aplica un raționament similar cu cel de la punctul a .

Rezultă că $\frac{p(x_i)p(y_j)}{p(x_i, y_j)} = 1$ și deci $p(x_i)p(y_j) = p(x_i, y_j)$ pentru orice i și j .

Aceasta echivalează cu a spune că variabilele X și Y sunt independente.

Computing the entropy of the exponential distribution

CMU, 2011 spring, R. Rosenfeld,
HW2, pr. 2.c



Pentru o distribuție de probabilitate continuă P , entropia se definește astfel:

$$H(P) = \int_{-\infty}^{+\infty} P(x) \log_2 \frac{1}{P(x)} dx$$

Calculați entropia *distribuției* continue *exponențiale* de parametru $\lambda > 0$.
Vă reamintim că definiția acestei distribuții este următoarea:

$$P(x) = \begin{cases} \lambda e^{-\lambda x}, & \text{dacă } x \geq 0; \\ 0, & \text{dacă } x < 0. \end{cases}$$

Answer

$$\begin{aligned}
 H(P) &= \int_{-\infty}^0 P(x) \log_2 \frac{1}{P(x)} dx + \int_0^{\infty} P(x) \log_2 \frac{1}{P(x)} dx \\
 &\stackrel{\text{def. } P}{=} \underbrace{\int_{-\infty}^0 0 \log_2 0 dx}_0 + \int_0^{\infty} \lambda e^{-\lambda x} \log_2 \frac{1}{\lambda e^{-\lambda x}} dx = \int_0^{\infty} \lambda e^{-\lambda x} \log_2 \frac{1}{\lambda e^{-\lambda x}} dx \\
 \Rightarrow H(P) &= \frac{1}{\ln 2} \int_0^{\infty} \lambda e^{-\lambda x} \ln \frac{1}{\lambda e^{-\lambda x}} dx = \frac{1}{\ln 2} \int_0^{\infty} \lambda e^{-\lambda x} \left(\ln \frac{1}{\lambda} + \ln \frac{1}{e^{-\lambda x}} \right) dx \\
 &= \frac{1}{\ln 2} \int_0^{\infty} \lambda e^{-\lambda x} (-\ln \lambda + \ln e^{\lambda x}) dx \\
 &= \frac{1}{\ln 2} \int_0^{\infty} \lambda e^{-\lambda x} (-\ln \lambda + \lambda x) dx \\
 &= \frac{1}{\ln 2} \int_0^{\infty} \lambda e^{-\lambda x} (-\ln \lambda) dx + \frac{1}{\ln 2} \int_0^{\infty} \lambda e^{-\lambda x} \lambda x dx \\
 &= \frac{-\ln \lambda}{\ln 2} \int_0^{\infty} \lambda e^{-\lambda x} dx + \frac{\lambda}{\ln 2} \int_0^{\infty} \lambda e^{-\lambda x} x dx \\
 &= \frac{\ln \lambda}{\ln 2} \int_0^{\infty} (e^{-\lambda x})' dx - \frac{\lambda}{\ln 2} \int_0^{\infty} (e^{-\lambda x})' x dx
 \end{aligned}$$

Prima integrală se rezolvă foarte ușor:

$$\int_0^{\infty} (e^{-\lambda x})' dx = e^{-\lambda x} \Big|_0^{\infty} = e^{-\infty} - e^0 = 0 - 1 = -1$$

Pentru a rezolva cea de-a doua integrală se poate folosi *formula de integrare prin părți*:

$$\int_0^{\infty} (e^{-\lambda x})' x dx = e^{-\lambda x} x \Big|_0^{\infty} - \int_0^{\infty} e^{-\lambda x} x' dx = e^{-\lambda x} x \Big|_0^{\infty} - \int_0^{\infty} e^{-\lambda x} dx$$

Integrala definită $e^{-\lambda x} x \Big|_0^{\infty}$ nu se poate calcula direct (din cauza conflictului $0 \cdot \infty$ care se produce atunci când lui x i se atribuie valoarea-limită ∞), ci se calculează folosind *regula lui l'Hôpital*:

$$\lim_{x \rightarrow \infty} x e^{-\lambda x} = \lim_{x \rightarrow \infty} \frac{x}{e^{\lambda x}} = \lim_{x \rightarrow \infty} \frac{x'}{(e^{\lambda x})'} = \lim_{x \rightarrow \infty} \frac{1}{\lambda e^{\lambda x}} = \frac{1}{\lambda} \lim_{x \rightarrow \infty} e^{-\lambda x} = e^{-\infty} = 0,$$

deci

$$e^{-\lambda x} x \Big|_0^{\infty} = 0 - 0 = 0.$$

Integrala $\int_0^\infty e^{-\lambda x} dx$ se calculează ușor:

$$\int_0^\infty e^{-\lambda x} dx = -\frac{1}{\lambda} \int_0^\infty (e^{-\lambda x})' dx = -\frac{1}{\lambda} e^{-\lambda x} \Big|_0^\infty = -\frac{1}{\lambda}(0 - 1) = \frac{1}{\lambda}$$

Prin urmare,

$$\int_0^\infty (e^{-\lambda x})' x dx = 0 - \frac{1}{\lambda} = -\frac{1}{\lambda},$$

ceea ce conduce la rezultatul final:

$$H(P) = \frac{\ln \lambda}{\ln 2}(-1) - \frac{\lambda}{\ln 2} \left(-\frac{1}{\lambda}\right) = -\frac{\ln \lambda}{\ln 2} + \frac{1}{\ln 2} = \frac{1 - \ln \lambda}{\ln 2}.$$

An upper bound for the entropy of a discrete distribution

CMU, 2003 fall, A. Moore, HW1, pr. 1.1

Fie X o variabilă aleatoare discretă care ia n valori și urmează distribuția probabilistă P . Conform definiției, entropia lui X este

$$H(X) = - \sum_{i=1}^n P(X = x_i) \log_2 P(X = x_i).$$

Arătați că $H(X) \leq \log_2 n$.

Sugestie: Puteți folosi inegalitatea $\ln x \leq x - 1$ care are loc pentru orice $x > 0$.

Aşadar,

$$H(X) = \frac{1}{\ln 2} \left(- \sum_{i=1}^n P(X = x_i) \ln P(X = x_i) \right)$$

$$H(X) \leq \log_2 n \Leftrightarrow \frac{1}{\ln 2} \left(- \sum_{i=1}^n P(X = x_i) \ln P(X = x_i) \right) \leq \log_2 n$$

$$\Leftrightarrow - \sum_{i=1}^n P(x_i) \ln P(x_i) \leq \ln n$$

$$\Leftrightarrow \sum_{i=1}^n P(x_i) \ln \frac{1}{P(x_i)} - \underbrace{\left(\sum_{i=1}^n P(x_i) \right)}_1 \ln n \leq 0$$

$$\Leftrightarrow \sum_{i=1}^n P(x_i) \ln \frac{1}{P(x_i)} - \sum_{i=1}^n P(x_i) \ln n \leq 0$$

$$\Leftrightarrow \sum_{i=1}^n P(x_i) \left(\ln \frac{1}{P(x_i)} - \ln n \right) \leq 0$$

$$\Leftrightarrow \sum_{i=1}^n P(x_i) \ln \frac{1}{n P(x_i)} \leq 0$$

Aplicând inegalitatea $\ln x \leq x - 1$ pentru $x = \frac{1}{n P(x_i)}$, vom avea:

$$\sum_{i=1}^n P(x_i) \ln \frac{1}{n P(x_i)} \leq \sum_{i=1}^n P(x_i) \left(\frac{1}{n P(x_i)} - 1 \right) = \sum_{i=1}^n \frac{1}{n} - \underbrace{\sum_{i=1}^n P(x_i)}_1 = 1 - 1 = 0$$

Observație: Această margine superioară chiar este „atinsă“. De exemplu, în cazul în care o variabilă aleatoare discretă X având n valori urmează distribuția uniformă, se poate verifica imediat că $H(X) = \log_2 n$.

Decision Trees

Some exercises

Exemplifying the application of the ID3 algorithm
on a toy *mushrooms* dataset

CMU, 2002(?) spring, Andrew Moore, midterm example questions, pr. 2

You are stranded on a deserted island. Mushrooms of various types grow widely all over the island, but no other food is anywhere to be found. Some of the mushrooms have been determined as poisonous and others as not (determined by your former companions' trial and error). You are the only one remaining on the island. You have the following data to consider:

Example	<i>IsNotHeavy</i>	<i>IsSmelly</i>	<i>IsSpotted</i>	<i>IsSmooth</i>	<i>IsComestible</i>
<i>A</i>	1	0	0	0	1
<i>B</i>	1	0	1	0	1
<i>C</i>	0	1	0	1	1
<i>D</i>	0	0	0	1	0
<i>E</i>	1	1	1	0	0
<i>F</i>	1	0	1	1	0
<i>G</i>	1	0	0	1	0
<i>H</i>	0	1	0	0	0
<i>U</i>	0	1	1	1	?
<i>V</i>	1	1	0	1	?
<i>W</i>	1	1	0	0	?

You know whether or not mushrooms *A* through *H* are poisonous, but you do not know about *U* through *W*.

For the *a–d* questions, consider only mushrooms *A* through *H*.

a. What is the entropy of *IsComestible*?

b. Which attribute should you choose as the root of a decision tree?

Hint: You can figure this out by looking at the data without explicitly computing the information gain of all four attributes.

c. What is the information gain of the attribute you chose in the previous question?

d. Build a ID3 decision tree to classify mushrooms as poisonous or not.





e. Classify mushrooms *U*, *V* and *W* using the decision tree as poisonous or not poisonous.

f. If the mushrooms *A* through *H* that you know are not poisonous suddenly became scarce, should you consider trying *U*, *V* and *W*? Which one(s) and why? Or if none of them, then why not?

a.

$$\begin{aligned}
 H_{Comestibilă} &= H[3+, 5-] \stackrel{def.}{=} -\frac{3}{8} \log_2 \frac{3}{8} - \frac{5}{8} \log_2 \frac{5}{8} = \frac{3}{8} \log_2 \frac{8}{3} + \frac{5}{8} \log_2 \frac{8}{5} = \\
 &= \frac{3}{8} 3 - \frac{3}{8} \log_2 3 + \frac{5}{8} 3 - \frac{5}{8} \log_2 5 = 3 - \frac{3}{8} \log_2 3 - \frac{5}{8} \log_2 5 \approx \\
 &\approx 0.9544
 \end{aligned}$$

b.



c.

$$\begin{aligned}
 H_{0/Netedă} &\stackrel{def.}{=} \frac{4}{8}H[2+, 2-] + \frac{4}{8}H[1+, 3-] = \frac{1}{2} \cdot 1 + \frac{1}{2} \left(\frac{1}{4} \log_2 \frac{4}{1} + \frac{3}{4} \log_2 \frac{4}{3} \right) \\
 &= \frac{1}{2} + \frac{1}{2} \left(\frac{1}{4} \cdot 2 + \frac{3}{4} \cdot 2 - \frac{3}{4} \log_2 3 \right) = \frac{1}{2} + \frac{1}{2} \left(2 - \frac{3}{4} \log_2 3 \right) \\
 &= \frac{1}{2} + 1 - \frac{3}{8} \log_2 3 = \frac{3}{2} - \frac{3}{8} \log_2 3 \approx 0.9056
 \end{aligned}$$

$$\begin{aligned}
 IG_{0/Netedă} &\stackrel{def.}{=} H_{Comestibilă} - H_{0/Netedă} \\
 &= 0.9544 - 0.9056 = 0.0488
 \end{aligned}$$

d.

$$\begin{aligned}
H_{0/U\text{\textit{\textit{șoară}}}} &\stackrel{def.}{=} \frac{3}{8}H[1+, 2-] + \frac{5}{8}H[2+, 3-] \\
&= \frac{3}{8} \left(\frac{1}{3} \log_2 \frac{3}{1} + \frac{2}{3} \log_2 \frac{3}{2} \right) + \frac{5}{8} \left(\frac{2}{5} \log_2 \frac{5}{2} + \frac{3}{5} \log_2 \frac{5}{3} \right) \\
&= \frac{3}{8} \left(\frac{1}{3} \log_2 3 + \frac{2}{3} \log_2 3 - \frac{2}{3} \cdot 1 \right) + \frac{5}{8} \left(\frac{2}{5} \log_2 5 - \frac{2}{5} \cdot 1 + \frac{3}{5} \log_2 5 - \frac{3}{5} \log_2 3 \right) \\
&= \frac{3}{8} \left(\log_2 3 - \frac{2}{3} \right) + \frac{5}{8} \left(\log_2 5 - \frac{3}{5} \log_2 3 - \frac{2}{5} \right) \\
&= \frac{3}{8} \log_2 3 - \frac{2}{8} + \frac{5}{8} \log_2 5 - \frac{3}{8} \log_2 3 - \frac{2}{8} \\
&= \frac{5}{8} \log_2 5 - \frac{4}{8} \approx 0.9512
\end{aligned}$$




$$\Rightarrow IG_{0/U\text{\textit{\textit{șoară}}}} \stackrel{def.}{=} H_{Comestibilă} - H_{0/U\text{\textit{\textit{șoară}}}} = 0.9544 - 0.9512 = 0.0032,$$


$$IG_{0/U\text{\textit{\textit{șoară}}}} = IG_{0/Mirositoare} = IG_{0/ArePete} = 0.0032 < IG_{0/Netedă} = 0.0488$$

Observație importantă

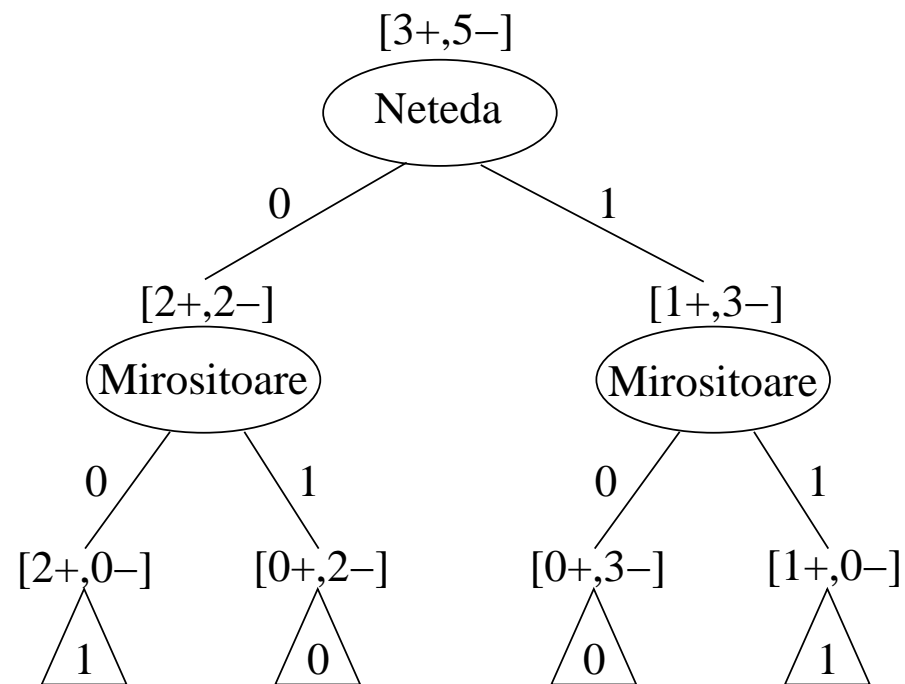
În loc să fi calculat efectiv aceste câștiguri de informație, pentru a determina atributul cel mai „bun“, ar fi fost suficient să comparăm valorile entropiilor condiționale medii $H_{0/Netedă}$ și $H_{0/Ușoară}$:

$$\begin{aligned}
 IG_{0/Netedă} > IG_{0/Ușoară} &\Leftrightarrow H_{0/Netedă} < H_{0/Ușoară} \\
 \Leftrightarrow \frac{3}{2} - \frac{3}{8} \log_2 3 &< \frac{5}{8} \log_2 5 - \frac{1}{2} \\
 \Leftrightarrow 12 - 3 \log_2 3 &< 5 \log_2 5 - 4 \\
 \Leftrightarrow 16 &< 5 \log_2 5 + 3 \log_2 3 \\
 \Leftrightarrow 16 &< 11.6096 + 4.7548 \text{ (adev.)}
 \end{aligned}$$

Nodul 1: Netedă = 0

Nodul 1: Neteďă = 1

Arborele ID3



<i>U</i>	$Netedă = 1, Mirositoare = 1 \Rightarrow Comestibilă = 1$
<i>V</i>	$Netedă = 1, Mirositoare = 1 \Rightarrow Comestibilă = 1$
<i>W</i>	$Netedă = 0, Mirositoare = 1 \Rightarrow Comestibilă = 0$

Exemplifying greedy character of the ID3 algorithm

CMU, 2003 fall, Andrew Moore, midterm, pr. 9.a




Fie attributele binare de intrare A, B, C , atributul de ieșire Y și următoarele exemple de antrenament:

A	B	C	Y
1	1	0	0
1	0	1	1
0	1	1	1
0	0	1	0


a. Determinați arborele de decizie calculat de algoritmul ID3. Este acest arbore de decizie consistent cu datele de antrenament?

Answer

Nodul 0: (rădăcina)



Nodul 1: Avem de clasificat instanțele cu $C = 1$, deci alegerea se face între attributele A și B .





Cele două entropii condiționale medii sunt egale:

$$H_{1/A} = H_{1/B} = \frac{2}{3}H[1+, 1-] + \frac{1}{3}H[1+, 0-]$$

Așadar, putem alege oricare dintre cele două attribute. Pentru fixare, îl alegem pe A .

Nodul 2: La acest nod nu mai avem decât atributul B , deci îl vom pune pe acesta. Arborele complet este reprezentat în slide-ul următor în partea stângă:



Prin construcție, algoritmul ID3 este consistent cu datele de antrenament dacă acestea sunt consistente (i.e., necontradictorii). În cazul nostru, se verifică imediat că datele de antrenament sunt consistente.

b. Există un arbore de decizie de adâncime mai mică (decât cea a arborelui ID3) consistent cu datele de mai sus? Dacă da, ce concept (logic) reprezintă acest arbore?

Se observă că atributul de ieșire Y reprezintă de fapt funcția logică $A \text{ XOR } B$.

Reprezentând această funcție ca arbore de decizie, vom obține arborele din slide-ul precedent, în partea dreaptă.

Acest arbore are cu un nivel mai puțin decât arborele construit cu algoritmul ID3.

Prin urmare, arborele ID3 nu este optim din punctul de vedere al numărului de niveluri.

Exemplifying


The application of the ID3 algorithm on continuous attributes;

Decision surfaces; decision boundaries;


The computation of the CVLOO error

CMU, 2002 fall, Andrew Moore, midterm, pr. 3


- training data:




- discretization / decision thresholds:




- compact representation of the ID3 tree:



- decision “surfaces”:




ID3 tree:




ID3: IG computations


Level 0:



Level 1:




Decision "surfaces":




ID3, CVLOO:
Decision surfaces

CVLOO error: 3/10



DT2




Decision "surfaces":

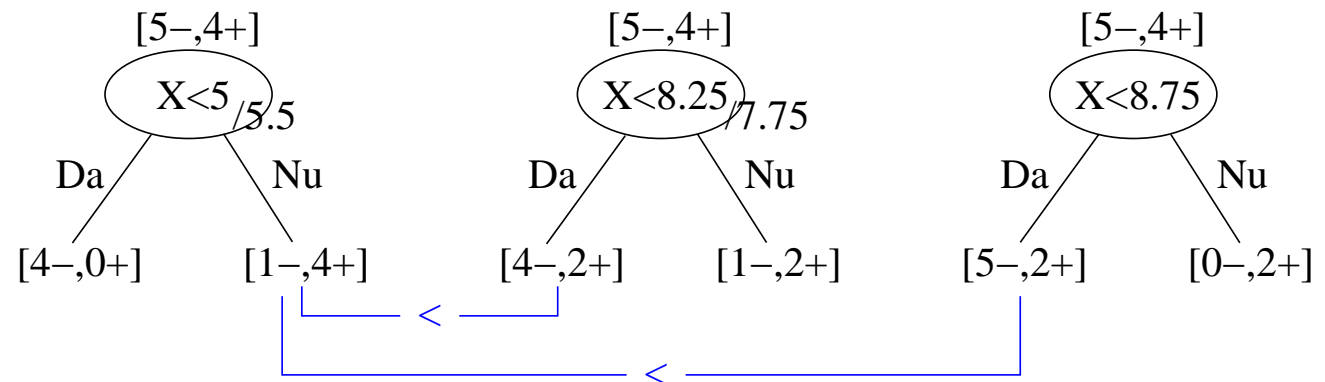


DT2, CVLOO IG computations

Case 1: $X=1, 2, 3, 4$




Case 2: $X=6, 7, 8$

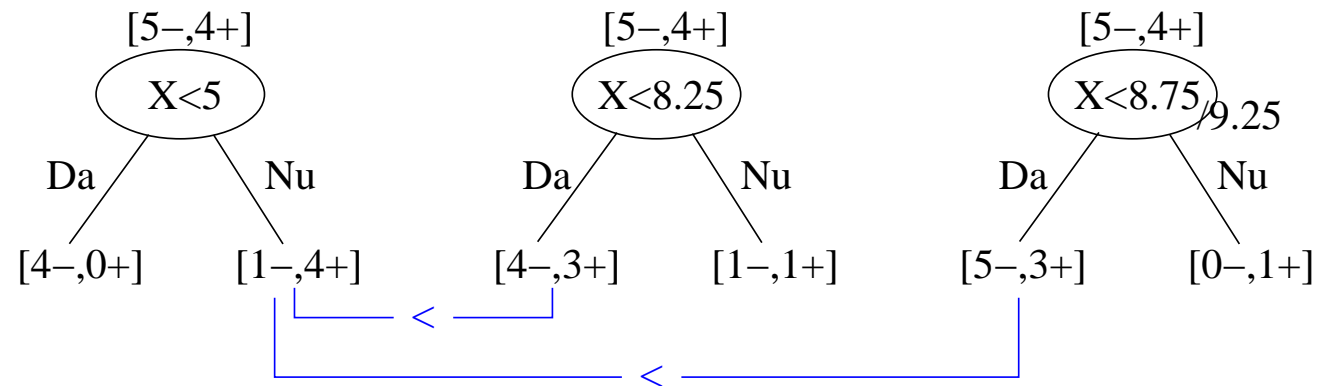


DT2, CVLOO IG computations (cont'd)

Case 3: $X=8.5$



Case 2: $X=9, 10$



CVLOO error: 1/10


Exemplifying

χ^2 -Based Pruning of Decision Trees

CMU, 2010 fall, Ziv Bar-Joseph, HW2, pr. 2.1

Input:

X_1	X_2	X_3	X_4	<i>Class</i>
1	1	0	0	0
1	0	1	0	1
0	1	0	0	0
1	0	1	1	1
0	1	1	1	1
0	0	1	0	0
1	0	0	0	1
0	1	0	1	1
1	0	0	1	1
1	1	0	1	1
1	1	1	1	1
0	0	0	0	0



Idea

While traversing the ID3 tree [usually in bottom-up manner], remove the nodes for which there is not enough (“significant”) **statistical evidence** that there is a **dependence** between the values of the input attribute tested in that node and the values of the output attribute (the labels), supported by the set of instances assigned to that node.

Contingency tables

O_{X_4}	$X_4 = 0$	$X_4 = 1$	
Class = 0	4	0	$N=12 \Rightarrow$
Class = 1	2	6	

$$\left\{ \begin{array}{l} P(\mathbf{Class} = 0) = \frac{4}{12} = \frac{1}{3}, \quad P(\mathbf{Class} = 1) = \frac{2}{3} \\ P(X_4 = 0) = \frac{6}{12} = \frac{1}{2}, \quad P(X_4 = 1) = \frac{1}{2} \end{array} \right.$$

$O_{X_1 X_4=0}$	$X_1 = 0$	$X_1 = 1$	
Class = 0	3	1	$N=6 \Rightarrow$
Class = 1	0	2	

$$\left\{ \begin{array}{l} P(\mathbf{Class} = 0 \mid X_4 = 0) = \frac{4}{6} = \frac{2}{3} \\ P(\mathbf{Class} = 1 \mid X_4 = 0) = \frac{1}{3} \\ P(X_1 = 0 \mid X_4 = 0) = \frac{3}{6} = \frac{1}{2} \\ P(X_1 = 1 \mid X_4 = 0) = \frac{1}{2} \end{array} \right.$$

$O_{X_2 X_4=0, X_1=1}$	$X_2 = 0$	$X_2 = 1$	
Class = 0	0	1	$N=3 \Rightarrow$
Class = 1	2	0	

$$\left\{ \begin{array}{l} P(\mathbf{Class} = 0 \mid X_4 = 0, X_1 = 1) = \frac{1}{3} \\ P(\mathbf{Class} = 1 \mid X_4 = 0, X_1 = 1) = \frac{2}{3} \\ P(X_2 = 0 \mid X_4 = 0, X_1 = 1) = \frac{2}{3} \\ P(X_2 = 1 \mid X_4 = 0, X_1 = 1) = \frac{1}{3} \end{array} \right.$$

The rationale behind the computation of the expected number of observations

$$P(C = i, A = j) = P(C = i) \cdot P(A = j)$$


$$P(C = i) = \frac{\sum_{k=1}^c O_{i,k}}{N} \text{ and } P(A = j) = \frac{\sum_{k=1}^r O_{k,j}}{N}$$

$$P(C = i, A = j) \stackrel{\text{indep.}}{=} \frac{(\sum_{k=1}^c O_{i,k}) (\sum_{k=1}^r O_{k,j})}{N^2}$$

$$E[C = i, A = j] = N \cdot P(C = i, A = j)$$

$$\chi^2 = \sum_{i=1}^r \sum_{j=1}^c \frac{(O_{i,j} - E_{i,j})^2}{E_{i,j}}$$

Chi Squared Pearson test statistics



Expected number of observations

E_{X_4}	$X_4 = 0$	$X_4 = 1$	$E_{X_1 X_4}$	$X_1 = 0$	$X_1 = 1$
<i>Class</i> = 0	2	2	<i>Class</i> = 0	2	2
<i>Class</i> = 1	4	4	<i>Class</i> = 1	1	1

$E_{X_2 X_4, X_1=1}$	$X_2 = 0$	$X_2 = 1$
<i>Class</i> = 0	$\frac{2}{3}$	$\frac{1}{3}$
<i>Class</i> = 1	$\frac{4}{3}$	$\frac{2}{3}$

$E_{X_4}(\text{Class} = 0, X_4 = 0):$

$$N = 12, P(\text{Class} = 0) = \frac{1}{3} \text{ si } P(X_4 = 0) = \frac{1}{2} \Rightarrow$$

$$N \cdot P(\text{Class} = 0, X_4 = 0) = N \cdot P(\text{Class} = 0) \cdot P(X_4 = 0) = 12 \cdot \frac{1}{3} \cdot \frac{1}{2} = 2$$

χ^2 Statistics

$$\chi^2_{X_4} = \frac{(4-2)^2}{2} + \frac{(0-2)^2}{2} + \frac{(2-4)^2}{4} + \frac{(6-4)^2}{4} = 2 + 2 + 1 + 1 = 6$$


$$\chi^2_{X_1|X_4=0} = \frac{(3-2)^2}{2} + \frac{(1-2)^2}{2} + \frac{(0-1)^2}{1} + \frac{(2-1)^2}{1} = 3$$

$$\chi^2_{X_2|X_4=0, X_1=1} = \frac{\left(0 - \frac{2}{3}\right)^2}{\frac{2}{3}} + \frac{\left(1 - \frac{1}{3}\right)^2}{\frac{1}{3}} + \frac{\left(2 - \frac{4}{3}\right)^2}{\frac{4}{3}} + \frac{\left(0 - \frac{2}{3}\right)^2}{\frac{2}{3}} = \frac{4}{9} \cdot \frac{27}{4} = 3$$

p -values: 0.0143, 0.0833, and respectively 0.0833.

The first one of these p -values is smaller than ε , therefore the root node (X_4) cannot be pruned.

Output (pruned tree) for 95% confidence level



Some exercises

Exemplifying

Text classification using the Naive Bayes algorithm

CMU, 2009 spring, Ziv Bar-Joseph, midterm, pr. 2

Training data:

'study'	'free'	'money'	Category	count
1	0	0	Regular	1
0	0	1	Regular	1
1	0	0	Regular	1
1	1	0	Regular	1
0	1	0	Spam	4
0	1	1	Spam	4

Estimating the parameters, by MLE and applying Laplace's rule ("add-one"):

$$P(\text{study}|\text{spam}) = \frac{0+1}{8+2} = \frac{1}{10}$$

$$P(\text{free}|\text{spam}) = \frac{8+1}{8+2} = \frac{9}{10}$$

$$P(\text{money}|\text{spam}) = \frac{4+1}{8+2} = \frac{1}{2}$$

$$P(\text{study}|\text{regular}) = \frac{3+1}{4+2} = \frac{2}{3}$$

$$P(\text{free}|\text{regular}) = \frac{1+1}{4+2} = \frac{1}{3}$$

$$P(\text{money}|\text{regular}) = \frac{1+1}{4+2} = \frac{1}{3}$$

Classification of the message

$s = \text{“money for psychology study”}$,
using the a priori probability $P(\text{spam}) = 0.1$:

$$P(\text{spam} \mid s) = P(\text{spam} \mid \text{study}, \neg \text{free}, \text{money})$$

$$\stackrel{F. Bayes}{=} \frac{P(\text{study}, \neg \text{free}, \text{money} \mid \text{spam}) \cdot P(\text{spam})}{P(\text{study}, \neg \text{free}, \text{money} \mid \text{spam})P(\text{spam}) + P(\text{study}, \neg \text{free}, \text{money} \mid \text{reg})P(\text{reg})}$$

$$P(\text{study}, \neg \text{free}, \text{money} \mid \text{spam}) \stackrel{\text{indep. cdt.}}{=} P(\text{study} \mid \text{spam}) \cdot P(\neg \text{free} \mid \text{spam}) \cdot P(\text{money} \mid \text{spam})$$

$$= \frac{1}{10} \cdot \frac{1}{10} \cdot \frac{1}{2} = \frac{1}{200}$$

$$P(\text{study}, \neg \text{free}, \text{money} \mid \text{reg}) \stackrel{\text{indep. cdt.}}{=} P(\text{study} \mid \text{reg}) \cdot P(\neg \text{free} \mid \text{reg}) \cdot P(\text{money} \mid \text{reg})$$

$$= \frac{2}{3} \cdot \frac{2}{3} \cdot \frac{1}{3} = \frac{4}{27}$$

Therefore,

$$P(\text{spam} \mid s) = \frac{\frac{1}{200} \cdot \frac{1}{10}}{\frac{1}{200} \cdot \frac{1}{10} + \frac{4}{27} \cdot \frac{9}{10}} \approx 0.0037$$

Exemplifying

The computation of the *error rate* for the Naive Bayes algorithm

CMU, 2010 fall, Aarti Singh, HW1, pr. 4.2

Consider a simple learning problem of determining whether Alice and Bob from CA will go to hiking or not $Y : Hike \in \{T, F\}$ given the weather conditions $X_1 : Sunny \in \{T, F\}$ and $X_2 : Windy \in \{T, F\}$ by a Naive Bayes classifier.

Using training data, we estimated the parameters

$$\begin{aligned} P(Hike) &= 0.5 \\ P(Sunny \mid Hike) &= 0.8, \quad P(Sunny \mid \neg Hike) = 0.7 \\ P(Windy \mid Hike) &= 0.4, \quad P(Windy \mid \neg Hike) = 0.5 \end{aligned}$$

Assume that the true distribution of X_1, X_2 , and Y satisfies the Naive Bayes assumption of conditional independence with the above parameters.

a. What is the joint probability that Alice and Bob go to hiking and the weather is sunny and windy, that is $P(Sunny, Windy, Hike)$?

Solution:

$$P(Sunny, Windy, Hike) \stackrel{cdt. \text{ indep.}}{=} P(Sunny|Hike) \cdot P(Windy|Hike) \cdot P(Hike) = 0.8 \cdot 0.4 \cdot 0.5 = 0.16.$$

b. What is the expected error rate of the Naive Bayes classifier?

(Informally, the expected error rate is the probability that an “observation”/instance randomly generated according to the *true* probabilistic distribution of data is incorrectly classified by the Naive Bayes algorithm.)

Solution:

X_1	X_2	Y	$P(X_1, X_2, Y) = P(X_1 Y) \cdot P(X_2 Y) \cdot P(Y)$	$Y_{NB}(X_1, X_2)$	$P_{NB}(Y X_1, X_2)$
F	F	F	$0.3 \cdot 0.5 \cdot 0.5 = 0.075$	F	0.555556
F	F	T	$0.2 \cdot 0.6 \cdot 0.5 = \mathbf{0.060}$	F	0.444444
F	T	F	$0.3 \cdot 0.5 \cdot 0.5 = 0.075$	F	0.652174
F	T	T	$0.2 \cdot 0.4 \cdot 0.5 = \mathbf{0.040}$	F	0.347826
T	F	F	$0.7 \cdot 0.5 \cdot 0.5 = \mathbf{0.175}$	T	0.421686
T	F	T	$0.8 \cdot 0.6 \cdot 0.5 = 0.240$	T	0.578314
T	T	F	$0.7 \cdot 0.5 \cdot 0.5 = 0.175$	F	0.522388
T	T	T	$0.8 \cdot 0.4 \cdot 0.5 = \mathbf{0.160}$	F	0.477612

Note:

Joint probabilities corresponding to incorrect predictions are shown in bold.

$$\begin{aligned}
 \text{error} &\stackrel{\text{def.}}{=} E_P[I_{Y_{NB}(X_1, X_2) \neq Y}] \\
 &= \sum_{X_1, X_2, Y} I[Y_{NB}(X_1, X_2) \neq Y] \cdot P(X_1, X_2, Y) \\
 &= \mathbf{0.060} + \mathbf{0.040} + \mathbf{0.175} + \mathbf{0.160} = \mathbf{0.435}
 \end{aligned}$$

Note:

I is the *indicator* function; its value is 1 whenever the associated condition (in our case, $f_{NB}(X_1, X_2) \neq Y$) is true, and 0 otherwise.

Next, suppose that we gather more information about weather conditions and introduce a new feature denoting whether the weather is $X_3 : \text{Rainy}$ or not. Assume that each day the weather in CA can be either *Rainy* or *Sunny*. That is, it can not be both *Sunny* and *Rainy*. (Similarly, it can not be $\neg \text{Sunny}$ and $\neg \text{Rainy}$).

c. In the above new case, are any of the Naive Bayes assumptions violated? Why (not)? What is the joint probability that Alice and Bob go to hiking and the weather is sunny, windy and not rainy, that is $P(\text{Sunny}, \text{Windy}, \neg \text{Rainy}, \text{Hike})$?

Solution:

The conditional independence of variables given the class label assumption of Naive Bayes is violated. Indeed, knowing if the weather is *Sunny* completely determines whether it is *Rainy* or not. Therefore, *Sunny* and *Rainy* are clearly NOT conditionally independent given *Hike*.

$$\begin{aligned}
 &P(\text{Sunny}, \text{Windy}, \neg \text{Rainy}, \text{Hike}) \\
 &= \underbrace{P(\neg \text{Rainy} | \text{Hike}, \text{Sunny}, \text{Windy})}_1 \cdot P(\text{Sunny}, \text{Windy} | \text{Hike}) \cdot P(\text{Hike}) \\
 &\stackrel{\text{cond. indep.}}{=} P(\text{Sunny} | \text{Hike}) \cdot P(\text{Windy} | \text{Hike}) \cdot P(\text{Hike}) \\
 &= 0.8 \cdot 0.4 \cdot 0.5 = 0.16.
 \end{aligned}$$

d. What is the expected error rate when the Naive Bayes classifier uses all three attributes? Does the performance of Naive Bayes improve by observing the new attribute Rainy? Explain why.

Solution:

X_1	X_2	X_3	Y	$P(X_1, X_2, Y)$	$P_{NB}(X_1, X_2, X_3, Y) = P(X_3) \cdot P(X_1 Y) \cdot P(X_2 Y) \cdot P(Y)$	$Y_{NB}(X_1, X_2, X_3)$	$P_{NB}(Y X_1, X_2, X_3)$
F	F	F	F	0	$0.075 \cdot 0.7 = 0.0525$	F	0.522388
F	F	F	T	0	$0.060 \cdot 0.8 = 0.0480$	F	0.477612
F	F	T	F	0.075	$0.075 \cdot 0.3 = 0.0225$	F	0.652174
F	F	T	T	0.060	$0.060 \cdot 0.2 = 0.0120$	F	0.347826
F	T	F	F	0	$0.075 \cdot 0.7 = 0.0525$	F	0.621302
F	T	F	T	0	$0.040 \cdot 0.8 = 0.0320$	F	0.378698
F	T	T	F	0.075	$0.075 \cdot 0.3 = 0.0225$	F	0.737705
F	T	T	T	0.040	$0.040 \cdot 0.2 = 0.0080$	F	0.262295
T	F	F	F	0.175	$0.175 \cdot 0.7 = 0.0525$	T	0.389507
T	F	F	T	0.240	$0.240 \cdot 0.8 = 0.1920$	T	0.610493
T	F	T	F	0	$0.175 \cdot 0.3 = 0.0525$	F	0.522388
T	F	T	T	0	$0.240 \cdot 0.2 = 0.0480$	F	0.477612
T	T	F	F	0.175	$0.175 \cdot 0.7 = 0.0525$	T	0.489022
T	T	F	T	0.160	$0.160 \cdot 0.8 = 0.1280$	T	0.510978
T	T	T	F	0	$0.175 \cdot 0.3 = 0.0225$	F	0.621302
T	T	T	T	0	$0.060 \cdot 0.2 = 0.0120$	F	0.378698

The new error rate is:

$$0.060 + 0.040 + 0.175 + 0.175 = 0.45 > 0.435 \text{ (see question } b\text{).}$$

The Naive Bayes classifier performance drops because the conditional independence assumptions do not hold for the correlated features.

Computing

The *sample complexity* of the Naive Bayes and Joint Bayes Classifiers

CMU, 2010 spring, Eric Xing, Tom Mitchell, Aarti Singh, HW2, pr. 1.1

A big reason we use the Naive Bayes classifier is that it requires less training data than the Joint Bayes Classifier. This exercise should give you a “feeling” for how great the disparity really is.

Imagine that each *instance* is an independent “*observation*” of the multi-variate random variable $\bar{X} = X_1, \dots, X_d$, where the X_i are i.i.d. and Bernoulli of parameter $p = 0.5$.

To train the Joint Bayes classifier, we need to see every value of \bar{X} “enough” times; training the Naive Bayes classifier only requires seeing both values of X_i “enough” times.

Main Question: How many “observations”/instances are needed until, with probability $1 - \varepsilon$, we have seen every variable we need to see at least once?

Note: To train the classifiers well would require more than this, but for this problem we only require one observation.

Hint: You may want to use the following *inequalities*:

- For any $k \geq 1$, $(1 - 1/k)^k \leq e^{-1}$
- For *any* events E_1, \dots, E_k , $Pr(E_1 \cup \dots \cup E_k) \leq \sum_{i=1}^k Pr(E_i)$.
(This is called the “union bounds” property.)

Consider the Naive Bayes classifier.

- a. Show that if N observations have been made, the probability that a given value of X_i (either 0 or 1) has *not* been seen is $\leq \frac{1}{2^{N-1}}$.
- b. Show that if more than $N_{NB} = 1 + \log_2 \left(\frac{d}{\varepsilon} \right)$ observations have been made, then the probability that *any* X_i has not been observed in both states is $\leq \varepsilon$.

Solution:

$$\text{a. } P(\text{component } X_i \text{ not seen in both states}) = \left(\frac{1}{2}\right)^N + \left(\frac{1}{2}\right)^N = \frac{2}{2^N} = \frac{1}{2^{N-1}}$$

$$\text{b. } P(\text{any component not seen in both states})$$

$$\leq \sum_{i=1}^d P(\text{component } X_i \text{ not seen in both states})$$

$$= \sum_{i=1}^d \frac{1}{2^{N_{NB}-1}} = d \cdot \frac{1}{2^{N_{NB}-1}} = d \cdot \frac{1}{2^{1+\log_2 \frac{d}{\varepsilon}-1}} = d \cdot \frac{1}{2^{\log_2 \frac{d}{\varepsilon}}} = d \cdot \frac{1}{\frac{d}{\varepsilon}} = d \cdot \frac{\varepsilon}{d} = \varepsilon$$

Consider the Joint Bayes classifier.

- c. Let \bar{x} be a particular value of \bar{X} . Show that after N observations, the probability that we have never seen \bar{x} is $\leq e^{-N/2^d}$.
- d. Using the “union bounds” property, show that if more than $N_{JB} = 2^d \ln \left(\frac{2^d}{\varepsilon} \right)$ observations have been made, then the probability that an arbitrarily chosen (but fixed) value of \bar{X} has not been seen is $\leq \varepsilon$.

Solution:

- c. $P(\bar{x} \text{ not seen in } N \text{ observations})$

$$= \left(1 - \frac{1}{2^d}\right)^N = \left[\left(1 - \frac{1}{2^d}\right)^{2^d}\right]^{N/2^d} \leq \left(\frac{1}{e}\right)^{N/2^d} = e^{-N/2^d}$$

- d. $P(\text{any } \bar{x} \text{ not seen in } N_{JB} \text{ observations})$

$$\begin{aligned} &\leq \sum_{\bar{x}} P(\bar{x} \text{ not seen in } N_{JB} \text{ observations}) \\ &= \sum_{\bar{x}} e^{-N_{JB}/2^d} = 2^d \cdot e^{-N_{JB}/2^d} = 2^d \cdot e^{-\ln \frac{2^d}{\varepsilon}} = 2^d \cdot \frac{1}{e^{\ln \frac{2^d}{\varepsilon}}} = \frac{2^d}{\frac{2^d}{\varepsilon}} = \varepsilon \end{aligned}$$

e. Let $d = 2$ and $\varepsilon = 0.1$. What are the values of N_{NB} and N_{JB} ?
 What about $d = 5$?
 And $d = 10$?

Solution:

$$\varepsilon = 0.1, d = 2 \Rightarrow \begin{cases} N_{NB} = 1 + \log_2 \frac{2}{0.1} = 1 + \log_2 20 \approx 5.32 \\ N_{JB} = 2^2 \cdot \ln \frac{2^2}{0.1} = 4 \cdot \ln 40 \approx 14.75 \end{cases}$$


$$\varepsilon = 0.1, d = 5 \Rightarrow \begin{cases} N_{NB} = 1 + \log_2 \frac{5}{0.1} = 1 + \log_2 50 \approx 6.64 \\ N_{JB} = 2^5 \cdot \ln \frac{2^5}{0.1} = 32 \cdot \ln 320 \approx 184.58 \end{cases}$$

$$\varepsilon = 0.1, d = 10 \Rightarrow \begin{cases} N_{NB} = 1 + \log_2 \frac{10}{0.1} = 1 + \log_2 100 \approx 7.64 \\ N_{JB} = 2^{10} \cdot \ln \frac{2^{10}}{0.1} = 1024 \cdot \ln 10240 \approx 9455.67 \end{cases}$$

Exemplifying

ML hypotheses and MAP hypotheses

CMU, 2009 spring, Tom Mitchell, midterm, pr. 2.3-4



Let's consider the 1-dimensional data set shown above, based on the single real-valued attribute X . Notice there are two classes (values of Y), and five data points.

Consider a special type of *decision trees* where leaves have *probabilistic labels*. Each leaf node gives the probability of each possible label, where the probability is the fraction of points at that leaf node with that label.

For *example*, a decision tree learned from the data set above with zero splits would say $P(Y = 1) = 3/5$ and $P(Y = 0) = 2/5$. A decision tree with one split (at $X = 2.5$) would say $P(Y = 1) = 1$ if $X < 2.5$, and $P(Y = 1) = 1/3$ if $X \geq 2.5$.


Solution:

- a. For the above data set, draw a tree that maximizes the *likelihood* of the data.

$T_{ML} = \operatorname{argmax}_T P_T(D)$, where

$$P_T(D) \stackrel{\text{def.}}{=} P(D|T) \stackrel{i.i.d.}{=} \prod_{i=1}^5 P(Y = y_i | X = x_i, T),$$

where y_i is the label/class of the instance x_i ($x_1 = 1.5$, $x_2 = 2$, $x_3 = 3$, $x_4 = 3.5$, $x_5 = 3.75$.)



- b. Consider a prior probability distribution $P(T)$ over trees that penalizes the number of splits in the tree.

$$P(T) \propto \left(\frac{1}{4}\right)^{\text{splits}(T)^2}$$


where T is a tree, $\text{splits}(T)$ is the number of splits in T , and \propto means “is proportional to”.

For the same data set, give the MAP tree when using this prior, $P(T)$, over trees.

Solution:


0 nodes:

$$P(T_0 \mid D) \propto \left(\frac{3}{5}\right)^3 \cdot \left(\frac{2}{5}\right)^2 \cdot \left(\frac{1}{4}\right)^0 = \frac{3^3 \cdot 2^2}{5^5} = \frac{108}{3125} = 0.0336$$



1 node:

$$P(T_1 \mid D) \propto 1^2 \cdot \left(\frac{2}{3}\right)^2 \cdot \frac{1}{3} \cdot \left(\frac{1}{4}\right)^1 = \frac{1}{27} = 0.037$$



2 nodes:

$$P(T_2) \propto \left(\frac{1}{4}\right)^4 \Rightarrow P(T_2 \mid D) \propto 1 \cdot \left(\frac{1}{4}\right)^4 = \frac{1}{256} = 0.0039 \Rightarrow \text{the MAP tree is } T_1.$$

The relationship between [the decision rules of]
Naive Bayes and Logistic Regression

CMU, 2005 fall, Tom Mitchell, HW2, pr. 2

CMU, 2009 fall, Carlos Guestrin, HW1, pr. 4.1.2

CMU, 2009 fall, Geoff Gordon, HW4, pr. 1.2

CMU, 2012 fall, Tom Mitchell, Ziv Bar-Joseph, HW2, pr. 3.a

a. [Equivalence of NB and LR]

In Tom's draft chapter (*Generative and discriminative classifiers: Naive Bayes and logistic regression*) it has been proved that when Y is Boolean and $X = (X_1, \dots, X_n)$ is a vector of continuous variables, then under certain assumptions the Gaussian Naive Bayes classifier implies that $P(Y|X)$ is given by the logistic function with appropriate parameters W . In particular:

$$P(Y = 1|X) = \frac{1}{1 + \exp(w_0 + \sum_{i=1}^n w_i X_i)}$$

and

$$P(Y = 0|X) = \frac{\exp(w_0 + \sum_{i=1}^n w_i X_i)}{1 + \exp(w_0 + \sum_{i=1}^n w_i X_i)}$$

Consider instead the case where Y is Boolean and $X = (X_1, \dots, X_n)$ is a vector of Boolean variables. Prove for this case also that $P(Y|X)$ follows this same form and hence that logistic regression is also the discriminative counterpart to a Naive Bayes generative classifier over Boolean features.

Note:

Discriminative classifiers learn the parameters of $P(Y|X)$ directly, whereas *generative classifiers* instead learn the parameters of $P(X|Y)$ and $P(Y)$.

Hints:

1. Simple notation will help. Since the X_i are Boolean variables, you need only one parameter to define $P(X_i|Y = y_k)$. Define $\theta_{i1} = P(X_i = 1|Y = 1)$, in which case $P(X_i = 0|Y = 1) = 1 - \theta_{i1}$. Similarly, use θ_{i0} to denote $P(X_i = 1|Y = 0)$.

2. Notice with the above notation you can represent $P(X_i|Y = 1)$ as follows:

$$P(X_i = 1|Y = 1) = \theta_{i1}^{X_i} (1 - \theta_{i1})^{(1-X_i)}$$

Note when $X_i = 1$ the second term is equal to 1 because its exponent is zero. Similarly, when $X_i = 0$ the first term is equal to 1 because its exponent is zero.

Solution

b. [Relaxing the conditional independence assumption]

To capture interactions between features, the Logistic Regression model can be supplemented with extra terms. For example, a term can be added to capture a dependency between X_1 and X_2 :

$$P(Y = 1|X) = \frac{1}{1 + \exp(w_0 + w_{1,2}X_1X_2 + \sum_{i=1}^n w_iX_i)}$$

Similarly, the conditional independence assumptions made by Naive Bayes can be relaxed so that X_1 and X_2 are not assumed to be conditionally independent. In this case, we can write:

$$P(Y|X) = \frac{P(Y) P(X_1, X_2|Y) \prod_{i=3}^n P(X_i|Y)}{P(X)}$$

Prove that for this case, that $P(Y|X)$ follows the same form as the logistic regression model supplemented with the extra term that captures the dependency between X_1 and X_2 (and hence that the supplemented Logistic Regression model is the discriminative counterpart to this generative classifier).

Hints:

1. Using simple notation will help here as well. You need more parameters than before to define $P(X_1, X_2, Y)$. Define $\beta_{ijk} = P(X_1 = i, X_2 = j, Y = k)$.
2. The above notation can be used to represent $P(X_1, X_2 | Y = k)$ as follows:


$$P(X_1, X_2, Y = k) = (\beta_{11k})^{X_1 X_2} (\beta_{10k})^{X_1 (1-X_2)} (\beta_{01k})^{(1-X_1) X_2} (\beta_{00k})^{(1-X_1)(1-X_2)}$$

Solution


Exemplifying the Gaussian [Naive] Bayes algorithm
CMU, 2014 fall, W. Cohen, Z. Bar-Joseph, HW2, pr. 5.c

In a two dimensional case, we can visualize how Gaussian Naive Bayes behaves when input features are correlated. A data set is shown in Figure (A), where red points are in Class 0, blue points are in Class 1. The conditional distributions are two-dimensional Gaussians. In (B), (C) and (D), the ellipses represent conditional distributions for each class. The centers of ellipses show the means, and the contours show the boundary of two standard deviations.


- a. Which of them is most likely to be the true conditional distribution?
- b. Which of them is most likely to be estimates by a Gaussian Naive Bayes model?
- c. If we assume the prior probabilities for both classes are equal, which model will achieve a higher accuracy on the training data?




(A) Data



(B)



(C)



(D)

Solution:

- a. (C) is the truth.
- b. (B) corresponds to the Gaussian Naive Bayes estimates. [LC: Here follows the explanation:]
Because the Gaussian Naive Bayes model assume independence of the two features conditioned on the class label, the estimated model should be aligned with the axes. Both (B) and (D) satisfy this, but only in (B) the width and height of the oval, which are proportional to the standard deviation of each axis, matched the data.
- c. (C) gives the lowest training error.

Proving

the relationship between the decision rules for

Gaussian Naive Bayes and the *Logistic Regression* algorithm
when the covariance matrices are diagonal and $\sigma_{i0}^2 = \sigma_{i1}^2$ for $i = 1, \dots, d$

CMU, 2009 spring, Ziv Bar-Joseph, HW2, pr. 2

Assume a two-class ($Y \in \{0, 1\}$) Naive Bayes model over the d -dimensional real-valued input space \mathbb{R}^d , where the input variables $X|Y = 0 \in \mathbb{R}^d$ are distributed as

$$\text{Gaussian}(\mu_0 = \langle \mu_{01}, \dots, \mu_{0d} \rangle, \sigma = \langle \sigma_1, \dots, \sigma_d \rangle)$$

and $X|Y = 1 \in \mathbb{R}^d$ as

$$\text{Gaussian}(\mu_1 = \langle \mu_{11}, \dots, \mu_{1d} \rangle, \sigma = \langle \sigma_1, \dots, \sigma_d \rangle)$$

i.e., the inputs given the class have different means but identical variance for both classes.

Prove that, given the conditions stated above, the conditional probability $P(Y = 1|X = x)$, where $X = (X_1, \dots, X_d)$ and $x = (x_1, \dots, x_d)$ can be written in a similar form to Logistic Regression:

$$\frac{1}{1 + \exp(w_0 + w \cdot x)}$$

with the parameters $w_0 \in \mathbb{R}$ and $w = (w_1, \dots, w_d) \in \mathbb{R}^d$ chosen in a suitable way.

As a consequence, the decision rule for the Gaussian Bayes classifier supported by this model the decision rule has a linear form.

Solution

$$\begin{aligned}
 P(Y = 1|X = x) &\stackrel{B.F.}{=} \frac{P(X = x|Y = 1) P(Y = 1)}{\sum_{y' \in \{0,1\}} P(X = x|Y = y') P(Y = y')} \\
 &= \frac{1}{1 + \frac{P(X = x|Y = 0)P(Y = 0)}{P(X = x|Y = 1)P(Y = 1)}} \\
 &= \frac{1}{1 + \exp \left(\ln \frac{P(X = x|Y = 0)P(Y = 0)}{P(X = x|Y = 1)P(Y = 1)} \right)} \\
 &= \frac{1}{1 + \exp \underbrace{\left(\ln \frac{P(X_1 = x_1, \dots, X_d = x_d|Y = 0)P(Y = 0)}{P(X_1 = x_1, \dots, X_d = x_d|Y = 1)P(Y = 1)} \right)}_{\text{exponent}}}
 \end{aligned}$$

$$\begin{aligned}
\text{exponent} & \stackrel{\text{cond. indep.}}{=} \ln \frac{P(Y=0)}{P(Y=1)} + \sum_{i=1}^d \ln \frac{P(X_i=x_i|Y=0)}{P(X_i=x_i|Y=1)} \\
& = \ln \frac{P(Y=0)}{P(Y=1)} + \sum_{i=1}^d \ln \left(\frac{\frac{1}{\sqrt{2\pi}\sigma_i} \exp\left(-\frac{(x_i-\mu_{i0})^2}{2\sigma_i^2}\right)}{\frac{1}{\sqrt{2\pi}\sigma_i} \exp\left(-\frac{(x_i-\mu_{i1})^2}{2\sigma_i^2}\right)} \right) \\
& = \ln \frac{P(Y=0)}{P(Y=1)} + \sum_{i=1}^d \left(\frac{(x_i-\mu_{i1})^2}{2\sigma_i^2} - \frac{(x_i-\mu_{i0})^2}{2\sigma_i^2} \right) \\
& = \ln \frac{P(Y=0)}{P(Y=1)} + \sum_{i=1}^d \frac{2x_i(\mu_{0i}-\mu_{1i}) + (\mu_{1i}^2 - \mu_{0i}^2)}{2\sigma_i^2} \\
& = \ln \frac{P(Y=0)}{P(Y=1)} + \sum_{i=1}^d \left(\frac{x_i(\mu_{0i}-\mu_{1i})}{\sigma_i^2} + \frac{(\mu_{1i}^2 - \mu_{0i}^2)}{2\sigma_i^2} \right) \\
& = \underbrace{\ln \frac{P(Y=0)}{P(Y=1)} + \sum_{i=1}^d \frac{(\mu_{1i}^2 - \mu_{0i}^2)}{2\sigma_i^2}}_{w_0} + \sum_{i=1}^d \underbrace{\frac{\mu_{0i}-\mu_{1i}}{\sigma_i^2}}_{w_i} x_i
\end{aligned}$$

In conclusion,

$$P(Y = 1|X = x) = \frac{1}{1 + e^{(w \cdot x + w_0)}}$$

with

$$w_0 = \ln \frac{P(Y = 0)}{P(Y = 1)} + \sum_{i=1}^d \frac{(\mu_{1i}^2 - \mu_{0i}^2)}{2\sigma_i^2} \text{ and } w_i = \frac{\mu_{0i} - \mu_{1i}}{\sigma_i^2}, i = 1, \dots, d$$

Note that

$$P(Y = 0|X = x) = \frac{e^{(w \cdot x + w_0)}}{1 + e^{(w \cdot x + w_0)}}$$

and

$$P(Y = 1|X = x) > P(Y = 0|X = x) \Leftrightarrow w \cdot x + w_0 < 0$$

Since the coefficients w_i for $i = 1, \dots, d$ do not depend on x_i , it follows that this *decision rule* of Gaussian Naive Bayes [in the conditions stated in the beginning of this problem] is a linear rule, like in Logistic Regression.

However, this relationship does not mean that there is a one-to-one correspondence between the parameters w_i of Gaussian Naive Bayes (GNB) and the parameters w_i of logistic regression (LR) because LR is discriminative and therefore doesn't model $P(X)$, while GNB does model $P(X)$.

To be more specific, note that the coefficients w_i in the GNB decision rules should be divided by $P(x_1, \dots, x_d)$ in order to correspond to $P(Y = 1|X = x)$, which means that then they will not anymore be independent of x_i , like the LR coefficients.

Estimating the parameters for the Gaussian Naive Bayes algorithm

CMU, 2014 fall, W. Cohen, Z. Bar-Joseph, HW2, pr. 5.ab

Consider a Gaussian Naive Bayes model, where the conditional distribution of each feature is a one-dimensional Gaussian, $X^{(j)}|Y \sim N(\mu_Y^{(j)}, (\sigma_Y^{(j)})^2)$, $j = 1, \dots, d$.

a. Given n independent training data points, $\{(X_1, Y_1), \dots, (X_n, Y_n)\}$, give a maximum-likelihood estimate (MLE) of the conditional distribution of feature $X^{(j)}$, $j = 1, \dots, d$.

Solution:

The likelihood of the samples in Class 0 is

$$\begin{aligned} L(X_{i,0}^{(j)} | \mu_0^{(j)}, (\sigma \mu_0^{(j)})^2) &= \prod_{i=1}^{n_0} \frac{1}{\sqrt{2\pi} \sigma_0^{(j)}} \exp \left(-\frac{(X_{i,0}^{(j)} - \mu_0^{(j)})^2}{2(\sigma_0^{(j)})^2} \right) \\ &= \left(\frac{1}{\sqrt{2\pi} \sigma_0^{(j)}} \right)^{n_0} \exp \left(-\sum_{i=1}^{n_0} \frac{(X_{i,0}^{(j)} - \mu_0^{(j)})^2}{2(\sigma_0^{(j)})^2} \right) \end{aligned}$$

and the log-likelihood is

$$\ln L = -n_0 \ln \sigma_0^{(j)} - \frac{1}{2(\sigma_0^{(j)})^2} \sum_{i=1}^{n_0} (X_{i,0}^{(j)} - \mu_0^{(j)})^2 + \text{constant}$$

Taking the partial derivatives of the log-likelihood, we have

$$\begin{aligned} \frac{\partial \ln L}{\partial \mu_0^{(j)}} = 0 &\Leftrightarrow \sum_{i=1}^{n_0} (X_{i,0}^{(j)} - \mu_0^{(j)}) = 0 \Leftrightarrow \mu_0^{(j)} = \frac{1}{n_0} \sum_{i=1}^{n_0} X_{i,0}^{(j)} \\ \frac{\partial \ln L}{\partial \sigma_0^{(j)}} = 0 &\Leftrightarrow -\frac{n_0}{\sigma_0^{(j)}} + \frac{1}{(\sigma_0^{(j)})^3} \sum_{i=1}^{n_0} (X_{i,0}^{(j)} - \mu_0^{(j)})^2 = 0 \Leftrightarrow (\sigma \mu_0^{(j)})^2 = \frac{1}{n_0} \sum_{i=1}^{n_0} (X_{i,0}^{(j)} - \hat{\mu}_0^{(j)})^2 \end{aligned}$$

Similarly, one can derive the MLE for the parameters in Class 1.

b. Suppose the prior of Y is already given. How many parameters do you need to estimate in Gaussian Naive Bayes model?

Solution:

For each class, there are 2 parameters (the mean and variance) for each feature, therefore there are $2 \cdot 2d = 4d$ parameters for all features in the two classes.

c. In a full/Joint Gaussian Bayes model, we assume that the conditional distribution $\Pr(X|Y)$ is a multidimensional Gaussian, $X|Y \sim \mathcal{N}(\mu_Y, \Sigma_Y)$, where $\mu \in \mathbb{R}^d$ is the mean vector and $\Sigma \in \mathbb{R}^{d \times d}$ is the covariance matrix.

Again, suppose the prior of Y is already given. How many parameters do you need to estimate in a full/Joint Gaussian Bayes model?

Solution:

For each class, there are d parameters for the mean, $d(d+1)/2$ parameters for the covariance matrix, because the covariance matrix is symmetric. Therefore, the number of parameters is $2 \cdot (d + d(d+1)/2) = d(d+3)$ in total for the two classes.

Proving
the relationship between
The full Gaussian Bayes algorithm and *Logistic Regression*
when $\Sigma_0 = \Sigma_1$

CMU, 2009 spring, A. Singh, T. Mitchell, HW2, pr. 2.2

Let's make the following *assumptions*:

1. Y is a boolean variable following a Bernoulli distribution, with parameter $\pi = P(Y = 1)$ and thus $P(Y = 0) = 1 - \pi$.
2. $X = \langle X_1, X_2, \dots, X_n \rangle$ is a vector of random variables *not* conditionally independent given Y , and $P(X|Y = k)$ follows a *multivariate normal distribution* $N(\mu_k, \Sigma)$.

Note that μ_k is the $n \times 1$ mean vector depending on the value of Y , and Σ is the $n \times n$ covariance matrix, which does not depend on Y . We will write/use the density of the multivariate normal distribution in vector/matrix notation.

Is the form of $P(Y|X)$ implied by such this not-so-naive Gaussian Bayes classifier [LC: similar to] the form used by logistic regression?
Derive the form of $P(Y|X)$ to prove your answer.

We start with:

$$\begin{aligned}
 P(Y = 1|X) &= \frac{P(X|Y = 1) P(Y = 1)}{P(X|Y = 1) P(Y = 1) + P(X|Y = 0) P(Y = 0)} \\
 &= \frac{1}{1 + \frac{P(Y = 0) P(X|Y = 0)}{P(Y = 1) P(X|Y = 1)}} = \frac{1}{1 + \exp\left(\ln \frac{P(Y = 0) P(X|Y = 0)}{P(Y = 1) P(X|Y = 1)}\right)} \\
 &= \frac{1}{1 + \exp\left(\ln \frac{P(Y = 0)}{P(Y = 1)} + \ln \frac{P(X|Y = 0)}{P(X|Y = 1)}\right)}
 \end{aligned}$$

Next we will focus on the term $\ln \frac{P(X|Y = 0)}{P(X|Y = 1)}$:

$$\ln \frac{P(X|Y = 0)}{P(X|Y = 1)} = \ln \frac{\frac{1}{(2\pi)^{d/2} |\Sigma|^{1/2}}}{\frac{1}{(2\pi)^{d/2} |\Sigma|^{1/2}}} + \ln \exp[(\star)] = \ln \exp[(\star)] = (\star)$$

where (\star) is the formulation obtained as the difference between the exponential parts of two multivariate Gaussian densities $P(X|Y = 0)$ and $P(X|Y = 1)$.

$$\begin{aligned}
(\star) &= \frac{1}{2}[(X - \mu_1)^\top \Sigma^{-1}(X - \mu_1) - (X - \mu_0)^\top \Sigma^{-1}(X - \mu_0)] \\
&= (\mu_0^\top - \mu_1^\top) \Sigma^{-1} X + \frac{1}{2} \mu_1^\top \Sigma^{-1} \mu_1 - \frac{1}{2} \mu_0^\top \Sigma^{-1} \mu_0
\end{aligned}$$

As a result, we have:

$$\begin{aligned}
P(Y = 1|X) &= \frac{1}{1 + \exp\left(\ln \frac{1 - \pi}{\pi} + \frac{1}{2} \mu_1^\top \Sigma^{-1} \mu_1 - \frac{1}{2} \mu_0^\top \Sigma^{-1} \mu_0 + (\mu_0^\top - \mu_1^\top) \Sigma^{-1} X\right)} \\
&= \frac{1}{1 + \exp(w_0 + w^\top X)}
\end{aligned}$$

where $w_0 = \ln \frac{1 - \pi}{\pi} + \frac{1}{2} \mu_1^\top \Sigma^{-1} \mu_1 - \frac{1}{2} \mu_0^\top \Sigma^{-1} \mu_0$ is a scalar,
and $w = \Sigma^{-1}(\mu_0 - \mu_1)$ is a $d \times 1$ a parameter vector.

Note that $((\mu_0^\top - \mu_1^\top) \Sigma^{-1})^\top = ((\mu_0 - \mu_1)^\top \Sigma^{-1})^\top = (\Sigma^{-1})^\top ((\mu_0 - \mu_1)^\top)^\top = \Sigma^{-1}(\mu_0 - \mu_1)$ **because** Σ^{-1} **is symmetric.**

(Σ is symmetric because it is a covariance matrix, and therefore Σ^{-1} is also symmetric.)

In conclusion, $P(Y|X)$ has the form of the logistic regression (in vector and matrix notation).

Some proofs

k -NN and the Curse of Dimensionality

Proving that the number of examples needed by k -NN
grows exponentially with the number of features

CMU, 2010 fall, Aarti Singh, HW2, pr. 2.2

[Slides originally drawn by Diana Mînzat, MSc student, FII, 2015 spring]

Consider a set of n points x_1, x_2, \dots, x_n independently and uniformly drawn from a p -dimensional zero-centered unit ball

$$B = \{x: \|x\|^2 \leq 1\} \subset \mathbb{R}^p,$$

where $\|x\| = \sqrt{x \cdot x}$ and \cdot is the inner product in \mathbb{R}^p .

In this problem we will study **the size of the 1-nearest neighbourhood of the origin O and how it changes in relation to the dimension p , thereby gain intuition about the downside of k -NN in a high dimension space.**

Formally, this size will be described as the distance from O to its nearest neighbour in the set $\{x_1, \dots, x_n\}$, denoted by d^* :

$$d^* := \min_{1 \leq i \leq n} \|x_i\|,$$

which is a random variable since the sample is random.

a. For $p = 1$, calculate $P(d^* \leq t)$, the *cumulative distribution function (c.d.f.)* of d^* , for $t \in [0, 1]$.

Solution:

In the one-dimensional space ($p = 1$), the unit ball is the interval $[-1, 1]$. The cumulative distribution function will have the following expression:

$$F_{n,1}(t) \stackrel{\text{not.}}{=} P(d^* \leq t) = 1 - P(d^* > t) = 1 - P(|x_i| > t, \text{ for } i = 1, 2, \dots, n)$$

Because the points x_1, \dots, x_n were generated independently, the c.d.f. can also be written as:

$$F_{n,1}(t) = 1 - \prod_{i=1}^n P(|x_i| > t) = 1 - (1 - t)^n$$

b. Find the formula of the *cumulative distribution function* of d^* for the general case, when $p \in \{1, 2, 3, \dots\}$.

Hint: You may find the following fact useful: the volume of a p -dimensional ball with radius r is

$$V_p(r) = \frac{(r\sqrt{\pi})^p}{\Gamma\left(\frac{p}{2} + 1\right)},$$

where Γ is Euler's Gamma function, defined by

$$\Gamma\left(\frac{1}{2}\right) = \sqrt{\pi}, \quad \Gamma(1) = 1, \quad \text{and} \quad \Gamma(x+1) = x\Gamma(x) \quad \text{for any } x > 1.$$

Note: It can be easily shown that $\Gamma(n+1) = n!$ for all $n \in \mathbb{N}^*$, therefore the Gamma function is a generalization of the factorial function.

Solution:

In the general case, i.e. considering a fixed $p \in \mathbb{N}^*$, it is obvious that the cumulative distribution function of d^* will have a similar form to the $p = 1$ case:

$$\begin{aligned} F_{n,p}(t) &\stackrel{\text{not.}}{=} P(d^* \leq t) = 1 - P(d^* > t) = 1 - P(\|x_i\| > t, i = 1, 2, \dots, n) \\ &= 1 - \prod_{i=1}^n P(\|x_i\| > t). \end{aligned}$$

Denoting the volume of the sphere of radius t by $V_p(t)$, and knowing that the points x_1, \dots, x_n follow a uniform distribution, we can rewrite the above formula as follows:

$$F_{n,p}(t) = 1 - \left(\frac{V_p(1) - V_p(t)}{V_p(1)} \right)^n = 1 - \left(1 - \frac{V_p(t)}{V_p(1)} \right)^n.$$

Using the suggested formula for the volume of the sphere, it follows immediately that $F_{n,p} = 1 - (1 - t^p)^n$.

c. What is the *median* of the random variable d^* (i.e., the value of t for which $P(d^* \leq t) = 1/2$) ? The answer should be a *function* of both the sample size n and the dimension p .

Fix $n = 100$ and plot the values of the median function for $p = 1, 2, 3, \dots, 100$ with the median values on the y -axis and the values of p on the x -axis. What do you see?

Solution:


In order to find the median value of the random variable d^* , we will solve the equation $P(d^* \leq t) = 1/2$ of variable t :

$$\begin{aligned} P(d^* \leq t) = \frac{1}{2} &\Leftrightarrow F_{n,p}(t) = \frac{1}{2} \Leftrightarrow 1 - (1 - t^p)^n = \frac{1}{2} \Leftrightarrow (1 - t^p)^n = \frac{1}{2} \\ &\Leftrightarrow 1 - t^p = \frac{1}{2^{1/n}} \Leftrightarrow t^p = 1 - \frac{1}{2^{1/n}} \end{aligned}$$

Therefore,

$$t_{med}(n, p) = \left(1 - \frac{1}{2^{1/n}}\right)^{1/p}.$$

The plot of the function $t_{med}(100, p)$ for $p = 1, 2, \dots, 100$:



Remark:

The minimal sphere containing the nearest neighbour of the origin in the set $\{x_1, x_2, \dots, x_n\}$ grows very fast as the value of p increases.

When p becomes greater than 10, most of the 100 training instances are closer to the surface of the unit ball than to the origin O .

d. Use the c.d.f. derived at point b to determine how large should the sample size n be such that with probability at least 0.9, the distance d^* from O to its nearest neighbor is less than $1/2$, i.e., half way from O to the boundary of the ball.

The answer should be a *function* of p .

Plot this function for $p = 1, 2, \dots, 20$ with the function values on the y -axis and values of p on the x -axis. What do you see?

Hint: You may find useful the Taylor series expansion of $\ln(1 - x)$:

$$\ln(1 - x) = - \sum_{i=1}^{\infty} \frac{x^i}{i} \text{ for } -1 \leq x < 1.$$

Solution:

$$\begin{aligned}
 P(d^* \leq 0.5) \geq 0.9 &\Leftrightarrow F_{n,p}(0.5) \geq \frac{9}{10} \stackrel{b.}{\Leftrightarrow} 1 - \left(1 - \frac{1}{2^p}\right)^n \geq \frac{9}{10} \Leftrightarrow \left(1 - \frac{1}{2^p}\right)^n \leq \frac{1}{10} \\
 &\Leftrightarrow n \cdot \ln \left(1 - \frac{1}{2^p}\right) \leq -\ln 10 \Leftrightarrow n \geq \frac{\ln 10}{-\ln \left(1 - \frac{1}{2^p}\right)}
 \end{aligned}$$

Using the decomposition of $\ln(1 - 1/2^p)$ into a Taylor series (with $x = 1/2^p$), we obtain:

$$\begin{aligned}
 P(d^* \leq 0.5) \geq 0.9 \\
 \Rightarrow n &\geq (\ln 10) 2^p \frac{1}{1 + \frac{1}{2} \cdot \frac{1}{2^p} + \frac{1}{3} \cdot \frac{1}{2^{2p}} + \dots + \frac{1}{n} \frac{1}{2^{(n-1)p}} + \dots} \\
 \Rightarrow n &\geq 2^{p-1} \ln 10.
 \end{aligned}$$

Note:

In order to obtain the last inequality in the above calculations, we considered the following two facts:

- i. $\frac{1}{3 \cdot 2^p} < \frac{1}{4}$ holds for any $p \geq 1$, and
- ii. $\frac{1}{n \cdot 2^{(n-1)p}} \leq \frac{1}{2^n} \Leftrightarrow 2^n \leq n \cdot 2^{(n-1)p}$ holds for any $p \geq 1$ and $n \geq 2$.

(This can be proven by induction on p).

So, we got:


$$1 + \frac{1}{2} \cdot \frac{1}{2^p} + \frac{1}{3} \cdot \frac{1}{2^{2p}} + \dots + \frac{1}{n} \frac{1}{2^{(n-1)p}} + \dots <$$

$$1 + \frac{1}{2} + \frac{1}{4} + \dots + \frac{1}{2^n} + \dots \rightarrow \frac{1}{1 - \frac{1}{2}} = 2.$$

The proven result

$$P(d^* \leq 0.5) \geq 0.9 \Rightarrow n \geq 2^{p-1} \ln 10$$

means that the sample size needed for the probability that $d^* < 0.5$ is large enough (9/10) grows exponentially with p .



e. Having solved the previous problems, what will you say about the downside of k -NN in terms of n and p ?

Solution:

The k -NN classifier works well when a test instance has a “dense” neighborhood in the training data.

However, the analysis here suggests that in order to provide a dense neighborhood, the size of the training sample should be exponential in the dimension p , which is clearly infeasible for a large p .

(Remember that p is the dimension of the space we work in, i.e. the number of features of the training instances.)

**An upper bound for the asymptotic error rate of 1-NN:
twice the error rate of Joint Bayes**

T. Cover and P. Hart (1967)

CMU, 2005 spring, Carlos Guestrin, HW3, pr. 1

Note: we will prove the *Covert & Hart' theorem* in the case of binary classification with real-values inputs.

Let x_1, x_2, \dots be the training examples in some fixed d -dimensional Euclidean space, and y_i be the corresponding binary class labels, $y_i \in \{0, 1\}$.

Let $p_y(x) \stackrel{not.}{=} P(X = x \mid Y = y)$ be the true conditional probability distribution for points in class y . We *assume* continuous and non-zero conditional probabilities: $0 < p_y(x) < 1$ for all x and y .

Let also $\theta \stackrel{not.}{=} P(Y = 1)$ be the probability that a random training example is in class 1. Again, *assume* $0 < \theta < 1$.

a. Calculate $q(x) \stackrel{\text{not.}}{=} p(Y = 1 | X = x)$, the true probability that a data point x belongs to class 1. Express $q(x)$ in terms of $p_0(x)$, $p_1(x)$, and θ .

Solution:

$$\begin{aligned}
 q(x) &\stackrel{\text{F. Bayes}}{=} \frac{P(X = x | Y = 1)P(Y = 1)}{P(X = x)} \\
 &= \frac{P(X = x | Y = 1)P(Y = 1)}{P(X = x | Y = 1)P(Y = 1) + P(X = x | Y = 0)P(Y = 0)} \\
 &= \frac{p_1(x) \theta}{p_1(x) \theta + p_0(x)(1 - \theta)}
 \end{aligned}$$

b. The Joint Bayes classifier (usually called the Bayes Optimal classifier) always assigns a data point x the most probable class: $\operatorname{argmax}_y P(Y = y \mid X = x)$.

Given some test data point x , what is the probability that example x will be misclassified using the Joint Bayes classifier, in terms of $q(x)$?

Solution:

The Joint Bayes classifier fails with probability $P(Y = 0 \mid X = x)$ when $P(Y = 1 \mid X = x) \geq P(Y = 0 \mid X = x)$, and respectively with probability $P(Y = 1 \mid X = x)$ when $P(Y = 0 \mid X = x) \geq P(Y = 1 \mid X = x)$. I.e.,

$$\begin{aligned} \operatorname{Error}_{\text{Bayes}}(x) &= \min\{P(Y = 0 \mid X = x), P(Y = 1 \mid X = x)\} \\ &= \min\{1 - q(x), q(x)\} \\ &= \begin{cases} q(x) & \text{if } q(x) \in [0, 1/2] \\ 1 - q(x) & \text{if } q(x) \in (1/2, 1]. \end{cases} \end{aligned}$$

c. The 1-nearest neighbor classifier assigns a test data point x the label of the closest training point x' .

Given some test data point x and its nearest neighbor x' , what is the *expected error* of the 1-nearest neighbor classifier, i.e., the probability that x will be misclassified, in terms of $q(x)$ and $q(x')$?

Solution:

$$\begin{aligned} \text{Error}_{1\text{-NN}}(x) &= P(Y = 1|X = x)P(Y = 0|X = x') + \\ &\quad P(Y = 0|X = x)P(Y = 1|X = x') \\ &= q(x)(1 - q(x')) + (1 - q(x))q(x'). \end{aligned}$$

d. In the asymptotic case, i.e. when the number of training examples of each class goes to infinity, and the training data fills the space in a dense fashion, the nearest neighbor x' of x has $q(x')$ converging to $q(x)$, i.e. $P(Y = 1|X = x') \rightarrow p(Y = 1|X = x)$.

(This is true due to *i.* the result obtained at the above point a , and *ii.* the assumed continuity of the function $p_y(x) \stackrel{not.}{=} p(X = x|Y = y)$ w.r.t. x .)

By performing this substitution in the expression obtained at point c , give the *asymptotic error* for the 1-nearest neighbor classifier at point x , in terms of $q(x)$.

Solution:

$$\lim_{x' \rightarrow x} Error_{1-NN}(x) = 2q(x)(1 - q(x))$$

e. Show that the asymptotic error obtained at point d is less than twice the Bayes Optimal error obtained at point b and subsequently that this inequality leads to the corresponding relationship between the expected error rates:

$$E[\lim_{n \rightarrow \infty} Error_{1-NN}] \leq 2E[Error_{Bayes}].$$

Solution:

$z(1 - z) \leq z$ for all z , in particular for $z \in [0, 1/2]$, and

$z(1 - z) \leq 1 - z$ for all z , in particular for $z \in [1/2, 1]$.


Therefore, for all x ,

$$q(x)(1 - q(x)) \leq \begin{cases} q(x) & \text{if } q(x) \in [0, 1/2] \\ 1 - q(x) & \text{if } q(x) \in (1/2, 1]. \end{cases}$$

The results obtained at points b and d lead to

$$\lim_{n \rightarrow \infty} Error_{1-NN} = 2q(x)(1 - q(x)) \leq 2Error_{Bayes}(x) \text{ for all } x.$$

By multiplication with $P(x)$ and summing upon all values of x , we get: $E[\lim_{n \rightarrow \infty} Error_{1-NN}] \leq 2E[Error_{Bayes}]$.



Remark

202.

from: *An Elementary Introduction to Statistical Learning Theory*,
S. Kulkarni, G. Harman, 2011, pp. 68-69

An even tighter bound exists for $E[\lim_{n \rightarrow \infty} Error_{1-NN}]$: $2E[Error_{Bayes}](1 - E[Error_{Bayes}])$

Proof:

From $\lim_{x' \rightarrow x} Error_{1-NN}(x) = 2q(x)(1 - q(x))$ (see point d) and $Error_{Bayes}(x) = \min\{1 - q(x), q(x)\}$ (see point b),

it follows that

$$\lim_{x' \rightarrow x} Error_{1-NN}(x) = 2Error_{Bayes}(x)(1 - Error_{Bayes}(x)).$$

By multiplying this last equality with $P(x)$ and summing on all x — in fact, integrating upon x —, we get

$$E[\lim_{x' \rightarrow x} Error_{1-NN}] = 2E[Error_{Bayes}(1 - Error_{Bayes})] = 2E[Error_{Bayes}] - 2E[(Error_{Bayes})^2].$$

Since $E[Z^2] \geq (E[Z])^2$ for any Z ($Var(Z) \stackrel{def.}{=} E[(Z - E[Z])^2] \stackrel{comp.}{=} E[Z^2] - (E[Z])^2 \geq 0$),
it follows that

$$E[\lim_{x' \rightarrow x} Error_{1-NN}] \leq 2E[Error_{Bayes}] - 2(E[Error_{Bayes}])^2 = 2E[Error_{Bayes}](1 - E[Error_{Bayes}]).$$

Remarks

- $E[\lim_{n \rightarrow \infty} Error_{1-NN}] \geq E[Error_{Bayes}]$

Proof:

$$2z - 2z^2 \geq z \quad \forall z \in [0, 1/2] \quad \text{and} \quad 2z - 2z^2 \geq 1 - z \quad \forall z \in [1/2, 1].$$

Therefore,

$$2q(x)(1 - q(x)) \geq Error_{Bayes}(x) \quad \text{for all } x,$$

and

$$\lim_{n \rightarrow \infty} Error_{1-NN}(x) = \lim_{x' \rightarrow x} Error_{1-NN}(x) \geq Error_{Bayes}(x) \quad \text{for all } x.$$

- The Cover & Hart' upper bound for the asymptotic error rate of 1-NN doesn't hold in the non-asymptotic case (where the number of training examples is finite).

Other Results

[from *An Elementary Introduction to Statistical Learning Theory*,
S. Kulkarni, G. Harman, 2011, pp. 69-70]

- When certain restrictions hold,

$$E\left[\lim_{n \rightarrow \infty} \text{Error}_{k\text{-NN}}\right] \leq \left(1 + \frac{1}{k}\right) E[\text{Error}_{\text{Bayes}}].$$

- However, it can be shown that there are some distributions for which 1-NN outperforms k -NN for any fixed $k > 1$.
- If $\frac{k_n}{n} \rightarrow 0$ for $n \rightarrow \infty$ (for instance, $k_n = \sqrt{n}$), then

$$E\left[\lim_{n \rightarrow \infty} \text{Error}_{k_n\text{-NN}}\right] = E[\text{Error}_{\text{Bayes}}].$$

Significance

The last result means that k_n -NN is

- a *universally consistent learner* (because when the amount of training data grows, its performance approaches that of Joint Bayes) and
- *non-parametric* (i.e., the underlying distribution of data can be arbitrary and we need no knowledge of its form).

Some other universally consistent learners exist.

However, the *convergence rate* is critical. For most learning methods, the convergence rate is very slow in high-dimensional spaces (due to “the curse of dimensionality”). It can be shown that *there is no “universal” convergence rate*, i.e. one can always find distributions for which the convergence rate is arbitrarily slow.

There is no one learning method which can universally beat out all other learning methods.

Conclusion

Such results make the ML field continue to be exciting, and makes the design of good learning algorithms and the understanding of their performance an important science and art!

**Exemplifying the application of
hierarchical agglomerative clustering
(single- and complete-linkage)**

CMU, 2012 fall, Tom Mitchell, Ziv Bar-Joseph, HW4, pr. 2.a

The table below is a distance matrix for 6 objects.

	<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>	<i>E</i>	<i>F</i>
<i>A</i>	0					
<i>B</i>	0.12	0				
<i>C</i>	0.51	0.25	0			
<i>D</i>	0.84	0.16	0.14	0		
<i>E</i>	0.28	0.77	0.70	0.45	0	
<i>F</i>	0.34	0.61	0.93	0.20	0.67	0

Show the final result of hierarchical clustering with single-linkage and complete-linkage by drawing the corresponding dendrograms.

Solution:


Single-linkage:

	<i>AB</i>	<i>C</i>	<i>D</i>	<i>E</i>	<i>F</i>
<i>AB</i>	0				
<i>C</i>	0.25	0			
<i>D</i>	0.16	0.14	0		
<i>E</i>	0.28	0.70	0.45	0	
<i>F</i>	0.34	0.93	0.20	0.67	0

	<i>AB</i>	<i>CD</i>	<i>E</i>	<i>F</i>
<i>AB</i>	0			
<i>CD</i>	0.16	0		
<i>E</i>	0.28	0.45	0	
<i>F</i>	0.34	0.20	0.67	0

	<i>ABCD</i>	<i>E</i>	<i>F</i>
<i>ABCD</i>	0		
<i>E</i>	0.28	0	
<i>F</i>	0.20	0.67	0

	<i>ABCD F</i>	<i>E</i>
<i>ABCD F</i>	0	
<i>E</i>	0.28	0




Complete-linkage:

	<i>AB</i>	<i>C</i>	<i>D</i>	<i>E</i>	<i>F</i>
<i>AB</i>	0				
<i>C</i>	0.51	0			
<i>D</i>	0.84	0.14	0		
<i>E</i>	0.77	0.70	0.45	0	
<i>F</i>	0.61	0.93	0.20	0.67	0

	<i>AB</i>	<i>CD</i>	<i>E</i>	<i>F</i>
<i>AB</i>	0			
<i>CD</i>	0.84	0		
<i>E</i>	0.77	0.70	0	
<i>F</i>	0.61	0.93	0.67	0

	<i>ABF</i>	<i>CD</i>	<i>E</i>
<i>ABF</i>	0		
<i>CD</i>	0.93	0	
<i>E</i>	0.77	0.70	0

	<i>ABF</i>	<i>CDE</i>
<i>ABF</i>	0	
<i>CDE</i>	0.93	0



Exemplifying
the application of hierarchical divisive clustering
and the relationship between single-linkage hierarchies and
Minimum Spanning Trees (MSTs)

CMU, 2009 spring, Ziv Bar-Joseph, final exam, pr. 9.3

Hierarchical clustering may be bottom-up or top-down.

In this problem we will see whether a top-down clustering algorithm can be exactly analogous to a bottom-up clustering algorithm.

Consider the following *top-down clustering algorithm*:

1. Calculate the pairwise distance $d(P_i, P_j)$ between every two objects P_i and P_j in the set of objects to be clustered, and build a complete graph on the set of objects with edge weights being the corresponding distances.
2. Generate the Minimum Spanning Tree of the graph, i.e. choose the subset of edges E' with minimum sum of weights such that $G' = (P, E')$ is a single connected tree.
3. Throw out the edge with the heaviest weight to generate two disconnected trees corresponding to top level clusters.
4. Repeat the previous step recursively on the lower level clusters to generate a top-down clustering on the set of n objects.

- a. Apply this algorithm on the dataset given in the nearby table, using the Euclidian distance.
- b. Does this top-down algorithm perform analogously to any bottom-up algorithm that you have encountered in class? Why?

Point	x	y
P_1	1	2
P_2	2	2
P_3	3	6
P_4	6	4
P_5	6	6
P_6	12	12

Solution:


a.

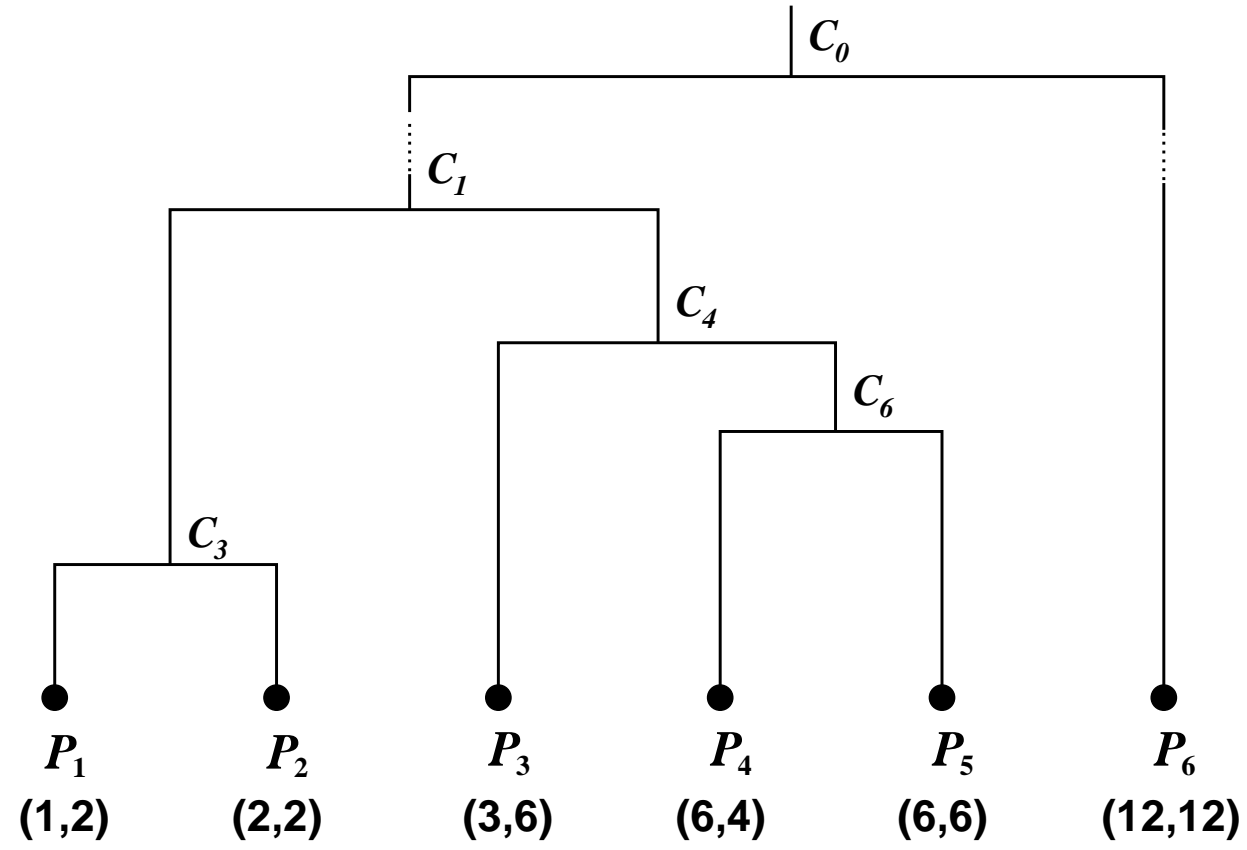
Kruskal algorithm:

1. (P_1, P_2) , cost 1,
2. (P_4, P_5) , cost 2,
3. (P_3, P_5) , cost 3,
4. (P_2, P_3) , cost $\sqrt{17}$
5. (P_5, P_6) , cost $6\sqrt{2}$.

Prim algorithm:

1. (P_1, P_2) , cost 1,
2. (P_2, P_3) , cost $\sqrt{17}$
3. (P_3, P_5) , cost 3,
4. (P_5, P_4) , cost 2,
5. (P_5, P_6) , cost $6\sqrt{2}$.





Note: If there is only one MST for the given dataset, then both Kruskal's and Prim's algorithm will find it. Otherwise, the two algorithms can produce different results.

One can see (both on this dataset and also in general) that Kruskal's algorithm is exactly analogous to the single-linkage bottom-up clustering algorithm.

Therefore, there is indeed a bottom-up equivalent to the top-down clustering algorithm presented in this exercise.

Exemplifying non-hierarchical clustering
using the K -means algorithm

T.U. Dresden, 2006 summer, Steffen Höldobler, Axel Grossmann, HW3


Folosiți algoritmul K -means și distanța euclidiană pentru a grupa următoarele 8 instanțe din \mathbb{R}^2 în 3 cluster:

$$A(2, 10), B(2, 5), C(8, 4), D(5, 8), E(7, 5), F(6, 4), G(1, 2), H(4, 9).$$

Se vor lua drept centroizi inițiali punctele A , D și G .

a. Rulați prima iterație a algoritmului K -means. Pe un grid de valori 10×10 veți marca instanțele date, pozițiile centroizilor la începutul primei iterații și componența fiecărui cluster la finalul acestei iterații. (Trasați mediatoarele segmentelor determinate de centroizi, ca separatori ai clusterelor.)

b. Câte iterații sunt necesare pentru ca algoritmul K -means să convergă? Desenați pe câte un grid rezultatul rulării fiecărei iterații.



Solution:

Iteration 0:





$$\left. \begin{array}{l} \mu_1^0 = (2, 10) \\ \mu_2^0 = (5, 8) \\ \mu_3^0 = (1, 2) \end{array} \right\} \Rightarrow$$

P_i	$d(\mu_1^0, P_i)$	$d(\mu_2^0, P_i)$	$d(\mu_3^0, P_i)$
$A(2, 10)$	0
$B(2, 5)$	5	$3\sqrt{2}$	$\sqrt{10}$
$C(8, 4)$			
$D(5, 8)$...	0	...
$E(7, 5)$			
$F(6, 4)$			
$G(1, 2)$	0
$H(4, 9)$			

$$\Rightarrow \left\{ \begin{array}{l} C_1^0 = \{A\} \\ C_2^0 = \{C, D, E, F, H\} \\ C_3^0 = \{B, G\} \end{array} \right.$$

Iteration 1:

$$\left. \begin{array}{l} \mu_1^1 = \mu_1^0 = (2, 10) \\ \mu_2^1 = \left(\frac{4+5+6+7+8}{5}, \frac{4+4+5+8+9}{5} \right) = (6, 6) \\ \mu_3^1 = \left(\frac{2+1}{2}, \frac{5+2}{2} \right) = (1.5, 3.5) \end{array} \right\} \Rightarrow \dots \Rightarrow \left\{ \begin{array}{l} C_1^0 = \{A, H\} \\ C_2^0 = \{C, D, E, F\} \\ C_3^0 = \{B, G\} \end{array} \right.$$

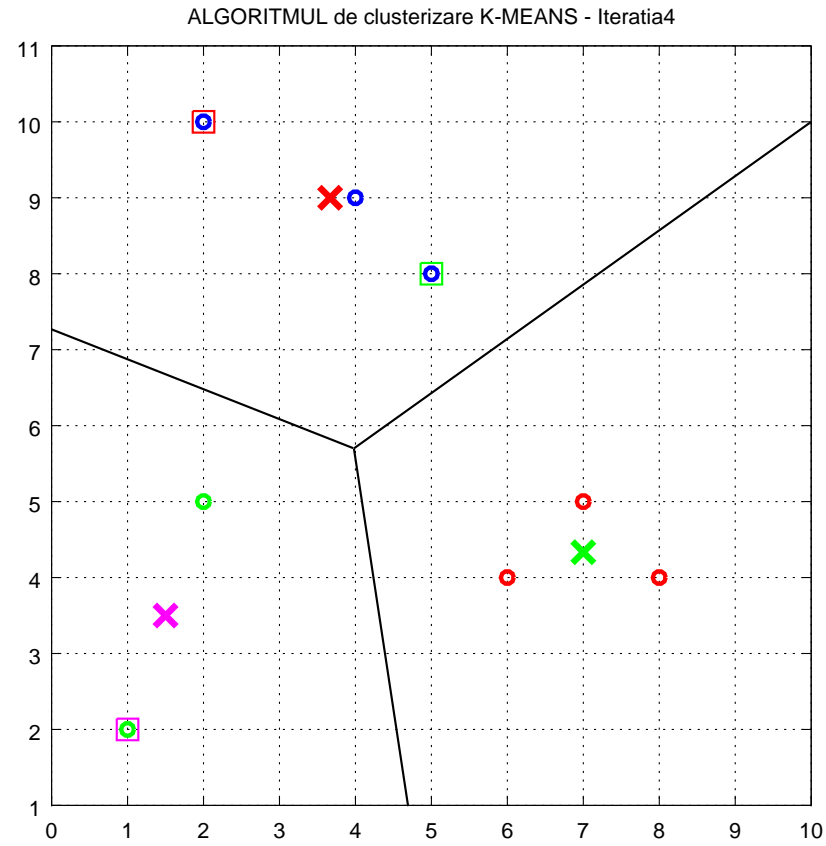


Iteration 2:

$$\left. \begin{array}{l} \mu_1^2 = (3, 9.5) \\ \mu_2^2 = \left(\frac{26}{4}, \frac{21}{4}\right) = (6.5, 5.25) \\ \mu_3^2 = \mu_3^1 = (1.5, 3.5) \end{array} \right\} \Rightarrow \dots \Rightarrow \left\{ \begin{array}{l} C_1^2 = \{A, D, H\} \\ C_2^2 = \{C, E, F\} \\ C_3^2 = \{B, G\} \end{array} \right.$$

Iteration 3:

$$\left. \begin{array}{l} \mu_1^3 = \left(\frac{2+4+5}{3}, \frac{8+9+10}{3}\right) = (11/3, 9) \\ \mu_2^3 = (7, 13/3) \\ \mu_3^3 = \mu_3^2 = (1.5, 3.5) \end{array} \right\} \Rightarrow \dots \Rightarrow \left\{ \begin{array}{l} C_1^3 = \{A, D, H\} = C_1^2 \\ C_2^3 = \{C, E, F\} = C_2^2 \\ C_3^3 = \{B, G\} = C_3^2 \end{array} \right\} \Rightarrow \textit{Stop}$$



Some proofs

K-means as an optimisation algorithm:

The monotonicity of the J_K criterion

[CMU, 2009 spring, Ziv Bar-Joseph, HW5, pr. 2.1]

Algoritmul K -means (S. P. Lloyd, 1957)

Input: $x_1, \dots, x_n \in \mathbb{R}^d$, cu $n \geq K$.

Output: o anumită K -partiție pentru $\{x_1, \dots, x_n\}$.

Procedură:

[*Inițializare/Iterația 0:*] $t \leftarrow 0$;

se fixează în mod arbitrar μ_1^0, \dots, μ_K^0 , centroizii inițiali ai clusterelor, și se asignează fiecare instanță x_i la centroidul cel mai apropiat, formând astfel clusterelor C_1^0, \dots, C_K^0 .

[*Recursivitate:*] Se execută iterația $++ t$:

Pasul 1: se calculează noile poziții ale centroizilor:

$$\mu_j^t = \frac{1}{|C_j^{t-1}|} \sum_{x_i \in C_j^{t-1}} x_i \text{ pentru } j = \overline{1, K};$$

Pasul 2:

se reasignează fiecare x_i la [clusterul cu] centroidul cel mai apropiat, adică se stabilește noua componență a clusterelor la iterația t : C_1^t, \dots, C_K^t ;

[*Terminare:*] până când o anumită condiție este îndeplinită

(de exemplu: până când pozițiile centroizilor — sau: componența clusterelor — nu se mai modifică de la o iterație la alta).

a. Demonstrați că, de la o iterație la alta, algoritmul K -means mărește *coeziunea de ansamblu* a clusterelor. I.e., considerând funcția

$$J(C^t, \mu^t) \stackrel{\text{def.}}{=} \sum_{i=1}^n \|x_i - \mu_{C^t(x_i)}^t\|^2 \stackrel{\text{def.}}{=} \sum_{i=1}^n (x_i - \mu_{C^t(x_i)}^t) \cdot (x_i - \mu_{C^t(x_i)}^t),$$

unde:

$C^t = (C_1^t, C_2^t, \dots, C_K^t)$ este colecția de clusterare (i.e., K -partiția) la momentul t ,

$\mu^t = (\mu_1^t, \mu_2^t, \dots, \mu_K^t)$ este colecția de centroizi ai clusterelor (K -configurația)

la momentul t ,

$C^t(x_i)$ desemnează clusterul la care este asignat elementul x_i la iterația t ,

operatorul \cdot desemnează produsul scalar al vectorilor din \mathbb{R}^d ,

arătați că $J(C^t, \mu^t) \geq J(C^{t+1}, \mu^{t+1})$ pentru orice t .

Ideea demonstrației

Inegalitatea de mai sus rezultă din două inegalități (care corespund pașilor 1 și 2 de la iterația t):

$$J(C^t, \mu^t) \stackrel{(1)}{\geq} J(C^t, \mu^{t+1}) \stackrel{(2)}{\geq} J(C^{t+1}, \mu^{t+1})$$


La prima inegalitate (cea corespunzătoare pasului 1) se poate considera că parametrul C^t este fixat iar μ este variabil, în vreme ce la a doua inegalitate (cea corespunzătoare pasului 2) se consideră μ^t fixat și C variabil.

Prima inegalitate se poate obține însumând o serie de inegalități, și anume câte una pentru fiecare cluster C_j^t . A doua inegalitate se demonstrează imediat.



Ilustrarea acestei idei, pe un exemplu particular:

Vezi următoarele 3 slide-uri



[Edinburgh, 2009 fall, C. Williams, V. Lavrenko, HW4, pr. 3]





init. / iter. 0





iter. 1



iter. 2



iter. 3



Pentru acest exemplu de aplicare a algoritmului K -means, scriem expresiile numerice pentru valoarea criteriului $J_2(C^t, \mu^t)$ pentru fiecare iterație ($t = 0, 1, 2, 3$).

iter.	$J_2(C^t, \mu^t)$
0.	$0 + \{(-9 - (-10))^2 + \dots + (-5 - (-10))^2 + 2[(5 - (-10))^2 + \dots + (9 - (-10))^2]\} \geq$
1.	$(-9 - (-20))^2 + \{(-8 - 7/3))^2 + \dots + (-5 - 7/3))^2 + 2[(5 - 7/3)^2 + \dots + (9 - 7/3)^2]\} \geq$
2.	$(-9 - (-9))^2 + \dots + (-5 - (-9))^2 + 2[(5 - 22/7)^2 + \dots + (9 - 22/7)^2] \geq$
3.	$(-9 - (-7))^2 + \dots + (-5 - (-7))^2 + 2[(5 - 7)^2 + \dots + (9 - 7)^2]$

Observație: La prima vedere, este greu să dovedim aceste inegalități ($J_2(C^{t-1}, \mu^{t-1}) \geq J_2(C^t, \mu^t)$, pentru $t = 1, 2, 3$) ...altfel decât calculând efectiv valoarea expresiilor care se compară. Însă, introducând niște termeni intermediari, inegalitățile acestea se vor demonstra într-un mod foarte elegant...

iter.	$J_2(C^{t-1}, \mu^t)$	$J_2(C^t, \mu^t)$
0.		$0 + \{(-9 - (-10))^2 + \dots + (-5 - (-10))^2 + 2[(5 - (-10))^2 + \dots + (9 - (-10))^2]\} \geq$
1.	$0 + \{(-9 - 7/3)^2 + (-8 - 7/3)^2 + \dots + (-5 - 7/3)^2 + 2[(5 - 7/3)^2 + \dots + (9 - 7/3)^2]\} \geq$	$(-9 - (-20))^2 + \{(-8 - 7/3)^2 + \dots + (-5 - 7/3)^2 + 2[(5 - 7/3)^2 + \dots + (9 - 7/3)^2]\} \geq$
2.	$(-9 - (-9))^2 + \{(-8 - 22/7)^2 + \dots + (-5 - 22/7)^2 + 2[(5 - 22/7)^2 + \dots + (9 - 22/7)^2]\} \geq$	$(-9 - (-9))^2 + (-8 - (-9))^2 + \dots + (-5 - (-9))^2 + 2[(5 - 22/7)^2 + \dots + (9 - 22/7)^2] \geq$
3.	$(-9 - (-7))^2 + \dots + (-5 - (-7))^2 + 2[(5 - 7)^2 + \dots + (9 - 7)^2] =$	$(-9 - (-7))^2 + \dots + (-5 - (-7))^2 + 2[(5 - 7)^2 + \dots + (9 - 7)^2]$

Explicații:

1. Inegalitățile pe orizontală ($J_2(C^{t-1}, \mu^t) \geq J_2(C^t, \mu^t)$, pentru $t = 1, 2, 3$) sunt ușor de demonstrat, pe baza corespondenței termen cu termen. (Ele corespund eventualelor micșorări ale distanțelor atunci când se face reassignarea instanțelor la centroizi.)

2. Restul inegalităților ($J_2(C^t, \mu^t) \geq J_2(C^t, \mu^{t+1})$, pentru $t = 1, 2, 3$) se rezolvă printr-o metodă de optimizare simplă. De exemplu, pentru $t = 1$ este imediat că funcția $(-9 - x)^2 + (-8 - x)^2 + \dots + (-5 - x)^2 + 2[(5 - x)^2 + \dots + (9 - x)^2]$ își atinge minimumul pentru $x = 7/3$, deci $J_2(C^1, \mu^2) \geq J_2(C^1, \mu^1)$.

Demonstrație, pentru cazul general

Observație: Pentru conveniență, ne vom limita la cazul $d = 1$. Extinderea demonstrației la cazul $d > 1$ nu comportă dificultăți.

Demonstrarea inegalității (1): $J(C^t, \mu^t) \geq J(C^t, \mu^{t+1})$

(Vezi pasul 1 al iterației t .)

Fixăm $j \in \{1, \dots, K\}$. Dacă notăm cu $C_j^t = \{x_{i_1}, x_{i_2}, \dots, x_{i_l}\}$, unde $l \stackrel{\text{not.}}{=} |C_j^t|$, atunci

$$J(C_j^t, \mu_j^t) = \sum_{p=1}^l (x_{i_p} - \mu_j^t)^2, \text{ deci } J(C^t, \mu^t) = \sum_{j=1}^K J(C_j^t, \mu_j^t).$$

Dacă se consideră C_j^t fixat, iar μ_j^t variabil, atunci putem minimiza imediat funcția

$$f(\mu) \stackrel{\text{def.}}{=} J(C_j^t, \mu) = l\mu^2 - 2\mu \sum_{p=1}^l x_{i_p} + \sum_{p=1}^l x_{i_p}^2 \Rightarrow \arg \min_{\mu} J(C_j^t, \mu) = \frac{1}{l} \sum_{p=1}^l x_{i_p} \stackrel{\text{def.}}{=} \mu_j^{t+1}.$$

Așadar, $J(C_j^t, \mu) \geq J(C_j^t, \mu_j^{t+1})$, pentru $\forall \mu$. În particular, pentru $\mu = \mu_j^t$ vom avea: $J(C_j^t, \mu_j^t) \geq J(C_j^t, \mu_j^{t+1})$. Inegalitatea aceasta este valabilă pentru toate clusterelor $j = 1, \dots, K$. Dacă sumăm toate aceste inegalități, rezultă: $J(C^t, \mu^t) \geq J(C^t, \mu^{t+1})$.

Demonstrarea inegalității (2): $J(C^t, \mu^{t+1}) \geq J(C^{t+1}, \mu^{t+1})$

(Vezi pasul 2 al iterației t .)

La acest pas, o instanță oarecare x_i , unde $i \in \{1, \dots, n\}$, este reassignată de la clusterul cu centroidul μ_j^{t+1} , la un alt centroid μ_q^{t+1} , dacă

$$\|x_i - \mu_{j'}^{t+1}\|^2 \geq \|x_i - \mu_q^{t+1}\|^2 \Leftrightarrow (x_i - \mu_{j'}^{t+1})^2 \geq (x_i - \mu_q^{t+1})^2, \text{ pentru orice } j' = 1, \dots, K.$$

În contextul iterației t , acest lucru implică

$$\left(x_i - \mu_{C^t(x_i)}^{t+1}\right)^2 \geq \left(x_i - \mu_{C^{t+1}(x_i)}^{t+1}\right)^2.$$

Sumând membru cu membru inegalitățile de acest tip obținute pentru $i = \overline{1, n}$, rezultă: $J(C^t, \mu^{t+1}) \geq J(C^{t+1}, \mu^{t+1})$, ceea ce era de demonstrat.


b. Ce puteți spune despre oprirea algoritmului K -means? (Termină oare acest algoritm într-un număr finit de pași, sau este posibil ca el să reviziteze o K -configurație anterioară $\mu = (\mu_1, \dots, \mu_K)$?)

Răspuns:

Dacă algoritmul revizitează o K -partiție, atunci rezultă că pentru un anumit t avem $J(C^{t-1}, \mu^t) = J(C^t, \mu^{t+1})$. Este posibil ca acest fapt să se întâmple, și anume atunci când:

- există instanțe multiple (i.e., $x_i = x_j$, deși $i \neq j$),
- criteriul de oprire al algoritmului K -means este de forma “până când componența clusterelor nu se mai modifică”,
- se presupune că, în cazul în care o instanță x_i este situată la egală distanță față de doi sau mai mulți centroizi, ea poate fi asignată în mod aleatoriu la oricare dintre ei.

Așa se întâmplă în *exemplul* din figura alăturată dacă se consideră că la o iterație t avem $x_2 = 0 \in C_1^t$ și $x_3 = 0 \in C_2^t$, iar la iterația următoare alegem ca $x_3 = 0 \in C_1^{t+1}$ și $x_2 = 0 \in C_2^{t+1}$ și, din nou, invers la iterația $t + 2$.



Observații

- Dacă se păstrează criteriul dat ca exemplu în enunțul problemei – adică se iterează până când centroizii “staționează” – algoritmul se poate opri fără ca la ultima iterație $J(C, \mu)$ să fi atins minimul posibil. În cazul exemplului de mai sus, vom avea $\frac{1}{4} + 2 \cdot \frac{1}{4} + \frac{1}{4} = 1 > \frac{2}{3}$.
- Dacă nu există instanțe multiple care să fie situate la distanțe egale față de doi sau mai mulți centroizi la o iterație oarecare a algoritmului K -means (precum sunt x_2 și x_3 în *exemplul* de mai sus), sau dacă se impune *restricția* ca în astfel de situații instanțele identice să fie asignate la un singur cluster, este evident că algoritmul K -means se oprește într-un număr finit de pași.

Concluzii

- Algoritmul K -means explorează — pornind de la o anumită inițializare a celor K centroizi —, doar un subset din totalul de K^n K -partiții, asigurându-ne însă că are loc proprietatea $J(C^0, \mu^1) \geq J(C^1, \mu^2) \geq \dots \geq J(C^{t-1}, \mu^t) \geq J(C^t, \mu^{t+1})$, conform punctului a al acestei probleme.
- **Atingerea minimului global** al funcției $J(C, \mu)$ — unde C este o variabilă care parcurge mulțimea tuturor K -partițiilor care se pot forma cu instanțele $\{x_1, \dots, x_n\}$ — **nu este garantată pentru algoritmul K -means**. Valoarea funcției J care se obține la oprirea algoritmului K -means este dependentă de plasarea inițială a centroizilor μ precum și de modul concret în care sunt alcătuite clusterelor în cazul în care o instanță oarecare se află la distanță egală de doi sau mai mulți centroizi, după cum am arătat în exemplul de mai sus.

K-means algorithm:

The “approximate” maximization of the “distance” between clusters

[CMU, 2010 fall, Aarti Singh, HW3, pr. 5.2]

Note: In this problem we will work with a version of the K -means algorithm which is slightly modified w.r.t. the one given in the problem CMU, 2009 spring, Ziv Bar-Joseph, HW5, pr. 2.1, where we have proved the monotonicity of the criterion J .

Let $X := \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\}$ be our sample points, and K denote the number of clusters to use. We represent the cluster assignments of the data points by an indicator matrix $\gamma \in \{0, 1\}^{n \times K}$ such that $\gamma_{ij} = 1$ means \mathbf{x}_i belongs to cluster j . We require that each point belongs to exactly one cluster, so $\sum_{j=1}^K \gamma_{ij} = 1$.

[We already know that] the K -means algorithm “estimates” γ by minimizing the following “cohesion criterion” (or, “measure of distortion”):

$$J(\gamma, \boldsymbol{\mu}_1, \boldsymbol{\mu}_2, \dots, \boldsymbol{\mu}_K) := \sum_{i=1}^n \sum_{j=1}^K \gamma_{ij} \|\mathbf{x}_i - \boldsymbol{\mu}_j\|^2,$$

where $\|\cdot\|$ denotes the vector 2-norm.

K -means alternates between estimating γ and re-computing $\boldsymbol{\mu}_j$'s.

- Initialize $\mu_1, \mu_2, \dots, \mu_K$, and let $C := \{1, \dots, K\}$.
- While the value of J is still decreasing, repeat the following:

1. Determine γ by

$$\gamma_{ij} \leftarrow \begin{cases} 1, & \|\mathbf{x}_i - \mu_j\|^2 \leq \|\mathbf{x}_i - \mu_{j'}\|^2, \forall j' \in C, \\ 0, & \text{otherwise.} \end{cases}$$

Break ties arbitrarily.

2. Recompute μ_j using the updated γ :
 For each $j \in C$, if $\sum_{i=1}^n \gamma_{ij} > 0$ set

$$\mu_j \leftarrow \frac{\sum_{i=1}^n \gamma_{ij} \mathbf{x}_i}{\sum_{i=1}^n \gamma_{ij}}.$$

Otherwise, don't change μ_j .

Let $\bar{\mathbf{x}}$ denote the sample mean.

Consider the following three quantities:

$$\text{Total variation:} \quad V(X) = \frac{\sum_{i=1}^n \|\mathbf{x}_i - \bar{\mathbf{x}}\|^2}{n}.$$

$$\text{Within-cluster variation:} \quad V_j(X) = \frac{\sum_{i=1}^n \gamma_{ij} \|\mathbf{x}_i - \boldsymbol{\mu}_j\|^2}{\sum_{i=1}^n \gamma_{ij}}.$$

$$\text{Between-cluster variation:} \quad \tilde{V}(X) = \sum_{j=1}^K \left(\frac{\sum_{i=1}^n \gamma_{ij}}{n} \right) \|\boldsymbol{\mu}_j - \bar{\mathbf{x}}\|^2.$$

What is the relation between these three quantities?

Based on this relation, show that K -means can be interpreted as minimizing a weighted average of within-cluster variations while approximately(!) maximizing the between-cluster variation. Note that the relation may contain an extra term that does not appear above.

Solution

To simplify the notation, we define $n_j = \sum_{i=1}^n \gamma_{ij}$.

We then have:

$$\begin{aligned}
 V(X) &= \frac{1}{n} \sum_{i=1}^n \|\mathbf{x}_i - \bar{\mathbf{x}}\|^2 = \frac{1}{n} \sum_{i=1}^n \sum_{j=1}^K \gamma_{ij} \|\mathbf{x}_i - \bar{\mathbf{x}}\|^2 = \frac{1}{n} \sum_{j=1}^K \sum_{i=1}^n \gamma_{ij} \|\mathbf{x}_i - \bar{\mathbf{x}}\|^2 \\
 &= \frac{1}{n} \sum_{j=1}^K \sum_{i=1}^n \gamma_{ij} \|\mathbf{x}_i - \boldsymbol{\mu}_j + \boldsymbol{\mu}_j - \bar{\mathbf{x}}\|^2 \\
 &= \frac{1}{n} \sum_{j=1}^K \sum_{i=1}^n \gamma_{ij} \left(\|\mathbf{x}_i - \boldsymbol{\mu}_j\|^2 + \|\boldsymbol{\mu}_j - \bar{\mathbf{x}}\|^2 + 2(\mathbf{x}_i - \boldsymbol{\mu}_j) \cdot (\boldsymbol{\mu}_j - \bar{\mathbf{x}}) \right) \\
 &= \sum_{j=1}^K \frac{n_j}{n} \frac{\sum_{i=1}^n \gamma_{ij} \|\mathbf{x}_i - \boldsymbol{\mu}_j\|^2}{n_j} + \sum_{j=1}^K \frac{n_j \|\boldsymbol{\mu}_j - \bar{\mathbf{x}}\|^2}{n} + \frac{2}{n} \sum_{j=1}^K (\boldsymbol{\mu}_j - \bar{\mathbf{x}}) \cdot \left(\sum_{i=1}^n \gamma_{ij} (\mathbf{x}_i - \boldsymbol{\mu}_j) \right) \\
 &= \sum_{j=1}^K \frac{n_j}{n} V_j(X) + \tilde{V}(X) - \frac{2}{n} \sum_{j=1}^K n_j (\boldsymbol{\mu}_j - \bar{\mathbf{x}}) \cdot (\boldsymbol{\mu}_j - \bar{\boldsymbol{\mu}}_j), \text{ where } \bar{\boldsymbol{\mu}}_j \stackrel{\text{not.}}{=} \frac{\sum_{i=1}^n \gamma_{ij} \mathbf{x}_i}{n_j}.
 \end{aligned}$$

Note: The last equality on the previous slide holds because

$$\begin{aligned}\sum_{i=1}^n \gamma_{ij}(x_i - \mu_j) &= \left(\sum_{i=1}^n \gamma_{ij} x_i \right) - n_j \mu_j = n_j \frac{\sum_{i=1}^n \gamma_{ij} x_i}{n_j} - n_j \mu_j \\ &= n_j \left(\frac{\sum_{i=1}^n \gamma_{ij} x_i}{n_j} - \mu_j \right) = n_j (\bar{\mu}_j - \mu_j)\end{aligned}$$

We already know — see CMU, 2009 spring, Ziv Bar-Joseph, HW5, pr. 2.1 — that K -means aims to minimize J , and consequently $\frac{1}{n}J$, which coincides with the first term in the expression we obtained for $V(X)$, namely $\sum_{j=1}^K \frac{n_j}{n} V_j(X)$.

Since the total variation $V(X)$ is constant, minimizing the first term is equivalent to maximizing the sum of the other two terms, which is expected to be dominated by the between-cluster variation $\tilde{V}(X)$ since a good μ_j should be close to $\bar{\mu}_j$, making the third term small in absolute value.

Exemplifying the application of a simple version of
EM/GMM

on data from \mathbb{R}

$$(\sigma_1 = \sigma_2 = 1, \pi_1 = \pi_2 = 1/2)$$

CMU, 2012 spring, Ziv Bar-Joseph, final exam, pr. 3.1
enhanced by Liviu Ciortuz

Suppose a GMM has two components with known variance and an equal prior distribution

$$\frac{1}{2}N(\mu_1, 1) + \frac{1}{2}N(\mu_2, 1)$$

The observed data are $x_1 = 0.5$ and $x_2 = 2$, and the current estimates of μ_1 and μ_2 are 1 and 2 respectively.

Execute the first iteration of the EM algorithm.

Hint: Normal densities for the standardized variable $y_{(\mu=0, \sigma=1)}$ at 0, 0.5, 1, 1.5, 2 are 0.4, 0.35, 0.24, 0.13, 0.05 respectively.

The E-step:

$$\begin{aligned}
E[Z_{i1}] &= P(Z_{i1} = 1|x_i, \mu) \stackrel{B.Th.}{=} \frac{P(x_i|Z_{i1} = 1, \mu_1)P(Z_{i1} = 1)}{P(x_i|Z_{i1} = 1, \mu_1)P(Z_{i1} = 1) + P(x_i|Z_{i2} = 1, \mu_2)P(Z_{i2} = 1)} \\
&= \frac{P(x_i|Z_{i1} = 1, \mu_1) \cdot \frac{1}{2}}{P(x_i|Z_{i1} = 1, \mu_1) \cdot \frac{1}{2} + P(x_i|Z_{i2} = 1, \mu_2) \cdot \frac{1}{2}} \\
&= \frac{P(x_i|Z_{i1} = 1, \mu_1)}{P(x_i|Z_{i1} = 1, \mu_1) + P(x_i|Z_{i2} = 1, \mu_2)} \text{ for } i \in \{1, 2\}.
\end{aligned}$$

Therefore,

$$\begin{aligned}
P(Z_{11} = 1|x_1, \mu) &= \frac{N(0.5; 1, 1)}{N(0.5; 1, 1) + N(0.5; 2, 1)} = \frac{N(0.5; 0, 1)}{N(0.5; 0, 1) + N(1.5; 0, 1)} = \frac{0.35}{0.35 + 0.13} = \frac{35}{48} \\
P(Z_{21} = 1|x_2, \mu) &= \frac{N(2; 1, 1)}{N(2; 1, 1) + N(2; 2, 1)} = \frac{N(1; 0, 1)}{N(1; 0, 1) + N(0; 0, 1)} = \frac{0.24}{0.24 + 0.4} = \frac{0.24}{0.64} = \frac{3}{8}
\end{aligned}$$

Similarly,

$$\begin{aligned}
P(Z_{12} = 1|x_1, \mu) &= P(Z_{11} = 0|x_1, \mu) = 1 - P(Z_{11} = 1|x_1, \mu) = \frac{13}{48} \\
P(Z_{22} = 1|x_2, \mu) &= P(Z_{21} = 0|x_2, \mu) = 1 - P(Z_{21} = 1|x_2, \mu) = \frac{5}{8}
\end{aligned}$$

The M-step:

$$\mu_j^{(t+1)} = \frac{\sum_{i=1}^2 E[Z_{ij}] x_i}{\sum_{i=1}^2 E[Z_{ij}]} = \frac{\sum_{i=1}^2 P(Z_{ij} = 1|x_i, \mu^{(t)}) x_i}{\sum_{i=1}^2 P(Z_{ij} = 1|x_i, \mu^{(t)})}$$

Therefore,

$$\mu_1^{(1)} = \frac{\frac{35}{48} \cdot 0.5 + \frac{3}{8} \cdot 2}{\frac{35}{48} + \frac{3}{8}} = \frac{107}{106} \approx 1.009 \quad \text{and} \quad \mu_2^{(1)} = \frac{\frac{13}{48} \cdot 0.5 + \frac{5}{8} \cdot 2}{\frac{13}{48} + \frac{5}{8}} = \frac{133}{86} \approx 1.54$$

Derivation of the EM algorithm for
a mixture of K uni-variate Gaussians:
the general case (i.e., when all parameters π, μ, σ^2 are free)

following Dahua Lin,
An Introduction to Expectation-Maximization
(MIT, ML 6768 course, 2012 fall)

Note

We will first consider $K = 2$.

Generalization to $K > 2$ will be shown afterwards.

Estimation (E) Step:

$$\begin{aligned}
 p_{ij} &\stackrel{\text{not.}}{=} P(Z_{ij} = 1 \mid X_i, \mu, \sigma, \pi) \stackrel{\text{calcul}}{=} E[Z_{ij} \mid X_i, \mu, \sigma, \pi] \\
 &= \frac{P(X_i = x_i \mid Z_{ij} = 1, \mu, \sigma, \pi) \cdot P(Z_{ij} = 1 \mid \mu, \sigma, \pi)}{\sum_{j'=1}^2 P(X_i = x_i \mid Z_{ij'} = 1, \mu, \sigma, \pi) \cdot P(Z_{ij'} = 1 \mid \mu, \sigma, \pi)} \\
 &= \frac{\frac{1}{\sqrt{2\pi}\sigma_j} \cdot \exp\left(-\frac{(x_i - \mu_j)^2}{2\sigma_j^2}\right) \cdot \pi_j}{\frac{1}{\sqrt{2\pi}\sigma_1} \cdot \exp\left(-\frac{(x_i - \mu_1)^2}{2\sigma_1^2}\right) \cdot \pi_1 + \frac{1}{\sqrt{2\pi}\sigma_2} \cdot \exp\left(-\frac{(x_i - \mu_2)^2}{2\sigma_2^2}\right) \cdot \pi_2}
 \end{aligned}$$

Therefore, for $t > 0$ we will have:

$$p_{ij}^{(t)} = \frac{\frac{\pi_j^{(t-1)}}{\sigma_j^{(t-1)}} \cdot \exp\left(-\frac{(x_i - \mu_j^{(t-1)})^2}{2(\sigma_j^{(t-1)})^2}\right)}{\frac{\pi_1^{(t-1)}}{\sigma_1^{(t-1)}} \cdot \exp\left(-\frac{(x_i - \mu_1^{(t-1)})^2}{2(\sigma_1^{(t-1)})^2}\right) + \frac{\pi_2^{(t-1)}}{\sigma_2^{(t-1)}} \cdot \exp\left(-\frac{(x_i - \mu_2^{(t-1)})^2}{2(\sigma_2^{(t-1)})^2}\right)}$$

The likelihood of a “complete” instance (x_i, z_{i1}, z_{i2}) :

$$\begin{aligned}
 & P(X_i = x_i, Z_{i1} = z_{i1}, Z_{i2} = z_{i2} \mid \mu, \sigma, \pi) \\
 &= P(X_i = x_i \mid Z_{i1} = z_{i1}, Z_{i2} = z_{i2}, \mu_i, \sigma_i, \pi_i) \cdot P(Z_{i1} = z_{i1}, Z_{i2} = z_{i2} \mid \mu_i, \sigma_i, \pi_i) \\
 &= \frac{1}{\sqrt{2\pi}\sigma_j} \cdot \exp\left(-\frac{(x_i - \mu_j)^2}{2\sigma_j^2}\right) \cdot \pi_j, \text{ where } z_{ij} = 1 \text{ and } z_{ij'} = 0 \text{ for } j' \neq j \\
 &= \frac{1}{\sqrt{2\pi} \sigma_1^{z_{i1}} \sigma_2^{z_{i2}}} \cdot \exp\left(-\frac{1}{2} \sum_{j \in \{1,2\}} z_{ij} \frac{(x_i - \mu_j)^2}{\sigma_j^2}\right) \cdot \pi_1^{z_{i1}} \pi_2^{z_{i2}}
 \end{aligned}$$

The log-likelihood of the same “complete” instance will be:

$$\begin{aligned}
 & \ln P(X_i = x_i, Z_{i1} = z_{i1}, Z_{i2} = z_{i2} \mid \mu, \sigma, \pi) \\
 &= -\frac{1}{2} \ln(2\pi) - \sum_{j=1}^2 z_{ij} \ln \sigma_j - \frac{1}{2} \sum_{j=1}^2 z_{ij} \frac{(x_i - \mu_j)^2}{\sigma_j^2} + \sum_{j=1}^2 z_{ij} \ln \pi_j
 \end{aligned}$$

Given the dataset $X = \{x_1, \dots, x_n\}$, the log-likelihood function will be:

$$\begin{aligned}
 l(\mu, \sigma, \pi) &\stackrel{\text{def.}}{=} \ln P(X, Z_1, Z_2 \mid \mu, \sigma, \pi) \stackrel{i.i.d.}{=} \ln \prod_{i=1}^n P(X_i = x_i, Z_{i1}, Z_{i2} \mid \mu, \sigma, \pi) \\
 &= \sum_{i=1}^n \ln P(X_i = x_i, Z_{i1}, Z_{i2} \mid \mu, \sigma, \pi) \\
 &= -\frac{n}{2} \ln(2\pi) - \sum_{i=1}^n \sum_{j=1}^2 Z_{ij} \ln \sigma_j - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^2 Z_{ij} \frac{(x_i - \mu_j)^2}{\sigma_j^2} + \sum_{i=1}^n \sum_{j=1}^2 Z_{ij} \ln \pi_j
 \end{aligned}$$

The expectation of the log-likelihood function:

$$E[\ln P(X, Z_1, Z_2 \mid \mu, \sigma, \pi)] =$$

$$-\frac{n}{2} \ln(2\pi) - \sum_{i=1}^n \sum_{j=1}^2 E[Z_{ij}] \ln \sigma_j - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^2 E[Z_{ij}] \frac{(x_i - \mu_j)^2}{\sigma_j^2} + \sum_{i=1}^n \sum_{j=1}^2 E[Z_{ij}] \ln \pi_j$$

Here above, the probability function w.r.t. which the expectation was computed was left unspecified. Now we will make it explicit:

$$Q(\mu, \sigma, \pi \mid \mu^{(t)}, \sigma^{(t)}, \pi^{(t)}) \stackrel{not.}{=} E_{Z \mid X, \mu^{(t)}, \sigma^{(t)}, \pi^{(t)}}[\ln P(X, Z_1, Z_2 \mid X, \mu, \sigma, \pi)]$$

$$= -\frac{n}{2} \ln(2\pi) - \sum_{i=1}^n \sum_{j=1}^2 E[Z_{ij} \mid X_i, \mu^{(t)}, \sigma^{(t)}, \pi^{(t)}] \ln \sigma_j$$

$$- \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^2 E[Z_{ij} \mid X_i, \mu^{(t)}, \sigma^{(t)}, \pi^{(t)}] \frac{(x_i - \mu_j)^2}{\sigma_j^2}$$

$$+ \sum_{i=1}^n \sum_{j=1}^2 E[Z_{ij} \mid X_i, \mu^{(t)}, \sigma^{(t)}, \pi^{(t)}] \ln \pi_j$$

$$p_{ij}^{(t)} \stackrel{not.}{=} E[Z_{ij} \mid X, \mu^{(t-1)}, \sigma^{(t-1)}, \pi^{(t-1)}] \Rightarrow$$

$$Q(\mu, \sigma, \pi \mid \mu^{(t)}, \sigma^{(t)}, \pi^{(t)}) =$$

$$-\frac{n}{2} \ln(2\pi) - \sum_{i=1}^n \sum_{j=1}^2 p_{ij}^{(t)} \ln \sigma_j - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^2 p_{ij}^{(t)} \frac{(x_i - \mu_j)^2}{\sigma_j^2} + \sum_{i=1}^n \sum_{j=1}^2 p_{ij}^{(t)} \ln \pi_j$$

Since $K = 2$ and $\pi_1 + \pi_2 = 1$, we get

$$Q(\mu, \sigma, \pi \mid \mu^{(t)}, \sigma^{(t)}, \pi^{(t)}) =$$

$$-\frac{n}{2} \ln 2\pi - \sum_{i=1}^n \sum_{j=1}^2 p_{ij}^{(t)} \ln \sigma_j - \frac{1}{2} \sum_{i=1}^n \sum_{j=1}^2 p_{ij}^{(t)} \frac{(x_i - \mu_j)^2}{\sigma_j^2} + \sum_{i=1}^n (p_{i1}^{(t)} \ln \pi_1 + p_{i2}^{(t)} \ln(1 - \pi_1))$$

Maximization (M) Step:

[For $K = 2$:]

$$\begin{aligned} \frac{\partial}{\partial \pi_1} Q(\mu, \sigma, \pi \mid \mu^{(t)}, \sigma^{(t)}, \pi^{(t)}) = 0 &\Leftrightarrow \frac{1}{\pi_1} \sum_{i=1}^n p_{i1}^{(t)} = \frac{1}{1 - \pi_1} \sum_{i=1}^n p_{i2}^{(t)} \Leftrightarrow \\ \sum_{i=1}^n p_{i1}^{(t)} = \pi_1 \left(\sum_{i=1}^n p_{i1}^{(t)} + \sum_{i=1}^n p_{i2}^{(t)} \right) &\Leftrightarrow \sum_{i=1}^n p_{i1}^{(t)} = \pi_1 \sum_{i=1}^n \underbrace{(p_{i1}^{(t)} + p_{i2}^{(t)})}_1 \Leftrightarrow \sum_{i=1}^n p_{i1}^{(t)} = n\pi_1 \\ \Rightarrow \pi_1^{(t+1)} &\leftarrow \frac{1}{n} \sum_{i=1}^n p_{i1}^{(t)} \end{aligned}$$

Taking into account that $\pi_1^{(t+1)} + \pi_2^{(t+1)} = 1$ and $p_{i1}^{(t)} + p_{i2}^{(t)} = 1$ for $i = 1, \dots, n$,

$$\Rightarrow \pi_2^{(t+1)} \leftarrow \frac{1}{n} \sum_{i=1}^n p_{i2}^{(t)}$$

$$\frac{\partial}{\partial \mu_1} Q(\mu, \sigma, \pi \mid \mu^{(t)}, \sigma^{(t)}, \pi^{(t)}) = 0 \Leftrightarrow \frac{1}{\sigma_1^2} \sum_{i=1}^n p_{i1}^{(t)} (x_i - \mu_1) = 0 \Leftrightarrow \sum_{i=1}^n p_{i1}^{(t)} (x_i - \mu_1) = 0$$

$$\Rightarrow \mu_1^{(t+1)} \leftarrow \frac{\sum_{i=1}^n p_{i1}^{(t)} x_i}{\sum_{i=1}^n p_{i1}^{(t)}}$$

Similarly,
$$\mu_2^{(t+1)} \leftarrow \frac{\sum_{i=1}^n p_{i2}^{(t)} x_i}{\sum_{i=1}^n p_{i2}^{(t)}}$$

$$\frac{\partial}{\partial \sigma_1} Q(\mu, \sigma, \pi \mid \mu^{(t)}, \sigma^{(t)}, \pi^{(t)}) = 0 \Leftrightarrow -\frac{1}{\sigma_1} \sum_{i=1}^n p_{i1}^{(t)} + \frac{1}{\sigma_1^3} \sum_{i=1}^n p_{i1}^{(t)} (x_i - \mu_1)^2 = 0,$$

$$\Rightarrow \left(\sigma_1^{(t+1)} \right)^2 \leftarrow \frac{\sum_{i=1}^n p_{i1}^{(t)} (x_i - \mu_1^{(t+1)})^2}{\sum_{i=1}^n p_{i1}^{(t)}}$$

$$\text{Similarly, } \left(\sigma_2^{(t+1)} \right)^2 \leftarrow \frac{\sum_{i=1}^n p_{i2}^{(t)} (x_i - \mu_2^{(t+1)})^2}{\sum_{i=1}^n p_{i2}^{(t)}}$$

Note: One could relatively easy prove that these solutions (namely, $\pi^{(t+1)}, \mu^{(t+1)}, \sigma^{(t+1)}$) of the partial derivatives of the *auxiliary function* Q designate the values for which Q reaches its *maximum*.

Generalization to $K > 2$

In this case, the Bernoulli distribution is replaced by a categorical one. The only one change needed in the above proof concerns updating the parameters of this distribution.

Since $\pi_1 + \dots + \pi_K = 1$, we must solve the following *constraint optimization problem*:

$$\begin{aligned} & \max_{\pi, \mu, \sigma} Q(\pi, \mu, \sigma | \pi^{(t)}, \mu^{(t)}, \sigma^{(t)}) \\ & \text{subject to } \sum_{i=1}^K \pi_j = 1 \text{ and } \pi_j \geq 0, \forall j = 1, \dots, K. \end{aligned}$$

By letting asside the \geq constraints, and using the *Lagrangean multiplier* $\lambda \in \mathbb{R}$, this problem becomes:

$$\max_{\pi, \mu, \sigma} \left(Q(\pi, \mu, \sigma | \pi^{(t)}, \mu^{(t)}, \sigma^{(t)}) + \lambda \left(1 - \sum_{i=1}^K \pi_j \right) \right).$$

For $j = 1, \dots, K$:

$$\frac{\partial}{\partial \pi_j} Q(\mu, \sigma, \pi \mid \mu^{(t)}, \sigma^{(t)}, \pi^{(t)}) = 0 \Leftrightarrow \sum_{i=1}^n p_{ij}^{(t)} \frac{1}{\pi_j} = \lambda \Leftrightarrow \pi_j^{(t+1)} = \frac{1}{\lambda} \sum_{i=1}^n p_{ij}^{(t)}.$$

Because $\sum_{j=1}^K \pi_j^{(t+1)} = 1$, it follows that

$$\lambda = \sum_{j=1}^K \sum_{i=1}^n \pi_j^{(t+1)} = \sum_{i=1}^n \underbrace{\sum_{j=1}^K \pi_j^{(t+1)}}_1 = \sum_{i=1}^n 1 = n.$$

Therefore,

$$\pi_j^{(t+1)} \leftarrow \frac{1}{n} \sum_{i=1}^n p_{ij}^{(t)}.$$

Note that indeed $\pi_j^{(t+1)} \geq 0$, because the $p_{ij}^{(t)}$ terms designate some probabilities (see E-step).

To summarize:

E Step:

$$p_{ij}^{(t)} \stackrel{\text{not.}}{=} P(z_{ij} = 1 \mid x_i; \mu^{(t)}, (\sigma^2)^{(t)}, \pi^{(t)}) = \frac{N(x_i \mid \mu_j^{(t)}, (\sigma_j^2)^{(t)}) \cdot \pi_j^{(t)}}{\sum_{l=1}^K N(x_i \mid \mu_l^{(t)}, (\sigma_l^2)^{(t)}) \cdot \pi_l^{(t)}}$$

where $N(x_i \mid \mu_j, \sigma_j^2) \stackrel{\text{def.}}{=} \frac{1}{\sqrt{2\pi}\sigma_j} \cdot \exp\left(-\frac{(x_i - \mu_j)^2}{2\sigma_j^2}\right).$

M Step:

$$\pi_j^{(t+1)} \leftarrow \frac{1}{n} \sum_{i=1}^n p_{ij}^{(t)}$$

$$\mu_j^{(t+1)} \leftarrow \frac{\sum_{i=1}^n p_{ij}^{(t)} x_i}{\sum_{i=1}^n p_{ij}^{(t)}}$$

$$\left(\sigma_j^{(t+1)}\right)^2 \leftarrow \frac{\sum_{i=1}^n p_{ij}^{(t)} (x_i - \mu_j^{(t+1)})^2}{\sum_{i=1}^n p_{ij}^{(t)}}$$

Exemplifying

some **methodological issues** regarding the application of
the **EM algorithmic schema**

(using a simple EM/GMM algorithm on data from \mathbb{R} ($\pi_1 = \pi_2 = 1/2$))

CMU, 2007 spring, Eric Xing, final exam, pr. 1.8

A long time ago there was a village amidst hundreds of lakes. Two types of fish lived in the region, but only one type in each lake.

These types of fish both looked exactly the same, smelled exactly the same when cooked, and had the exact same delicious taste – except one was poisonous and would kill any villager who ate it. The only other difference between the fish was their effect on the pH (acidity) of the lake they occupy.

The pH for lakes occupied by the non-poisonous type of fish was distributed according to a Gaussian with unknown mean (μ_{safe}) and variance (σ_{safe}^2) and the pH for lakes occupied by the poisonous type was distributed according to a different Gaussian with unknown mean (μ_{deadly}) and variance (σ_{deadly}^2). (Poisonous fish tended to cause slightly more acidic conditions).

Naturally, the villagers turned to machine learning for help. However, there was much debate about the right way to apply EM to their problem. For each of the following procedures, indicate whether it is an accurate implementation of Expectation-Maximization and will provide a reasonable estimate for parameters μ and σ^2 for each class.

a.

Guess initial values of μ and σ^2 for each class.

(1) For each lake, find the most likely class of fish for the lake.

(2) Update the μ and σ^2 values using their maximum likelihood estimates based on these predictions.

Iterate (1) and (2) until convergence.

b.

For each lake, guess an initial probability that it is safe.

(1) Using these probabilities, find the maximum likelihood estimates for the μ and σ values for each class.

(2) Use these estimates of μ and σ to reestimate lake safety probabilities.

Iterate (1) and (2) until convergence.

c.

Compute the mean and variance of the pH levels across all lakes.

Use these values for the μ and σ^2 value of each class of fish.

(1) Use the μ and σ^2 values of each class to compute the belief that each lake contains poisonous fish.

(2) Find the maximum likelihood values for μ and σ^2 .

Iterate (1) and (2) until convergence.

Solution

- a. It'll do ok if we give sensible enough μ and σ^2 initial values.
- b. Ok, this is the same as *a* after the first M-step.
(See the general EM algorithmic schema on the next slide.)
- c. This will be stuck at the initial μ and σ^2 :


In the E-step we'll get:

$$P(\text{safe}|x) = \frac{P(x|\text{safe}) \cdot P(\text{safe})}{P(x|\text{safe}) \cdot P(\text{safe}) + P(x|\text{deadly}) \cdot P(\text{deadly})} = \frac{1}{2}$$

since on one side we assume $P(\text{safe}) = P(\text{poison}) = \frac{1}{2}$ and on the other side $P(x|\text{safe}) = P(x|\text{poison})$ because $\mu_{\text{safe}} = \mu_{\text{deadly}}$ and $\sigma_{\text{safe}}^2 = \sigma_{\text{deadly}}^2$.

In the M-step μ and σ^2 will not change since we are again letting them be calculated from all lakes (weighted equally).

The [general] EM algorithmic schema



The EM algorithm for modeling
mixtures of multi-variate Gaussians

Stanford University, Prof. Andrew Ng
ML course, 2009, lecture notes, parts VIII and IX
[adapted by Liviu Ciortuz]

Suppose that we are given the *instances* $x_1, \dots, x_n \in \mathbb{R}^d$ (all seen as column-vectors). We wish to *model* these data by specifying a joint distribution $p(x_i, z_i) = p(x_i|z_i) \cdot p(z_i)$. Here,

$z_i \sim \text{Categorical}(\pi)$,

K denotes the number of values that the z_i 's can take on, namely $\pi_j \stackrel{\text{not.}}{=} p(z_i = j)$ for $j = 1, \dots, K$, with $\sum_{j=1}^K \pi_j = 1$, and the [conditional] distribution $x_i|z_i = j$ is a Gaussian of mean vector μ_j and covariance matrix Σ_j .

Thus, our model posits that each x_i was *generated* by randomly choosing z_i from $\{1, \dots, K\}$, and then x_i was drawn from one of the K Gaussians, depending on z_i . This is called *the mixture of [multi-variate] Gaussians* model. Remember that

$$\mathcal{N}(x; \mu, \sigma) = \frac{1}{(2\pi)^{d/2} |\Sigma|^{1/2}} \exp \left(-\frac{1}{2} (x - \mu)^\top \Sigma^{-1} (x - \mu) \right)$$

Note that the z_i 's are *latent* random variables, meaning that they're hidden/unobserved.

[Use the EM general scheme (see Tom Mitchell's *Machine Learning* book, 1997, pag. 194-195) to] prove that the EM algorithm for estimating the parameters π, μ and Σ of our mixture of multi-variate Gaussian distributions has the following *update rules*:

E-step:

$$w_{ij} \stackrel{not.}{=} E[z_i = j | x_i; \pi', \mu', \Sigma'] = \frac{\mathcal{N}(x_i; \mu', \Sigma') \pi_j}{\sum_{l=1}^K \mathcal{N}(x_i; \mu', \Sigma') \pi_l} \quad (4)$$

M-step:

$$\pi_j = \frac{1}{n} \sum_{i=1}^n w_{ij}, \quad (5)$$

$$\mu_j = \frac{\sum_{i=1}^n w_{ij} x_i}{\sum_{i=1}^n w_{ij}}. \quad (6)$$

$$\Sigma_j = \frac{\sum_{i=1}^n w_{ij} (x_i - \mu_j)(x_i - \mu_j)^\top}{\sum_{i=1}^n w_{ij}}. \quad (7)$$

where π', μ' and Σ' represent the values of our parameters at initialization, and respectively the previous iteration of the EM algorithm.

Hint: You may find useful the following formulas (from *Matrix Identities*, by Sam Roweis, 1999):

$$(1e) \quad (A^{-1})^\top = (A^\top)^{-1}$$

$$(2b) \quad |A^{-1}| = \frac{1}{|A|}$$

$$(4a) \quad \frac{\partial}{\partial X} |AXB| = |AXB|(X^{-1})^\top = |AXB|(X^\top)^{-1}$$

$$(4b) \quad \frac{\partial}{\partial X} \ln |X| = (X^{-1})^\top = (X^\top)^{-1}$$

$$(5a) \quad \frac{\partial}{\partial X} a^\top X = \frac{\partial}{\partial X} X^\top a = a$$

$$(5b) \quad \frac{\partial}{\partial X} X^\top AX = (A + A^\top)X$$

$$(5c) \quad \frac{\partial}{\partial X} a^\top Xb = ab^\top$$

$$(5e) \quad \frac{\partial}{\partial X} a^\top Xa = \frac{\partial}{\partial X} a^\top X^\top a = aa^\top$$

$$(5g) \quad \frac{\partial}{\partial X} (Xa + b)^\top C(Xa + b) = (C + C^\top)(Xa + b)a^\top$$

Solution

The *E-step* is easy (use Bayes rule):

$$\begin{aligned}
 w_{ij} &\stackrel{not.}{=} E[z_i = j | x_i; \pi', \mu', \Sigma'] = p(z_i = j | x_i; \pi', \mu', \Sigma') = \\
 &= \frac{p(x_i | z_i = j; \mu', \Sigma') p(z_i = j; \pi')}{\sum_{l=1}^K p(x_i | z_i = l; \mu', \Sigma') p(z_i = l; \pi')} = \frac{\mathcal{N}(x_i; \mu', \Sigma') \pi'_j}{\sum_{l=1}^K \mathcal{N}(x_i; \mu', \Sigma') \pi_l}
 \end{aligned}$$

We will now concentrate on the *M-step*:

According to *the general EM scheme*, we need to maximize, with respect to our parameters π, μ, Σ , the “auxiliary” function

$$\begin{aligned}
 Q(\pi, \mu, \Sigma | \pi', \mu', \Sigma') &\stackrel{\text{def.}}{=} E_{p(z_i=j|x_i; \pi', \mu', \Sigma')} \ln p(x, z; \mu, \Sigma, \pi) \\
 &= \sum_{i=1}^n \sum_{j=1}^K p(z_i = j | x_i; \pi', \mu', \Sigma') \ln p(x_i, z_i = j; \mu, \Sigma, \pi) \\
 &= \sum_{i=1}^n \sum_{j=1}^K p(z_i = j | x_i; \pi', \mu', \Sigma') \ln(p(x_i | z_i = j; \mu, \Sigma) p(z_i = j; \pi)) \\
 &= \sum_{i=1}^n \sum_{j=1}^K w_{ij} \ln \frac{1}{(2\pi)^{d/2} |\Sigma_j|^{1/2}} \cdot \exp\left(-\frac{1}{2}(x_i - \mu_j)^\top \Sigma_j^{-1} (x_i - \mu_j)\right) \cdot \pi_j \\
 &= \sum_{i=1}^n \sum_{j=1}^K w_{ij} \left[-\ln((2\pi)^{d/2} |\Sigma_j|^{1/2}) - \frac{1}{2}(x_i - \mu_j)^\top \Sigma_j^{-1} (x_i - \mu_j) + \ln \pi_j \right] \quad (8)
 \end{aligned}$$

First, let's derive the M-step update rule for μ_l , with $l = 1, \dots, K$.

We have to maximize (8) with respect to μ_l , so let's compute the corresponding derivative:

$$\begin{aligned}
 & \frac{\partial}{\partial \mu_l} \sum_{i=1}^n \sum_{j=1}^K w_{ij} \left[-\ln((2\pi)^{d/2} |\Sigma_j|^{1/2}) - \frac{1}{2} (x_i - \mu_j)^\top \Sigma_j^{-1} (x_i - \mu_j) + \ln \pi_j \right] \\
 &= -\frac{\partial}{\partial \mu_l} \sum_{i=1}^n \sum_{j=1}^K w_{ij} \frac{1}{2} (x_i - \mu_j)^\top \Sigma_j^{-1} (x_i - \mu_j) = -\frac{1}{2} \sum_{i=1}^n w_{ij} \frac{\partial}{\partial \mu_l} \sum_{i=1}^n (x_i - \mu_j)^\top \Sigma_j^{-1} (x_i - \mu_j) \\
 &\stackrel{(5g)}{=} -\frac{1}{2} \sum_{i=1}^n w_{ij} \frac{\partial}{\partial \mu_l} \sum_{i=1}^n (\Sigma_j^{-1} + (\Sigma_j^{-1})^\top) (x_i - \mu_j) \\
 &\stackrel{(1e)}{=} \frac{1}{2} \sum_{i=1}^n w_{ij} \sum_{i=1}^n (\Sigma_j^{-1} + \underbrace{(\Sigma_j^\top)^{-1}}_{\Sigma_j}) (x_i - \mu_j) = \frac{1}{2} \sum_{i=1}^n w_{ij} \sum_{i=1}^n 2 \Sigma_j^{-1} (x_i - \mu_j) \\
 &= \sum_{i=1}^n w_{il} (\Sigma_l^{-1} x_i - \Sigma_l^{-1} \mu_l) = \sum_{i=1}^n w_{il} \Sigma_l^{-1} x_i - \sum_{i=1}^n w_{il} \Sigma_l^{-1} \mu_l.
 \end{aligned}$$

Setting this to zero and solving for μ_l therefore yields the update rule

$$\mu_l = \frac{\sum_{i=1}^n w_{il} x_i}{\sum_{i=1}^n w_{il}}.$$

Secondly, we'll derive the M-step updates to Σ_j , for $j = 1, \dots, K$. Grouping together only the terms that depend on Σ_j in (8), we find that we need to maximize

$$\begin{aligned} & \sum_{i=1}^n \sum_{j=1}^K w_{ij} \left[\ln \frac{1}{|\Sigma_j|^{1/2}} - \frac{1}{2} (x_i - \mu_j)^\top \Sigma_j^{-1} (x_i - \mu_j) \right] \\ & \stackrel{(2b)}{=} \sum_{i=1}^n \sum_{j=1}^K w_{ij} \left[\ln |\Sigma_j^{-1}|^{1/2} - \frac{1}{2} (x_i - \mu_j)^\top \Sigma_j^{-1} (x_i - \mu_j) \right]. \end{aligned}$$

We use the usual trick of working with the precision matrix $\Lambda_j \stackrel{not.}{=} \Sigma_j^{-1}$, where Σ_j is assumed invertible.

When maximizing the above quantity with respect to Λ_j by taking derivatives, we find:

$$\begin{aligned}
& \frac{\partial}{\partial \Lambda_j} \sum_{i=1}^n w_{ij} \left[\ln |\Lambda_j|^{1/2} - \frac{1}{2} (x_i - \mu_j)^\top \Lambda_j (x_i - \mu_j) \right] \\
&= \frac{1}{2} \sum_{i=1}^n w_{ij} \frac{\partial}{\partial \Lambda_j} \ln |\Lambda_j| - \frac{1}{2} \sum_{i=1}^n w_{ij} \frac{\partial}{\partial \Lambda_j} [(x_i - \mu_j)^\top \Lambda_j (x_i - \mu_j)] \\
&\stackrel{(4b),(5c)}{=} \frac{1}{2} \sum_{i=1}^n w_{ij} \Lambda_j^{-1} - \frac{1}{2} \sum_{i=1}^n w_{ij} (x_i - \mu_j)(x_i - \mu_j)^\top \\
&= \frac{1}{2} \Lambda_j^{-1} \sum_{i=1}^n w_{ij} - \frac{1}{2} \sum_{i=1}^n w_{ij} (x_i - \mu_j)(x_i - \mu_j)^\top.
\end{aligned}$$

Setting this to zero and solving, we get:

$$\Sigma_j = \Lambda_j^{-1} = \frac{\sum_{i=1}^n w_{ij} (x_i - \mu_j)(x_i - \mu_j)^\top}{\sum_{i=1}^n w_{ij}}.$$

Finally, let's derive the M-step update for the parameters π_j . Grouping together only the terms that depend on π_j in (8), we find that we need to maximize

$$\sum_{i=1}^n \sum_{j=1}^K w_{ij} \ln \pi_j.$$

However, there is an additional constraint that the π_j 's sum to 1, since they represent the probabilities $\pi_j = p(z_i = j; \pi)$. To deal with the constraint that $\sum_{j=1}^K \pi_j = 1$, we construct the Lagrangian

$$\mathcal{L}(\pi) = \sum_{i=1}^n \sum_{j=1}^K w_{ij} \ln \pi_j + \beta \left(\sum_{j=1}^K \pi_j - 1 \right),$$

where β is the Lagrange multiplier.

Note: We don't need to worry about the constraint that $\pi_j \geq 0$, because as we'll shortly see, the solution we'll find from this derivation will automatically satisfy that anyway.

Taking derivatives of $\mathcal{L}(\pi)$, we find:

$$\frac{\partial}{\partial \pi_j} \mathcal{L}(\pi) = \sum_{i=1}^n \frac{w_{ij}}{\pi_j} + \beta = \frac{1}{\pi_j} \sum_{i=1}^n w_{ij} + \beta.$$

Setting this to zero and solving, we get $\pi_j = \frac{\sum_{i=1}^n w_{ij}}{-\beta}$.

By using the constraint $\sum_j \pi_j = 1$, and given the fact that $\sum_j w_{ij} = 1$ since $w_{ij} \stackrel{\text{not.}}{=} p(z_i = j | x_i; \pi', \mu', \Sigma')$, we easily find

$$-\beta = \sum_{i=1}^n \sum_{j=1}^K w_{ij} = \sum_{i=1}^n 1 = n.$$

We therefore have our M-step derivation for the parameters π_j :

$$\pi_j = \frac{1}{n} \sum_{i=1}^n w_{ij},$$

and, obviously, $\pi_j \geq 0$.

Remarks

1. Let's contrast the update rules in the M-step with the formulas we would have when the z_i s were known exactly (see the MLE of the parameters of a single multi-variate Gaussian distribution, CMU, 2010 fall, Aarti Singh, HW1, pr. 3.2.1):

$$\begin{aligned}\pi_j &= \frac{1}{n} \sum_{i=1}^n 1\{z_i = j\}, \\ \mu_j &= \frac{\sum_{i=1}^n 1\{z_i = j\} x_i}{\sum_{i=1}^n 1\{z_i = j\}}, \\ \Sigma_j &= \frac{\sum_{i=1}^n 1\{z_i = j\} (x_i - \mu_j)(x_i - \mu_j)^\top}{\sum_{i=1}^n 1\{z_i = j\}},\end{aligned}$$

with $1\{z_i = j\}$ (“indicator functions”) indicating from which Gaussian each datapoint had come.

They are identical, except that instead of the indicator functions $1\{z_i = j\}$ indicating from which Gaussian each datapoint had come, we now have the w_{ij} s.

Remarks (cont'd)

2. The EM-algorithm is reminiscent of the K -means clustering algorithm, except that instead of the “hard” cluster assignments $c(i)$, we have the “soft” assignments w_{ij} .
3. Similar to K -means, the EM algorithm is also susceptible to local optima, so reinitializing at several different initial parameters may be a good idea.
4. It's clear that the EM algorithm has a very natural interpretation of repeatedly trying to guess the unknown z_i 's.

A link between
K-means and EM/GMM (the multi-variate case)

CMU, 2008 fall, Eric Xing, HW4, pr. 2.2

(see also CMU, 2010 fall, Aarti Singh, HW4, pr. 1.2)

Given N data points x_i , ($i = 1, \dots, N$), K -means will group them into K clusters by minimizing the *distortion* function

$$J = \sum_{i=1}^N \sum_{j=1}^K \gamma_{ij} \|x_i - \mu_j\|^2,$$

where μ_j is the centroid of the j -th cluster, and $\gamma_{ij} = 1$ if x_i belongs to the j -th cluster and $\gamma_{ij} = 0$ otherwise.

In this exercise, we will use the following procedure for K -means:

- Initialize [randomly] the cluster centroids μ_j , $j = 1, \dots, K$;
- Iterate until *convergence*:
 - for every data point x_n , update its *cluster assignment*: $\gamma_{ij} = 1$ if $j = \arg \min_k \|x_i - \mu_k\|^2$, and $\gamma_{ij} = 0$ otherwise.
 - for each cluster j , update its centroid: $\mu_j = \frac{\sum_{i=1}^N \gamma_{ij} x_i}{\sum_{i=1}^N \gamma_{ij}}$

Remember that in GMM, $p(x) = \sum_{j=1}^K \pi_j \mathcal{N}(x|\mu_j, \Sigma_j)$, where π_j is the prior [probability] for the j^{th} component, μ_j and Σ_j are the mean and covariance matrix for the j^{th} component respectively. In the E-step of the EM algorithm, we will update

$$p(z_{ij} = 1|x_i) = \frac{\pi_j \mathcal{N}(x_i|\mu_j, \Sigma_j)}{\sum_{k=1}^K \pi_k \mathcal{N}(x_i|\mu_k, \Sigma_k)}$$

Now, suppose that

- i. $\Sigma_k = \sigma^2 I$, for some $\sigma > 0$, and for all $k = 1, \dots, K$
- ii. $\pi_k \neq 0$ for $k = 1, \dots, K$ [LC: at any iteration of the EM algorithm], and
- iii. $\|x_i - \mu_{k'}\| \neq \|x_i - \mu_k\|$ for any $k' \neq k$ [at any iteration of the EM algorithm].

Under the above assumptions, prove that when $\sigma \rightarrow 0_+$ we will get $p(z_{ij} = 1|x_i) \rightarrow \gamma_{ij}$, where γ_{ij} is the *cluster assignment* used in K -means.

Answer:

$$\begin{aligned}
 p(z_{ij} = 1|x_i) &= \frac{\pi_j \mathcal{N}(x_i|\mu_j, \Sigma_j)}{\sum_{k=1}^K \pi_k \mathcal{N}(x_i|\mu_k, \Sigma_k)} = \frac{\pi_j \exp\left(-\frac{1}{2\sigma^2}\|x_i - \mu_j\|^2\right)}{\sum_{k=1}^K \pi_k \exp\left(-\frac{1}{2\sigma^2}\|x_i - \mu_k\|^2\right)} \\
 &= \frac{1}{1 + \sum_{k \neq j} \frac{\pi_k}{\pi_j} \exp\left(\frac{1}{2\sigma^2}(\|x_i - \mu_j\|^2 - \|x_i - \mu_k\|^2)\right)}
 \end{aligned}$$

Case 1:

If $\|x_i - \mu_j\| = \min_k \|x_i - \mu_k\|$, then for each $k \neq j$ we have $\|x_i - \mu_j\|^2 - \|x_i - \mu_k\|^2 < 0$. Since $\sigma \rightarrow 0_+$, it will follow that $\exp\left(\frac{1}{2\sigma^2}(\|x_i - \mu_j\|^2 - \|x_i - \mu_k\|^2)\right) \rightarrow 0$. So, $p(z_{ij} = 1|x_i) \rightarrow 1$.

Case 2:

If $\|x_i - \mu_j\| \neq \min_k \|x_i - \mu_k\|$, then

- **for all k such that $\|x_i - \mu_j\| < \|x_i - \mu_k\|$ it will follow (exactly as above) that $\exp\left(\frac{1}{2\sigma^2}(\|x_i - \mu_j\|^2 - \|x_i - \mu_k\|^2)\right) \rightarrow 0$;**
- **for all k such that $\|x_i - \mu_j\| > \|x_i - \mu_k\|$ we will have $\exp\left(\frac{1}{2\sigma^2}(\|x_i - \mu_j\|^2 - \|x_i - \mu_k\|^2)\right) \rightarrow +\infty$.**

Therefore, $p(z_{ij} = 1|x_i) \rightarrow \frac{1}{1 + \infty} = 0$.

**Algoritmul EM, fundamentare teoretică:
pasul E [și pasul M]**

CMU, 2008 fall, Eric Xing, HW4, pr. 1.1-3

Algoritmul EM (Expectation-Maximization) permite crearea unor modele probabiliste care pe de o parte depind de un set de parametri θ iar pe de altă parte includ pe lângă variabilele obișnuite (“observabile” sau “vizibile”) x și variabile necunoscute (“neobservabile”, “ascunse” sau “latente”) z .

În general, în astfel de situații/modele, nu se poate face în mod direct o estimare a parametrilor modelului (θ), în așa fel încât să se garanteze atingerea maximului verosimilității datelor observabile x .

În schimb, algoritmul EM procedează în manieră iterativă, constituind astfel o modalitate foarte convenabilă de estimare a parametrilor θ .

Definim log-verosimilitatea datelor observabile (x) ca fiind $\log P(x \mid \theta)$, iar log-verosimilitatea datelor *complete* (observabile, x , și neobservabile, z) ca fiind $\log P(x, z \mid \theta)$.

Observație: Pe tot parcursul acestui exercițiu se va considera funcția log ca având baza supraunitară, fixată.

a. Log-verosimilitatea datelor *observabile* (x) se poate exprima în funcție de datele neobservabile (z), astfel:^a

$$\ell(\theta) \stackrel{not.}{=} \log P(x \mid \theta) = \log \left(\sum_z P(x, z \mid \theta) \right)$$

În continuare vom nota cu q o funcție / distribuție de probabilitate definită peste variabilele ascunse/neobservabile z .

Folosiți *inegalitatea lui Jensen* pentru a demonstra că are loc următoarea inegalitate:

$$\log P(x \mid \theta) \geq \sum_z q(z) \log \left(\frac{P(x, z \mid \theta)}{q(z)} \right) \quad (9)$$

pentru orice x (fixat), pentru orice valoare a parametrului θ și pentru orice distribuție probabilistă q definită peste variabilele neobservabile z .

^aRețineți că x , vectorul de date observabile, este fixat (dat), în vreme ce z , vectorul de date neobservabile, este liber (variabil).

Observație (1)

Semnificația inegalității (9) este următoarea:

Funcția (de fapt, orice funcție de forma)

$$F(q, \theta) \stackrel{\text{def.}}{=} \sum_z q(z) \log \left(\frac{P(x, z | \theta)}{q(z)} \right)$$

constituie o margină inferioară pentru funcția de log-verosimilitate a datelor incomplete / observabile, $\ell(\theta) \stackrel{\text{not.}}{=} \log P(x | \theta)$.

Remarcați faptul că F este o funcție de două variabile, iar prima variabilă nu este de tip numeric (cum este θ), ci este de tip funcțional.

Mai mult, se observă că expresia funcției F este de fapt o medie,

$$E_{q(z)} \left[\log \frac{P(x, z | \theta)}{q(z)} \right],$$

atunci când x , q și parametrul θ se consideră fixați, iar z este lăsat să varieze.

Soluție

Inegalitatea lui Jensen, în contextul teoriei probabilităților:

considerând X este o variabilă aleatoare (unară),

dacă φ este o funcție (reală) convexă, atunci $\varphi(E[X]) \leq E[\varphi(X)]$;

dacă φ este funcție concavă, atunci $\varphi(E[X]) \geq E[\varphi(X)]$.

Aici vom folosi funcția \log cu bază supraunitară (funcție concavă), deci aplicând inegalitatea lui Jensen vom obține: $\log(E[X]) \geq E[\log(X)]$.

Log-verosimilitatea datelor observabile este:

$$\begin{aligned} \ell(\theta) \stackrel{not.}{=} \log P(x | \theta) &= \log \left(\sum_z P(x, z | \theta) \right) = \log \left(\sum_z q(z) \frac{P(x, z | \theta)}{q(z)} \right) \\ &\stackrel{def.}{=} \log \left(E_{q(z)} \left[\frac{P(x, z | \theta)}{q(z)} \right] \right) \end{aligned}$$

Conform inegalității lui Jensen (înlocuind X de mai sus cu $\frac{P(x, z | \theta)}{q(z)}$), rezultă:

$$\ell(\theta) \stackrel{not.}{=} \log P(x | \theta) \geq E_{q(z)} \left[\log \frac{P(x, z | \theta)}{q(z)} \right] \stackrel{def.}{=} \sum_z q(z) \log \frac{P(x, z | \theta)}{q(z)},$$

Notăție

În continuare, pentru a vă aduce mereu aminte că distribuția q se referă la datele neobservabile z , vom folosi notația $q(z)$ în loc de q .

În consecință, în cele ce urmează, în funcție de context, $q(z)$ va desemna fie la distribuția q , fie la valoarea acestei distribuții pentru o valoare oarecare [a variabilei neobservabile] z .

(Este adevărat că această lejeră ambiguitate poate induce în eroare cititorul neexperimentat.)

b. Vă reamintim definiția entropiei relative (numită și divergența Kullback-Leibler):

$$KL(q(z) \parallel P(z \mid x, \theta)) = - \sum_z q(z) \log \left(\frac{P(z \mid x, \theta)}{q(z)} \right)$$

Arătați că

$$\log P(x \mid \theta) = F(q(z), \theta) + KL(q(z) \parallel P(z \mid x, \theta)).$$

Observație (2):

Semnificația egalității care trebuie demonstrată la acest punct este foarte interesantă: diferența dintre funcția obiectiv $\ell(\theta) \stackrel{not.}{=} \log P(x \mid \theta)$ și marginea sa inferioară $F(q(z), \theta)$ — a se vedea punctul a — este $KL(q(z) \parallel P(z \mid x, \theta))$. Tocmai pe această chestiune se va “construi” punctul final, și cel mai important, al problemei noastre.

Observație (3)

Ideile de bază ale algoritmului EM sunt două:

1. În loc să calculeze maximul funcției de log-verosimilitate $\log P(x \mid \theta)$ în raport cu θ , algoritmul EM va maximiza marginea sa inferioară, $F(q(z), \theta)$, în raport cu ambele argumente, $q(z)$ și θ .
2. Pentru a căuta maximul (de fapt, un maxim local al) marginii inferioare $F(q(z), \theta)$, algoritmul EM aplică metoda *creșterii pe coordonate* (engl., coordinate ascent): după ce inițial se fixează $\theta^{(0)}$ eventual aleatoriu, se maximizează *iterativ* funcția $F(q(z), \theta)$, în mod *alternativ*: mai întâi în raport cu distribuția $q(z)$ și apoi în raport cu parametrul θ .

$$\text{Pasul E:} \quad q^{(t)}(z) = \operatorname{argmax}_{q(z)} F(q(z), \theta^{(t)})$$

$$\text{Pasul M:} \quad \theta^{(t+1)} = \operatorname{argmax}_{\theta} F(q^{(t)}(z), \theta)$$

Soluție

$$\begin{aligned}
 F(q(z), \theta) &\stackrel{\text{def.}}{=} \sum_z q(z) \log \left(\frac{P(x, z \mid \theta)}{q(z)} \right) \\
 &= \sum_z q(z) \log \left(\frac{P(z \mid x, \theta) \cdot P(x \mid \theta)}{q(z)} \right) \\
 &= \sum_z q(z) \left[\log \frac{P(z \mid x, \theta)}{q(z)} + \log P(x \mid \theta) \right] \\
 &= \sum_z q(z) \log \left(\frac{P(z \mid x, \theta)}{q(z)} \right) + \sum_z q(z) \log P(x \mid \theta) \\
 &= -KL(q(z) \parallel P(z \mid x, \theta)) + \log P(x \mid \theta) \cdot \underbrace{\sum_z q(z)}_{=1}
 \end{aligned}$$

$$\Rightarrow \quad \log P(x \mid \theta) = F(q(z), \theta) + KL(q(z) \parallel P(z \mid x, \theta)).$$

Observație (4)

Conform proprietății $KL(p \parallel q) \geq 0$ pentru $\forall p, q$, rezultă $KL(q(z) \parallel P(z \mid x, \theta)) \geq 0$.

Așadar, din egalitatea care tocmai a fost demonstrată la punctul b obținem (din nou!, după rezultatul de la punctul a) că $F(q(z), \theta)$ este o margine inferioară pentru log-verosimilitatea datelor observabile, $\ell(\theta) \stackrel{not.}{=} \log P(x \mid \theta)$.

c. Fie $\theta^{(t)}$ valoarea obținută pentru parametrul / parametrii θ la iterația t a algoritmului EM. Considerând această valoare fixată, arătați că maximul lui F în raport cu argumentul / distribuția $q(z)$ este atins pentru distribuția $P(z \mid x, \theta^{(t)})$, iar valoarea maximului este:

$$\max_{q(z)} F(q(z), \theta^{(t)}) = E_{P(z|x, \theta^{(t)})} [\log P(x, z \mid \theta^{(t)})] + H(P(z \mid x, \theta^{(t)}))$$

Soluție

Trebuie să maximizăm $F(q(z), \theta^{(t)})$ — marginea inferioară a log-verosimilității datelor observabile x — în raport cu distribuția $q(z)$.

Pe de o parte, rezultatul de la punctul a ne spune că $F(q(z), \theta) \leq \log P(x | \theta)$, pentru orice valoare a lui θ ; în particular, pentru $\theta^{(t)}$ avem

$$\log P(x | \theta^{(t)}) \geq F(q(z), \theta^{(t)})$$

Pe de altă parte, dacă în egalitatea demonstrată la punctul b se înlocuiește θ cu $\theta^{(t)}$, rezultă:

$$\log P(x | \theta^{(t)}) = F(q(z), \theta^{(t)}) + KL(q(z) || P(z | x, \theta^{(t)}))$$

În fine, dacă alegem $q(z) = P(z | x, \theta^{(t)})$, atunci termenul $KL(q(z) || P(z | x, \theta^{(t)}))$ din dreapta egalității de mai sus devine zero (vezi exercițiul [PS-31] de la capitolul de Probabilități și statistică).

Așadar, valoarea $\max_{q(z)} F(q(z), \theta^{(t)})$ se obține pentru distribuția $q(z) = P(z | x, \theta^{(t)})$.

Acum vom calcula această valoare maximă:

$$\begin{aligned}
 \max_{q(z)} F(q(z), \theta^{(t)}) &= \log P(x | \theta^{(t)}) \stackrel{\text{def.}}{=} \sum_z P(z | x, \theta^{(t)}) \log \left(\frac{P(x, z | \theta^{(t)})}{P(z | x, \theta^{(t)})} \right) \\
 &= E_{P(z|x, \theta^{(t)})} \left[\log \frac{P(x, z | \theta^{(t)})}{P(z | x, \theta^{(t)})} \right] \\
 &= E_{P(z|x, \theta^{(t)})} [\log P(x, z | \theta^{(t)}) - \log P(z | x, \theta^{(t)})] \\
 &= E_{P(z|x, \theta^{(t)})} [\log P(x, z | \theta^{(t)})] - E_{P(z|x, \theta^{(t)})} [\log P(z | x, \theta^{(t)})] \\
 &= E_{P(z|x, \theta^{(t)})} [\log P(x, z | \theta^{(t)})] + H[P(z | x, \theta^{(t)})] \\
 &= Q(\theta^{(t)} | \theta^{(t)}) + H[P(z | x, \theta^{(t)})]
 \end{aligned}$$

unde $Q(\theta | \theta^{(t)}) \stackrel{\text{not.}}{=} E_{P(z|x, \theta^{(t)})} [\log P(x, z | \theta)]$.

Observație (5)

Notând $G_t(\theta) \stackrel{\text{def.}}{=} F(P(z \mid x, \theta^{(t)}), \theta)$, din calculul de mai sus rezultă că $G_t(\theta^{(t)}) = \log P(x \mid \theta^{(t)}) = Q(\theta^{(t)} \mid \theta^{(t)}) + H(P(z \mid x, \theta^{(t)}))$. Se poate demonstra ușor — procedând similar cu calculul de mai sus — egalitatea

$$G_t(\theta) = Q(\theta \mid \theta^{(t)}) + H[P(z \mid x, \theta^{(t)})]$$

Observând că termenul $H[P(z \mid x, \theta^{(t)})]$ din această ultimă egalitate nu depinde de θ , rezultă imediat că

$$\operatorname{argmax}_{\theta} G_t(\theta) = \operatorname{argmax}_{\theta} Q(\theta \mid \theta^{(t)})$$

În consecință,

$$\theta^{(t+1)} \stackrel{\text{def.}}{=} \operatorname{argmax}_{\theta} F(P(z \mid x, \theta^{(t)}), \theta) = \operatorname{argmax}_{\theta} G_t(\theta) = \operatorname{argmax}_{\theta} Q(\theta \mid \theta^{(t)})$$

Egalitatea precedentă este responsabilă pentru următoarea reformulare (cea uzuală!) a algoritmului EM:

Pasul E': calculează $Q(\theta \mid \theta^{(t)}) = E_{P(z|x, \theta^{(t)})}[\log P(x, z \mid \theta)]$

Pasul M': calculează $\theta^{(t+1)} = \operatorname{argmax}_{\theta} Q(\theta \mid \theta^{(t)})$

Algoritmul EM: corectitudine / convergență

prelucrare de Liviu Ciortuz, după
en.wikipedia.org/wiki/Expectation-maximization

Pentru a maximiza funcția de log-verosimilitate a datelor observabile, i.e. $\ell(\theta) \stackrel{\text{def.}}{=} \log P(x \mid \theta)$, unde baza logaritmului (nespecificată) este considerată supraunitară, algoritmul EM procedează în mod iterativ, optimizând la pasul M al fiecărei iterații (t) o funcție “auxiliară”

$$Q(\theta \mid \theta^{(t)}) \stackrel{\text{def.}}{=} E_{P(z \mid x, \theta^{(t)})}[\log P(x, z \mid \theta)],$$

reprezentând media log-verosimilității datelor complete (observabile și neobservabile) în raport cu distribuția condițională $P(z \mid x, \theta^{(t)})$.

Vom considera iterațiile $t = 0, 1, \dots$ și $\theta^{(t+1)} = \operatorname{argmax}_{\theta} Q(\theta \mid \theta^{(t)})$, cu $\theta^{(0)}$ ales în mod arbitrar.

Demonstrați că pentru orice t fixat (arbitrar) și pentru orice θ astfel încât $Q(\theta \mid \theta^{(t)}) \geq Q(\theta^{(t)} \mid \theta^{(t)})$ are loc inegalitatea:

$$\log P(x \mid \theta) - \log P(x \mid \theta^{(t)}) \geq Q(\theta \mid \theta^{(t)}) - Q(\theta^{(t)} \mid \theta^{(t)}) \quad (10)$$

Observații

1. Semnificația imediată a relației (10):

Orice îmbunătățire a valorii funcției auxiliare $Q(\theta \mid \theta^{(t)})$ conduce la o îmbunătățire cel puțin la fel de mare a valorii funcției obiectiv, $\ell(\theta)$.


2. Dacă în inegalitatea (10) se înlocuiește θ cu $\theta^{(t+1)} = \operatorname{argmax}_{\theta} Q(\theta \mid \theta^{(t)})$, va rezulta

$$\log P(x \mid \theta^{(t+1)}) \geq \log P(x \mid \theta^{(t)}).$$

În final, vom avea

$$\ell(\theta^{(0)}) \leq \ell(\theta^{(t)}) \leq \ell(\theta^{(t+1)}) \leq \dots$$

Șirul acesta (monoton) este mărginit superior de 0 (vezi definiția lui ℓ), deci converge la o anumită valoare ℓ^* . În anumite cazuri / condiții, această valoare este un maxim (în general, local) al funcției de log-verosimilitate.



Observații (cont.)

3. Conform aceleiași inegalități (10), la pasul M de la iterația t a algoritmului EM, este suficient ca în loc să se ia $\theta^{(t+1)} = \operatorname{argmax}_{\theta} Q(\theta \mid \theta^{(t)})$, să se aleagă $\theta^{(t+1)}$ astfel încât $Q(\theta^{(t+1)} \mid \theta^{(t)}) > Q(\theta^{(t)} \mid \theta^{(t)})$. Aceasta constituie *versiunea “generalizată”* a algoritmului EM.

Demonstrația relației (10)

$$P(x, z \mid \theta) = P(z \mid x, \theta) \cdot P(x \mid \theta) \Rightarrow \log P(x \mid \theta) = \log P(x, z \mid \theta) - \log P(z \mid x, \theta) \Rightarrow$$

$$\underbrace{\sum_z P(z \mid x, \theta^{(t)}) \cdot \log P(x \mid \theta)}_1 =$$

$$\sum_z P(z \mid x, \theta^{(t)}) \cdot \log P(x, z \mid \theta) - \sum_z P(z \mid x, \theta^{(t)}) \cdot \log P(z \mid x, \theta) \Rightarrow$$

$$\log P(x \mid \theta) = Q(\theta \mid \theta^{(t)}) - \sum_z P(z \mid x, \theta^{(t)}) \cdot \log P(z \mid x, \theta)$$

Ultimul termen din egalitatea aceasta reprezintă o cross-entropie, pe care o vom nota cu $CH(\theta \mid \theta^{(t)})$. Așadar,

$$\log P(x \mid \theta) = Q(\theta \mid \theta^{(t)}) + CH(\theta \mid \theta^{(t)})$$

Această egalitate este valabilă pentru toate valorile posibile ale parametrului θ . În particular pentru $\theta = \theta^{(t)}$, vom avea:

$$\log P(x \mid \theta^{(t)}) = Q(\theta^{(t)} \mid \theta^{(t)}) + CH(\theta^{(t)} \mid \theta^{(t)})$$


Demonstrația relației (10), cont.

Scăzând membru cu membru ultimele două egalități, obținem:

$$\log P(x \mid \theta) - \log P(x \mid \theta^{(t)}) = Q(\theta \mid \theta^{(t)}) - Q(\theta^{(t)} \mid \theta^{(t)}) + CH(\theta \mid \theta^{(t)}) - CH(\theta^{(t)} \mid \theta^{(t)})$$

Conform inegalității lui Gibbs, avem $CH(\theta \mid \theta^{(t)}) \geq CH(\theta^{(t)} \mid \theta^{(t)})$, deci în final rezultă:


$$\log P(x \mid \theta) - \log P(x \mid \theta^{(t)}) \geq Q(\theta \mid \theta^{(t)}) - Q(\theta^{(t)} \mid \theta^{(t)})$$



From:

What is the expectation maximization algorithm?

Chuong B. Do, Serafim Batzoglou,
Nature Biotechnology,
vol. 26, no. 8, 2008, pp. 897-899



**The EM algorithm for solving
a Bernoulli mixture model**

CMU, 2008 fall, Eric Xing, HW4, pr. 1.4-7

Suppose I have two unfair coins. The first lands on heads with probability p , and the second lands on heads with probability q .

Imagine n tosses, where for each toss I choose to use the first coin with probability π and choose to use the second with probability $1 - \pi$. The outcome of each toss i is $x_i \in \{0, 1\}$.

Suppose I tell you the outcomes of the n tosses, $x \stackrel{\text{not.}}{=} \{x_1, x_2, \dots, x_n\}$, but I don't tell you which coins I used on which toss.

Given only the outcomes, x , your job is to compute estimates for θ which is the set of all parameters, $\theta = \{p, q, \pi\}$ using the EM algorithm.

To compute these estimates, we will create a latent variable Z , where $z_i \in \{0, 1\}$ indicates the coin used for the n^{th} toss. For example $z_2 = 1$ indicates the first coin was used on the second toss.

We define the “incomplete” data log-likelihood as $\log P(x|\theta)$ and the “complete” data log-likelihood as $\log P(x, z|\theta)$.

- a. Show that $E[z_i | x_i, \theta] = P(z_i = 1 | x_i, \theta)$.
- b. Use Bayes rule to compute $P(z_i = 1 | x_i, \theta^{(t)})$, where $\theta^{(t)}$ denotes the parameters at iteration t .
- c. Write down the complete log-likelihood, $\log P(x, z | \theta)$.
- d. ***E-Step:*** Show that the expected log-likelihood of the complete data $Q(\theta | \theta^{(t)}) \stackrel{not.}{=} E_{P(z|x, \theta^{(t)})}[\log P(x, z | \theta)]$ is given by

$$\begin{aligned}
 Q(\theta | \theta^{(t)}) &= \sum_{i=1}^n E[z_i | x_i, \theta^{(t)}] \cdot (\log \pi + x_i \log p + (1 - x_i) \log(1 - p)) + \\
 &\quad + (1 - E[z_i | x_i, \theta^{(t)}]) \cdot (\log(1 - \pi) + x_i \log q + (1 - x_i) \log(1 - q))
 \end{aligned}$$

- e. ***M-Step:*** Describe the process you would use to obtain the update equations for $p^{(t+1)}$, $q^{(t+1)}$ si $\pi^{(t+1)}$.

Solution

a.

$$\begin{aligned}
 E[z_i \mid x_i, \theta] &= \sum_{z \in \{0,1\}} z_i P(z_i \mid x_i, \theta) = 0 \cdot P(z_i = 0 \mid x_i, \theta) + 1 \cdot P(z_i = 1 \mid x_i, \theta) \\
 &\Rightarrow E[z_i \mid x_i, \theta] = P(z_i = 1 \mid x_i, \theta)
 \end{aligned}$$

b.

$$\begin{aligned}
 &P(z_i = 1 \mid x_i, \theta) \\
 &= \frac{P(x_i \mid z_i = 1, \theta)P(z_i = 1 \mid \theta)}{P(x_i \mid z_i = 1, \theta)P(z_i = 1 \mid \theta) + P(x_i \mid z_i = 0, \theta)P(z_i = 0 \mid \theta)} \\
 &= \frac{p^{x_i} \cdot (1 - p)^{1-x_i} \cdot \pi}{p^{x_i} \cdot (1 - p)^{1-x_i} \cdot \pi + q^{x_i} \cdot (1 - q)^{1-x_i} \cdot (1 - \pi)}
 \end{aligned}$$

c.

$$\begin{aligned}
 \log P(x, z \mid \theta) &\stackrel{i.i.d.}{=} \log \prod_{i=1}^n P(x_i, z_i \mid \theta) = \log \prod_{i=1}^n P(x_i \mid z_i, \theta) \cdot P(z_i \mid \theta) \\
 &= \log \prod_{i=1}^n \left(p^{x_i} (1-p)^{1-x_i} \pi \right)^{z_i} \left(q^{x_i} (1-q)^{1-x_i} (1-\pi) \right)^{1-z_i} \\
 &= \sum_{i=1}^n \log \left(\left(p^{x_i} (1-p)^{1-x_i} \pi \right)^{z_i} \left(q^{x_i} (1-q)^{1-x_i} (1-\pi) \right)^{1-z_i} \right) \\
 &= \sum_{i=1}^n \left[z_i \log \left(p^{x_i} (1-p)^{1-x_i} \pi \right) + (1-z_i) \log \left(q^{x_i} (1-q)^{1-x_i} (1-\pi) \right) \right]
 \end{aligned}$$

d.

$$\begin{aligned}
 Q(\theta \mid \theta^{(t)}) &\stackrel{\text{not.}}{=} E_{P(z|x, \theta^{(t)})} [\log P(x, z \mid \theta)] \\
 &= E_{P(z|x, \theta^{(t)})} \left[\sum_{i=1}^n \left[z_i \log (p^{x_i} (1-p)^{1-x_i} \pi) + \right. \right. \\
 &\quad \left. \left. (1-z_i) \log (q^{x_i} (1-q)^{1-x_i} (1-\pi)) \right] \right] \\
 &= \sum_{i=1}^n \left[E[z_i \mid x_i, \theta^{(t)}] \cdot \log p^{x_i} (1-p)^{1-x_i} \pi + \right. \\
 &\quad \left. + (1 - E[z_i \mid x_i, \theta^{(t)}]) \cdot \log q^{x_i} (1-q)^{1-x_i} (1-\pi) \right] \\
 &= \sum_{i=1}^n \left[E[z_i \mid x_i, \theta^{(t)}] \cdot (\log \pi + x_i \log p + (1-x_i) \log(1-p)) + \right. \\
 &\quad \left. + (1 - E[z_i \mid x_i, \theta^{(t)}]) \cdot (\log(1-\pi) + x_i \log q + (1-x_i) \log(1-q)) \right]
 \end{aligned}$$

e.

$$\mu_i^{(t)} \stackrel{not.}{=} E[z_i \mid x_i, \theta^{(t)}] = \frac{(p^{(t)})^{x_i} \cdot (1 - p^{(t)})^{1-x_i} \cdot \pi^{(t)}}{(p^{(t)})^{x_i} \cdot (1 - p^{(t)})^{1-x_i} \cdot \pi^{(t)} + (q^{(t)})^{x_i} \cdot (1 - q^{(t)})^{1-x_i} \cdot (1 - \pi^{(t)})}$$

$$\begin{aligned} Q(\theta \mid \theta^{(t)}) &= \sum_{i=1}^n \left[\mu_i^{(t)} (\log \pi + x_i \log p + (1 - x_i) \log(1 - p)) + \right. \\ &\quad \left. + (1 - \mu_i^{(t)}) \cdot (\log(1 - \pi) + x_i \log q + (1 - x_i) \log(1 - q)) \right] \end{aligned}$$

$$\begin{aligned} \frac{\partial Q(\theta \mid \theta^{(t)})}{\partial p} = 0 &\Leftrightarrow \sum_{i=1}^n \mu_i^{(t)} \left(\frac{x_i}{p} - \frac{1 - x_i}{1 - p} \right) = 0 \Leftrightarrow \frac{1}{p} \sum_{i=1}^n \mu_i^{(t)} x_i = \frac{1}{1 - p} \sum_{i=1}^n \mu_i^{(t)} (1 - x_i) \\ &\Leftrightarrow (1 - p) \sum_{i=1}^n \mu_i^{(t)} x_i = p \sum_{i=1}^n \mu_i^{(t)} (1 - x_i) \Leftrightarrow \sum_{i=1}^n \mu_i^{(t)} x_i = p \left(\sum_{i=1}^n \mu_i^{(t)} (1 - x_i) + \sum_{i=1}^n \mu_i^{(t)} x_i \right) \\ &\Leftrightarrow \sum_{i=1}^n \mu_i^{(t)} x_i = p \sum_{i=1}^n \mu_i^{(t)} \Rightarrow p^{(t+1)} = \frac{\sum_{i=1}^n \mu_i^{(t)} x_i}{\sum_{i=1}^n \mu_i^{(t)}} \end{aligned}$$

$$\frac{\partial Q(\theta \mid \theta^{(t)})}{\partial q} = 0 \Leftrightarrow \sum_{i=1}^n (1 - \mu_i^{(t)}) \left(\frac{x_i}{q} - \frac{1 - x_i}{1 - q} \right) = 0 \quad \Rightarrow \quad q^{(t+1)} = \frac{\sum_{i=1}^n (1 - \mu_i^{(t)}) x_i}{\sum_{i=1}^n (1 - \mu_i^{(t)})}$$

$$\frac{\partial Q(\theta \mid \theta^{(t)})}{\partial \pi} = 0 \Leftrightarrow \sum_{i=1}^n \left(\frac{\mu_i^{(t)}}{\pi} - \frac{1 - \mu_i^{(t)}}{1 - \pi} \right) = 0 \Leftrightarrow \frac{1}{\pi} \sum_{i=1}^n \mu_i^{(t)} = \frac{1}{1 - \pi} \sum_{i=1}^n (1 - \mu_i^{(t)})$$

$$\Leftrightarrow (1 - \pi) \sum_{i=1}^n \mu_i^{(t)} = \pi \sum_{i=1}^n (1 - \mu_i^{(t)}) \Leftrightarrow \sum_{i=1}^n \mu_i^{(t)} = \pi \left(\sum_{i=1}^n (1 - \mu_i^{(t)}) + \sum_{i=1}^n \mu_i^{(t)} \right)$$

$$\Leftrightarrow \sum_{i=1}^n \mu_i^{(t)} = \pi \sum_{i=1}^n 1 \Leftrightarrow \sum_{i=1}^n \mu_i^{(t)} = n\pi \quad \Rightarrow \quad \pi^{(t+1)} = \frac{1}{n} \sum_{i=1}^n \mu_i^{(t)}$$

Estimarea parametrilor unei mixturi de distribuții binomiale

- A. când toate variabilele sunt observabile: MLE;
- B. când unele variabile sunt neobservabile: algoritmul EM

prelucrare de Liviu Ciortuz, după

“What is the expectation maximization algorithm?”,
Chuong B. Do, Serafim Batzoglou,
Nature Biotechnology, vol. 26, no. 8, 2008, pag. 897-899

Fie următorul *experiment probabilist*:

Disponem de două monede, A și B .

Efectuăm 5 serii de operațiuni de tipul următor:

- Alegem în mod aleatoriu una dintre monedele A și B , cu probabilitate egală ($1/2$);
- Aruncăm de 10 ori moneda care tocmai a fost aleasă (Z) și notăm rezultatul (X), ca număr de fețe ‘head’ (ro. ‘stemă’) obținute în urma aruncării.

A. La acest punct vom considera că s-a obținut următorul rezultat pentru experimentul nostru:

i	Z_i	X_i
1	B	$5H \ (5T)$
2	A	$9H \ (1T)$
3	A	$8H \ (2T)$
4	B	$4H \ (6T)$
5	A	$7H \ (3T)$

Semnificația variabilelor aleatoare Z_i și X_i pentru $i = 1, \dots, 5$ din tabelul de mai sus este imediată.

i. Calculați $\hat{\theta}_A$ și $\hat{\theta}_B$, probabilitățile de apariție a feței ‘head’ pentru cele două monede, folosind *definiția clasică a probabilităților*, și anume raportul dintre numărul de cazuri favorabile și numărul de cazuri posibile, relativ la întregul experiment.

Răspuns:

Analizând datele din tabelul din enunț, rezultă imediat

$$\hat{\theta}_A = \frac{24}{24 + 6} = 0.8 \text{ și } \hat{\theta}_B = \frac{9}{9 + 11} = 0.45.$$

De observat că termenii 6 și respectiv 11 de la numitorii acestor fracții reprezintă numărul de fețe ‘tail’ care au fost obținute la aruncarea monedei *A* și respectiv *B*: $6T = 1T + 2T + 3T$, $11T = 5T + 6T$.

Observație

Dacă în locul variabilelor binare $Z_i \in \{A, B\}$ pentru $i = 1, \dots, 5$ introducem în mod natural variabilele-indicator $Z_{i,A} \in \{0, 1\}$ și $Z_{i,B} \in \{0, 1\}$ tot pentru $i = 1, \dots, 5$, definite prin $Z_{i,A} = 1$ iff $Z_i = A$, și $Z_{i,A} = 0$ iff $Z_{i,B} = 0$, atunci procesările necesare pentru calculul probabilităților/parametrilor $\hat{\theta}_A$ și $\hat{\theta}_B$ pot fi prezentate în mod sintetizat ca în tabelul de mai jos.^a

i	$Z_{i,A}$	$Z_{i,B}$	X_i
1	0	1	5H
2	1	0	9H
3	1	0	8H
4	0	1	4H
5	1	0	7H

\Rightarrow

^aPrezentăm acest “artificiu” ca pregătire pentru rezolvarea (ulterioară a) punctului B al prezentei probleme.

$$\Rightarrow$$

$X_i \cdot Z_{i,A}$	$X_i \cdot Z_{i,B}$
$0H \ (0T)$	$5H \ (5T)$
$9H \ (1T)$	$0H \ (0T)$
$8H \ (2T)$	$0H \ (0T)$
$0H \ (0T)$	$4H \ (6T)$
$7H \ (3T)$	$0H \ (0T)$
$\sum_{i=1}^5 X_i \cdot Z_{i,A} = 24H$	$\sum_{i=1}^5 X_i \cdot Z_{i,B} = 9H$
$\sum_{i=1}^5 (10 - X_i) \cdot Z_{i,A} = 6T$	$\sum_{i=1}^5 (10 - X_i) \cdot Z_{i,B} = 11T$

$$\Rightarrow$$

$$\Rightarrow \begin{cases} \hat{\theta}_A = \frac{24}{24 + 6} = 0.8 \\ \hat{\theta}_B = \frac{9}{9 + 11} = 0.45 \end{cases}$$

ii. Calculați $l_1(\theta_A, \theta_B) \stackrel{not.}{=} P(X, Z \mid \theta_A, \theta_B)$, funcția de verosimilitate a datelor “complete” — unde $X \stackrel{not.}{=} \langle X_1, \dots, X_5 \rangle$ sunt datele “observabile”, iar $Z \stackrel{not.}{=} \langle Z_1, \dots, Z_5 \rangle$ sunt datele “neobservabile” —, în raport cu parametrii θ_A și respectiv θ_B ai distribuțiilor binomiale care modelează aruncarea celor două monede.

Răspuns:

Calculul verosimilității datelor complete:

$$\begin{aligned}
 l_1(\theta_A, \theta_B) &\stackrel{\text{def.}}{=} P(X, Z_A, Z_B \mid \theta_A, \theta_B) \stackrel{\text{indep. cdt.}}{=} \prod_{i=1}^5 P(X_i, Z_{i,A}, Z_{i,B} \mid \theta_A, \theta_B) \\
 &= \prod_{i=1}^5 P(X_i \mid Z_{i,A}, Z_{i,B}; \theta_A, \theta_B) \cdot P(Z_{i,A}, Z_{i,B} \mid \theta_A, \theta_B) \\
 &= P(X_1 \mid Z_{B,1} = 1, \theta_B) \cdot 1/2 \cdot \\
 &\quad P(X_2 \mid Z_{A,2} = 1, \theta_A) \cdot 1/2 \cdot \\
 &\quad P(X_3 \mid Z_{A,3} = 1, \theta_A) \cdot 1/2 \cdot \\
 &\quad P(X_4 \mid Z_{B,4} = 1, \theta_B) \cdot 1/2 \cdot \\
 &\quad P(X_5 \mid Z_{A,5} = 1, \theta_A) \cdot 1/2 \\
 &= \theta_B^5 (1 - \theta_B)^5 \cdot \theta_A^9 (1 - \theta_A) \cdot \theta_A^8 (1 - \theta_A)^2 \cdot \theta_B^4 (1 - \theta_B)^6 \cdot \theta_A^7 (1 - \theta_A)^3 \cdot \frac{1}{2^5} \\
 &= \frac{1}{2^5} \theta_A^{24} (1 - \theta_A)^6 \theta_B^9 (1 - \theta_B)^{11}
 \end{aligned}$$

iii. Calculați $\hat{\theta}_A \stackrel{not.}{=} \arg \max_{\theta_A} \log l_1(\theta_A, \theta_B)$ și $\hat{\theta}_B \stackrel{not.}{=} \arg \max_{\theta_B} \log l_1(\theta_A, \theta_B)$ folosind derivatele parțiale de ordinul întâi.

Observații:

1. Baza logaritmului, fixată dar lăsată mai sus nespecificată, se va considera supraunitară (de exemplu 2, e sau 10).
2. Lucrând corect, veți obține același rezultat ca la punctul i .

Răspuns:

Funcția de log-verosimilitate a datelor complete se exprimă astfel:

$$\log l_1(\theta_A, \theta_B) = -5 \log 2 + 24 \log \theta_A + 6 \log(1 - \theta_A) + 9 \log \theta_B + 11 \log(1 - \theta_B)$$

Prin urmare, maximul acestei funcții în raport cu parametrul θ_A se calculează astfel:

$$\frac{\partial \log l_1(\theta_A, \theta_B)}{\partial \theta_A} = 0 \Leftrightarrow \frac{24}{\theta_A} - \frac{6}{1 - \theta_A} = 0 \Leftrightarrow \frac{4}{\theta_A} = \frac{1}{1 - \theta_A} \Leftrightarrow 4 - 4\theta_A = \theta_A \Leftrightarrow \hat{\theta}_A = 0.8$$

Similar, se face calculul și pentru $\frac{\partial \log l_1(\theta_A, \theta_B)}{\partial \theta_B}$ și se obține $\hat{\theta}_B = 0.45$.^a

Cele două valori obținute, $\hat{\theta}_A$ și $\hat{\theta}_B$, reprezintă estimarea de verosimilitate maximă (MLE) a probabilităților de apariție a feței ‘head’ (‘stema’) pentru moneda A și respectiv moneda B .

^aSe verifică ușor faptul că într-adevăr rădăcinile derivatelor parțiale de ordinul întâi pentru funcția de log-verosimilitate reprezintă puncte de maxim. Pentru aceasta se studiază semnele acestor derivate.

Observații

1. Am arătat pe acest caz particular că metoda de calculare a probabilităților ($\hat{\theta}_A$ și $\hat{\theta}_B$) direct din datele observate (așa cum o știm din liceu) corespunde de fapt metodei de estimare în sensul verosimilității maxime (MLE).
2. La punctul B vom arăta cum anume se poate face estimarea aceluiași parametri θ_A și θ_B în cazul în care o parte din variabile, și anume $Z_{i,A}$ și $Z_{i,B}$, sunt neobservabile.

B. La acest punct se va relua experimentul de la punctul A, însă de data aceasta vom considera că valorile variabilelor Z_i nu sunt cunoscute.

i	Z_i	X_i
1	?	$5H \ (5T)$
2	?	$9H \ (1T)$
3	?	$8H \ (2T)$
4	?	$4H \ (6T)$
5	?	$7H \ (3T)$

iv. Pentru conveniență, pentru $i = 1, \dots, 5$ vom considera variabilele-indicator “neobservabile” $Z_{i,A}, Z_{i,B} \in \{0, 1\}$, cu $Z_{i,A} = 1$ iff $Z_{i,B} = 0$ și $Z_{i,B} = 1$ iff $Z_{i,A} = 0$.

Folosind teorema lui Bayes, calculați mediile variabilelor neobservabile $Z_{i,A}$ și $Z_{i,B}$ condiționate de variabilele observabile X_i . Veți considera că parametrii acestor distribuții binomiale care modelează aruncarea monedelor A și B au valorile $\theta_A^{(0)} = 0.6$ și respectiv $\theta_B^{(0)} = 0.5$.

Așadar, se cer: $E[Z_{i,A} \mid X_i, \theta_A^{(0)}, \theta_B^{(0)}]$ și $E[Z_{i,B} \mid X_i, \theta_A^{(0)}, \theta_B^{(0)}]$ pentru $i = 1, \dots, 5$. Ca și mai înainte, probabilitățile a priori $P(Z_{i,A} = 1)$ și $P(Z_{i,B} = 1)$ se vor considera $1/2$.

Răspuns:

Algoritmul EM ne permite să facem în mod iterativ estimarea parametrilor θ_A și θ_B în funcție de valorile variabilelor observabile, X_i , și de valorile inițiale atribuite parametrilor (în cazul nostru, $\theta_A^{(0)} = 0.6$ și $\theta_B^{(0)} = 0.5$).

Vom face o sinteză a calculelor de la prima iterație a algoritmului EM — detaliate la punctele *iv*, *v* și *vi* de mai jos — sub forma următoare, care seamănă într-o anumită măsură cu tabelele de la punctul A:

i	$Z_{i,A}$	$Z_{i,B}$	X_i		$E[Z_{i,A}]$	$E[Z_{i,B}]$	
1	—	—	$5H$		$0.45H$	$0.55H$	
2	—	—	$9H$	\xRightarrow{E}	$0.80H$	$0.20H$	\xRightarrow{M}
3	—	—	$8H$		$0.73H$	$0.27H$	
4	—	—	$4H$		$0.35H$	$0.65H$	
5	—	—	$7H$		$0.65H$	$0.35H$	

$$\begin{array}{ccc}
& \begin{array}{c|c} X_i \cdot E[Z_{i,A}] & X_i \cdot E[Z_{i,B}] \\ \hline 2.2H \ (2.2T) & 2.8H \ (2.8T) \\ 7.2H \ (0.8T) & 1.8H \ (0.2T) \\ 5.9H \ (1.5T) & 2.1H \ (0.5T) \\ 1.4H \ (2.1T) & 2.6H \ (3.9T) \\ 4.5H \ (1.9T) & 2.5H \ (1.1T) \\ \hline \sum_{i=1}^5 X_i \cdot E[Z_{i,A}] = 21.3H & \sum_{i=1}^5 X_i \cdot E[Z_{i,B}] = 11.7H \\ \sum_{i=1}^5 (10 - X_i) \cdot E[Z_{i,A}] = 8.3T & \sum_{i=1}^5 (10 - X_i) \cdot E[Z_{i,B}] = 8.3T \end{array} & \Rightarrow
\end{array}$$

$$\Rightarrow \begin{cases} \hat{\theta}_A^{(1)} = \frac{21.3}{21.3 + 8.7} \approx 0.71 \\ \hat{\theta}_B^{(1)} = \frac{11.7}{11.7 + 8.3} \approx 0.58 \end{cases}$$

Observații:

- este posibilă calcularea mediilor variabilelor neobservabile $Z_{i,A}$ și $Z_{i,B}$, condiționate de variabilele observabile X_i și în funcție de valorile asig-nate inițial (0.6 și 0.5 în enunț, dar în general ele se pot asigna în mod aleatoriu) pentru parametrii θ_A și θ_B ;
- față de tabloul de sinteză de la punctul precedent, când toate vari-abilele erau observabile și se calculau produsele $X_i \cdot Z_{i,A}$ și $X_i \cdot Z_{i,B}$, aici se înlocuiesc variabilele $Z_{i,A}$ și $Z_{i,B}$ cu mediile $E[Z_{i,A}]$ și $E[Z_{i,B}]$ în pro-dusele respective. De fapt, în loc să se calculeze $\sum_i X_i \cdot Z_{i,A}$ se calculează media $E[\sum_i X_i \cdot Z_{i,A}]$, și similar pentru B .

Justificarea concretă pentru cele două observații de mai sus va fi dezvoltată la punctele mai jos.

Notăție: Pentru simplitate, în cele de mai sus (inclusiv în tabelele prece-dente), prin $E[Z_{i,A}]$ am notat $E[Z_{i,A} \mid X_i, \theta^{(0)}]$, iar prin $E[Z_{i,B}]$ am notat $E[Z_{i,B} \mid X_i, \theta^{(0)}]$, unde $\theta^{(0)} \stackrel{not.}{=} (\theta_A^{(0)}, \theta_B^{(0)})$.

Întrucât variabilele $Z_{i,A}$ au valori boolene (0 sau 1), rezultă că

$$\begin{aligned} E[Z_{i,A} | X_i, \theta^{(0)}] &= 0 \cdot P(Z_{i,A} = 0 | X_i, \theta^{(0)}) + 1 \cdot P(Z_{i,A} = 1 | X_i, \theta^{(0)}) \\ &= P(Z_{i,A} = 1 | X_i, \theta^{(0)}) \end{aligned}$$

Probabilitățile $P(Z_{i,A} = 1 | X, \theta^{(0)}) = P(Z_{i,A} = 1 | X_i, \theta^{(0)})$, pentru $i = 1, \dots, 5$, se pot calcula folosind teorema lui Bayes:

$$\begin{aligned} P(Z_{i,A} = 1 | X_i, \theta^{(0)}) &= \frac{P(X_i | Z_{i,A} = 1, \theta^{(0)}) \cdot P(Z_{i,A} = 1 | \theta^{(0)})}{\sum_{j \in \{0,1\}} P(X_i | Z_{i,A} = j, \theta^{(0)}) \cdot P(Z_{i,A} = j | \theta^{(0)})} \\ &= \frac{P(X_i | Z_{i,A} = 1, \theta_A^{(0)})}{P(X_i | Z_{i,A} = 1, \theta_A^{(0)}) + P(X_i | Z_{i,B} = 1, \theta_B^{(0)})} \end{aligned}$$

S-a ținut cont că $P(Z_{i,A} = 1 | \theta^{(0)}) = P(Z_{i,B} = 1 | \theta^{(0)}) = 1/2$ (a se vedea enunțul).

De exemplu, pentru $i = 1$ vom avea:

$$E[Z_{A,1} \mid X_1, \theta^{(0)}] = \frac{0.6^5(1 - 0.6)^5}{0.6^5(1 - 0.6)^5 + 0.5^5(1 - 0.5)^5} = \frac{1}{1 + \left(\frac{0.25}{0.24}\right)^5} \approx 0.45$$

Similar cu $E[Z_{A,1} \mid X_1, \theta^{(0)}]$ se calculează și celelalte medii $E[Z_{i,A} \mid X_i, \theta^{(0)}]$ pentru $i = 2, \dots, 5$ și $E[Z_{i,B} \mid X_i, \theta^{(0)}]$ pentru $i = 1, \dots, 5$.

Observație: Se poate ține cont că, de îndată ce s-a calculat $E[Z_{i,A} \mid X_i, \theta^{(0)}]$, se poate obține imediat și $E[Z_{i,B} \mid X_i, \theta^{(0)}] = 1 - E[Z_{i,A} \mid X_i, \theta^{(0)}]$, fiindcă $Z_{i,A} + Z_{i,B} = 1$.

v. Calculați media funcției de log-verosimilitate a datelor complete, X (observabile) și Z (neobservabile):

$$l_2(\theta_A, \theta_B) \stackrel{def.}{=} E_{P(Z|X, \theta^{(0)})}[\log P(X, Z | \theta)],$$

unde $\theta \stackrel{not.}{=} (\theta_A, \theta_B)$ și $\theta^{(0)} \stackrel{not.}{=} (\theta_A^{(0)}, \theta_B^{(0)})$.

Semnificația notației de mai sus este următoarea:

Funcția $l_2(\theta_A, \theta_B)$ este o medie a variabilei aleatoare reprezentată de log-verosimilitatea datelor complete (observabile și, respectiv, neobservabile), iar această medie se calculează în raport cu distribuția probabilistă condițională a datelor neobservabile, $P(Z | X, \theta^{(0)})$.

Observație: La elaborarea calculului, veți folosi mai întâi proprietatea de liniaritate a mediilor variabilelor aleatoare, și apoi rezultatele de la punctul *iv*.

Răspuns:

Media funcției de log-verosimilitate a datelor complete, $l_2(\theta_A, \theta_B)$, se calculează astfel:

$$\begin{aligned}
 l_2(\theta_A, \theta_B) &\stackrel{def.}{=} E_{P(Z|X, \theta^{(0)})} [\log P(X, Z \mid \theta)] \\
 &\stackrel{indep. \text{ cdt.}}{=} E_{P(Z|X, \theta^{(0)})} \left[\log \prod_{i=1}^5 P(X_i, Z_{i,A}, Z_{i,B} \mid \theta_A, \theta_B) \right] \\
 &= E_{P(Z|X, \theta^{(0)})} \left[\log \prod_{i=1}^5 P(X_i \mid Z_{i,A}, Z_{i,B}; \theta_A, \theta_B) \cdot P(Z_{i,A}, Z_{i,B} \mid \theta_A, \theta_B) \right]
 \end{aligned}$$

În continuare, omițând din nou distribuția probabilistă în raport cu care se calculează media aceasta întrucât ea poate fi subînțeleasă, vom scrie:

$$\begin{aligned}
& l_2(\theta_A, \theta_B) \\
&= E\left[\log \prod_{i=1}^5 \cdot (\theta_A^{Z_{i,A}})^{X_i} \cdot [(1 - \theta_A)^{Z_{i,A}}]^{10-X_i} \cdot (\theta_B^{Z_{i,B}})^{X_i} \cdot [(1 - \theta_B)^{Z_{i,B}}]^{10-X_i} \cdot \frac{1}{2} \right] \\
&= E\left[\sum_{i=1}^5 [X_i \cdot Z_{i,A} \cdot \log \theta_A + (10 - X_i) \cdot Z_{i,A} \cdot \log(1 - \theta_A) + \right. \\
&\quad \left. X_i \cdot Z_{i,B} \cdot \log \theta_B + (10 - X_i) \cdot Z_{i,B} \cdot \log(1 - \theta_B) - \log 2] \right] \\
&= \sum_{i=1}^5 [X_i \cdot E[Z_{i,A}] \cdot \log \theta_A + (10 - X_i) \cdot E[Z_{i,A}] \cdot \log(1 - \theta_A) + \\
&\quad X_i \cdot E[Z_{i,B}] \cdot \log \theta_B + (10 - X_i) \cdot E[Z_{i,B}] \cdot \log(1 - \theta_B) - \log 2] \\
&= \sum_{i=1}^5 \log[\theta_A^{X_i \cdot E[Z_{i,A}]} \cdot (1 - \theta_A)^{(10-X_i) \cdot E[Z_{i,A}]} \cdot \theta_B^{X_i \cdot E[Z_{i,B}]} \cdot (1 - \theta_B)^{(10-X_i) \cdot E[Z_{i,B}]} \cdot \frac{1}{2}] \\
&= \log(\theta_A^{2.2} (1 - \theta_A)^{2.2} \theta_B^{2.8} (1 - \theta_B)^{2.8} \cdot \dots \cdot \theta_A^{4.5} (1 - \theta_A)^{1.9} \theta_B^{2.5} (1 - \theta_B)^{1.1} \cdot \frac{1}{2}).
\end{aligned}$$

La ultima egalitate de mai sus, cantitățile fracționare provin din calculele simple $X_1 \cdot E[Z_{1,A} \mid X_1, \theta] \approx 2.2$, $X_1 \cdot E[Z_{1,B} \mid X_1, \theta] \approx 2.8$, ..., $X_5 \cdot E[Z_{1,A} \mid X_5, \theta] \approx 4.5$, $X_5 \cdot E[Z_{1,B} \mid X_5, \theta] \approx 2.5$ (a se vedea tabelele precedente).

vi. Calculați $\theta_A^{(1)} \stackrel{not.}{=} \arg \max_{\theta_A} l_2(\theta_A, \theta_B)$ și $\theta_B^{(1)} \stackrel{not.}{=} \arg \max_{\theta_B} l_2(\theta_A, \theta_B)$.

Răspuns:

Valorile parametrilor θ_A și θ_B pentru care se atinge maximul mediei funcției de log-verosimilitate a datelor complete se obțin cu ajutorul derivatelor parțiale de ordinul întâi:

$$\begin{aligned} \frac{\partial l_2(\theta_A, \theta_B)}{\partial \theta_A} &= 0 \\ \Rightarrow \frac{\partial}{\partial \theta_A} (2.2 \log \theta_A + 2.2 \log(1 - \theta_A) + \dots + 4.5 \log \theta_A + 1.9 \log(1 - \theta_A)) &= 0 \\ \Rightarrow \frac{2.2}{\theta_A} - \frac{2.2}{1 - \theta_A} + \dots + \frac{4.5}{\theta_A} - \frac{1.9}{1 - \theta_A} &= 0 \Rightarrow \dots \Rightarrow \theta_A^{(1)} \approx 0.71. \end{aligned}$$

Similar, vom obține $\theta_B^{(1)} \approx 0.58$.

C. Formalizați pașii E și M ai algoritmului EM pentru estimarea parametrilor θ_A și θ_B în condițiile de la punctul B.

Răspuns:

Formulele care se folosesc în cadrul algoritmului EM pentru rezolvarea problemei date — i.e. estimarea parametrilor θ_A și θ_B atunci când variabilele Z_i sunt neobservabile —, se elaborează/deduc astfel:

Pasul E:

$$\begin{aligned}
 E[Z_{i,A} \mid X, \theta] &= P(Z_{i,A} = 1 \mid X, \theta) = P(Z_{i,A} = 1 \mid X_i, \theta) \\
 &\stackrel{T. B.}{=} \frac{P(X_i \mid Z_{i,A} = 1; \theta) \cdot \overbrace{P(Z_{i,A} = 1 \mid \theta)}^{1/2}}{P(X_i \mid Z_{i,A} = 1; \theta) \cdot P(Z_{i,A} = 1 \mid \theta) + P(X_i \mid Z_{i,B} = 1; \theta) \cdot \underbrace{P(Z_{i,B} = 1 \mid \theta)}_{1/2}} \\
 &= \frac{P(X_i \mid Z_{i,A} = 1; \theta)}{P(X_i \mid Z_{i,A} = 1; \theta) + P(X_i \mid Z_{i,B} = 1; \theta)} \\
 &= \frac{\theta_A^{X_i} (1 - \theta_A)^{10-X_i}}{\theta_A^{X_i} (1 - \theta_A)^{10-X_i} + \theta_B^{X_i} (1 - \theta_B)^{10-X_i}}
 \end{aligned}$$

Similar, vom obține:

$$E[Z_{i,B} \mid X, \theta] = \frac{\theta_B^{X_i} (1 - \theta_B)^{10-X_i}}{\theta_A^{X_i} (1 - \theta_A)^{10-X_i} + \theta_B^{X_i} (1 - \theta_B)^{10-X_i}}$$

Notând cu

- x_i valoarea variabilei X_i ,
- $\theta_A^{(t)}$ și respectiv $\theta_B^{(t)}$ estimările parametrilor θ_A și θ_B la iterația t a algoritmului EM,
- $\mu_{i,A}^{(t+1)}$ și respectiv $\mu_{i,B}^{(t+1)}$, mediile $E[Z_{i,A} \mid X_i, \theta_A^{(t)}]$ și $E[Z_{i,B} \mid X_i, \theta_B^{(t)}]$,

vom avea:

$$\mu_{i,A}^{(t+1)} = \frac{(\theta_A^{(t)})^{x_i} (10 - \theta_A^{(t)})^{10-x_i}}{(\theta_A^{(t)})^{x_i} (10 - \theta_A^{(t)})^{10-x_i} + (\theta_B^{(t)})^{x_i} (10 - \theta_B^{(t)})^{10-x_i}}$$

$$\mu_{i,B}^{(t+1)} = \frac{(\theta_B^{(t)})^{x_i} (10 - \theta_B^{(t)})^{10-x_i}}{(\theta_A^{(t)})^{x_i} (10 - \theta_A^{(t)})^{10-x_i} + (\theta_B^{(t)})^{x_i} (10 - \theta_B^{(t)})^{10-x_i}}$$

Pasul M:

Ca și mai înainte, în formulele de mai jos vom folosi notațiile simplificate $E[Z_{i,A}] \stackrel{not.}{=} E[Z_{i,A} | X_i, \theta^{(t)}]$ și $E[Z_{i,B}] \stackrel{not.}{=} E[Z_{i,B} | X_i, \theta^{(t)}]$.

Cu aceste notații, procedând similar cu calculul de la partea B, punctul v , vom avea:

$$l_2(\theta_A, \theta_B) = \log \prod_{i=1}^5 \theta_A^{x_i E[Z_{i,A}]} (1 - \theta_A)^{(10-x_i) E[Z_{i,A}]} \theta_B^{x_i E[Z_{i,B}]} (1 - \theta_B)^{(10-x_i) E[Z_{i,B}]}$$

Prin urmare,

$$\begin{aligned}
 \frac{\partial}{\partial \theta_A} l_2(\theta_A, \theta_B) = 0 &\Rightarrow \frac{1}{\theta_A} \sum_{i=1}^5 x_i E[Z_{i,A}] = \frac{1}{1 - \theta_A} \sum_{i=1}^5 (10 - x_i) E[Z_{i,A}] \\
 &\Rightarrow (1 - \theta_A) \sum_{i=1}^5 x_i E[Z_{i,A}] = \theta_A \sum_{i=1}^5 (10 - x_i) E[Z_{i,A}] \\
 &\Rightarrow \sum_{i=1}^5 x_i E[Z_{i,A}] = 10 \theta_A \sum_{i=1}^5 E[Z_{i,A}] \\
 &\Rightarrow \theta_A = \frac{\sum_{i=1}^5 x_i E[Z_{i,A}]}{10 \sum_{i=1}^5 E[Z_{i,A}]} \quad \text{și, similar, } \theta_B = \frac{\sum_{i=1}^5 x_i E[Z_{i,B}]}{10 \sum_{i=1}^5 E[Z_{i,B}]}
 \end{aligned}$$

Așadar, la pasul M al algoritmului EM vom avea:

$$\theta_A^{(t+1)} = \frac{\sum_{i=1}^5 x_i \mu_{i,A}^{(t+1)}}{10 \sum_{i=1}^5 \mu_{i,A}^{(t+1)}} \quad \text{și} \quad \theta_B^{(t+1)} = \frac{\sum_{i=1}^5 x_i \mu_{i,B}^{(t+1)}}{10 \sum_{i=1}^5 \mu_{i,B}^{(t+1)}}$$

Observație

Implementând algoritmul EM cu relațiile obținute pentru pasul E și pasul M, după execuția a 10 iterații se vor obține valorile $\theta_A^{(10)} \approx 0.80$ și $\theta_B^{(10)} \approx 0.52$.

Este interesant de observat că estimarea obținută pentru parametrul θ_A este acum la același nivel cu cea obținută prin metoda verosimilității maxime (MLE, vezi punctul A) în cazul observării tuturor variabilelor (0.80, vezi rezolvarea de la partea A, punctul i), iar estimarea obținută pentru parametrul θ_B a coborât de la valoarea 0.58 care a fost obținută la prima iterație a algoritmului EM la o valoare (0.52) care este considerabil mai apropiată de estimarea prin metoda MLE (0.45).

Maximum likelihood




Coin A	Coin B
	5 H, 5 T
9 H, 1 T	
8 H, 2 T	
	4 H, 6 T
7 H, 3 T	
24 H, 6 T	9 H, 11 T

$$\hat{\theta}_A = \frac{24}{24 + 6} = 0.80$$

$$\hat{\theta}_B = \frac{9}{9 + 11} = 0.45$$

Expectation maximization



The EM algorithm for solving
a mixture of K categorical distributions,
applied to the problem of Word Sense Disambiguation,
i.e. identifying the semantic domains associated to words in a
text document

CMU, 2012 fall, Eric Xing, Aarti Singh, HW3, pr. 3

The objective of this exercise is to derive the update equations of the EM algorithm for optimizing the latent variables [designating the semantic domains] involved in generating a text document.

Each *word* will be seen as a random variable w that can take values $1, \dots, V$ from the *vocabulary* of words. In fact, we will denote each w by an array of V components such that $w(i) = 1$ if w takes the value of the i -th word in the vocabulary. Hence, $\sum_{i=1}^V w(i) = 1$.

Given a *document* containing words w_j , $j = 1, \dots, N$, where N is the length of the document, we will assume that these words are generated from a mixture of K discrete *topics*:

$$P(w) = \sum_{m=1}^K \pi_m P(w|\mu_m) \text{ and } P(w|\mu_m) = \prod_{i=1}^V \mu_m(i)^{w(i)},$$

where

π_m denotes the prior [probability] for the latent topic variable $t = m$,

$\mu_k \stackrel{\text{not.}}{=} (\mu_k(1), \dots, \mu_k(i), \dots, \mu_k(V))$, with $\mu_k(i) \geq 0$ for $i = 1, \dots, V$ and $\sum_{i=1}^V \mu_m(i) = 1$ for each $k = 1, \dots, K$, and

$\mu_m(i) \stackrel{\text{not.}}{=} P(w(i) = 1 | t = m)$.

a. In the expectation step, for each word w_j , compute

$$F_j(t) \stackrel{not.}{=} P(t|w_j; \theta),$$

the probability that w_j belongs to each of the K topics, where θ is the set of parameters of this mixture model.

Answer:

$$\begin{aligned} F_j(t_j = m) &\stackrel{not.}{=} P(t_j = m|w_j; \theta) \stackrel{Bayes T.}{=} \frac{P(w_j|t_j = m; \theta) P(t_j = m|\theta)}{P(w_j|\theta)} \\ &= \frac{\pi_m P(w_j|\mu_m)}{\sum_{m'=1}^K \pi_{m'} P(w_j|\mu_{m'})} = \frac{\pi_m \prod_{l=1}^V \mu_m(l)^{w_j(l)}}{\sum_{m'=1}^K \pi_{m'} \prod_{l=1}^V \mu_{m'}(l)^{w_j(l)}} \end{aligned}$$

b. In the maximization step, compute θ that maximizes the log-likelihood of the data

$$l(w; \theta) \stackrel{not.}{=} \log \prod_{j=1}^N P(w_j; \theta)$$

Hint: Summing over the latent topic variable, we can write $l(w; \theta)$ as

$$l(w; \theta) = \sum_{j=1}^N \log \sum_t P(w_j, t; \theta) = \sum_{j=1}^N \log F_j(t) \frac{\sum_t P(w_j, t; \theta)}{F_j(t)}$$

Further on, using Jensen's inequality (see pr. EM-2) we get:

$$l(w; \theta) \geq \sum_{j=1}^N \sum_t F_j(t) \log \frac{P(w_j, t; \theta)}{F_j(t)} = \sum_{j=1}^N \sum_t F_j(t) \log P(w_j; \theta) - H(F_j)$$

Hence compute θ as:

$$\operatorname{argmax}_{\theta} \sum_{j=1}^N \sum_t F_j(t) \log P(w_j, t; \theta)$$

Answer:

$$\begin{aligned}
 \theta &= \operatorname{argmax}_{\theta} \sum_{j=1}^N \sum_{m=1}^K F_j(t_j = m) \log P(w_j, t_j = m | \theta) \\
 &= \operatorname{argmax}_{\theta} \sum_{j=1}^N \sum_{m=1}^K F_j(t_j = m) \log P(w_j | t_j = m; \theta) P(t_j = m | \theta) \\
 &= \operatorname{argmax}_{\theta} \sum_{j=1}^N \sum_{m=1}^K F_j(t_j = m) \log \left(\pi_m \prod_{l=1}^V \mu_m(l)^{w_j(l)} \right) \\
 &= \operatorname{argmax}_{\theta} \sum_{j=1}^N \sum_{m=1}^K [F_j(t_j = m) \log \pi_m + F_j(t_j = m) \sum_{l=1}^V \log \mu_m(l)^{w_j(l)}] \\
 &= \operatorname{argmax}_{\theta} \sum_{j=1}^N \sum_{m=1}^K [F_j(t_j = m) \log \pi_m + F_j(t_j = m) \sum_{l=1}^V w_j(l) \log \mu_m(l)] \quad (11)
 \end{aligned}$$

To optimize $\mu_m(l)$:

After eliminating from (11) the terms which are constant with respect to μ_m we get:

$$\sum_{j=1}^N F_j(t_j = m) \sum_{l=1}^V w_j(l) \log \mu_m(l)$$

We will use a Lagrangian to constrain μ_m to be a probability distribution:

$$\mathcal{L}(\mu_m(l)) = \sum_{j=1}^N F_j(t_j = m) \sum_{l=1}^V w_j(l) \log \mu_m(l) + \beta \left(\sum_{l=1}^V \mu_m(l) - 1 \right)$$

Then, solving for $\mu_m(l)$:

$$\begin{aligned} \frac{\partial}{\partial \mu_m(l)} \mathcal{L}(\mu_m(l)) &= 0 \Leftrightarrow \sum_{j=1}^N F_j(t_j = m) \frac{w_j(l)}{\mu_m(l)} + \beta = 0 \\ \Leftrightarrow \frac{1}{\mu_m(l)} \sum_{j=1}^N F_j(t_j = m) w_j(l) + \beta &= 0 \Leftrightarrow \frac{1}{\mu_m(l)} = \frac{-\beta}{\sum_{j=1}^N F_j(t_j = m) w_j(l)} \\ \Leftrightarrow \mu_m(l) &= \frac{\sum_{j=1}^N F_j(t_j = m) w_j(l)}{-\beta} \end{aligned} \tag{12}$$

Knowing that $\sum_{l=1}^V \mu_m(l) = 1$, we have:

$$\sum_{l=1}^V \frac{\sum_{j=1}^N F_j(t_j = m) w_j(l)}{-\beta} = 1 \Leftrightarrow -\beta = \sum_{l=1}^V \sum_{j=1}^N F_j(t_j = m) w_j(l)$$

Hence, substituting for $-\beta$ in (12), we get:

$$\begin{aligned} \mu_m(l) &= \frac{\sum_{j=1}^N F_j(t_j = m) w_j(l)}{\sum_{l=1}^V \sum_{j=1}^N F_j(t_j = m) w_j(l)} = \frac{\sum_{j=1}^N F_j(t_j = m) w_j(l)}{\sum_{j=1}^N \sum_{l=1}^V F_j(t_j = m) w_j(l)} \\ &= \frac{\sum_{j=1}^N F_j(t_j = m) w_j(l)}{\sum_{j=1}^N F_j(t_j = m) \underbrace{\sum_{l=1}^V w_j(l)}_1} = \frac{\sum_{j=1}^N F_j(t_j = m) w_j(l)}{\sum_{j=1}^N F_j(t_j = m)} \end{aligned}$$

Note: Intuitively the last expression can be interpreted as the portion [of word[occurrence]s] that had $w(l) = 1$ among all words [in the given document] which are deemed to belong to cluster m .

To optimize π_m , we proceed similarly:

We begin by removing from (11) the terms that are constant with respect to π_m , thus getting:

$$\sum_{j=1}^N F_j(t_j = m) \log \pi_m$$

So, using the Lagrangian with the constraint that $\sum_{m=1}^K \pi_m = 1$

$$\mathcal{L}(\pi_m) = \sum_{j=1}^N F_j(t_j = m) \log \pi_m + \beta \left(\sum_{m=1}^K \pi_m - 1 \right),$$

and solving for π_m , we have:

$$\begin{aligned} \frac{\partial}{\partial \pi_m} \mathcal{L}(\pi_m) = 0 &\Leftrightarrow \sum_{j=1}^N \frac{F_j(t_j = m)}{\pi_m} + \beta = 0 \Leftrightarrow \frac{1}{\pi_m} \sum_{j=1}^N F_j(t_j = m) = -\beta \\ &\Leftrightarrow \pi_m = \frac{\sum_{j=1}^N F_j(t_j = m)}{-\beta} \end{aligned} \tag{13}$$

Since $\sum_{m=1}^K \pi_m = 1$, it gives us:

$$\begin{aligned} \sum_{m=1}^K \frac{\sum_{j=1}^N F_j(t_j = m)}{-\beta} = 1 &\Leftrightarrow \frac{1}{-\beta} \sum_{m=1}^K \sum_{j=1}^N F_j(t_j = m) = 1 \\ &\Leftrightarrow -\beta = \sum_{m=1}^K \sum_{j=1}^N F_j(t_j = m) \end{aligned}$$

Substituting for $-\beta$ in (13), we get:

$$\pi_m = \frac{\sum_{j=1}^N F_j(t_j = m)}{\sum_{m=1}^K \sum_{j=1}^N F_j(t_j = m)} = \frac{\sum_{j=1}^N F_j(t_j = m)}{\sum_{j=1}^N \underbrace{\sum_{m=1}^K F_j(t_j = m)}_1} = \frac{\sum_{j=1}^N F_j(t_j = m)}{N}$$

Note: Intuitively the last expression can be interpreted as the portion [of word[occurrence]s] that belongs to cluster m among the total of N words.

To summarize:

E step:

$$F_j(t_j = m) = \frac{\pi_m \prod_{l=1}^V \mu_m(l)^{w_j(l)}}{\sum_{m'=1}^K \pi_{m'} \prod_{l=1}^V \mu_{m'}(l)^{w_j(l)}} \text{ for } j = 1, \dots, N \text{ and } m = 1, \dots, K$$

M step:

$$\mu_m(l) = \frac{\sum_{j=1}^N F_j(t_j = m) w_j(l)}{\sum_{j=1}^N F_j(t_j = m)} \text{ for } m = 1, \dots, K \text{ and } l = 1, \dots, V$$

$$\pi_m = \frac{\sum_{j=1}^N F_j(t_j = m)}{N} \text{ for } m = 1, \dots, K$$

Using the EM algorithm for
estimating the *selection probability* for a mixture
of two (arbitrary) distributions

CMU, 2006 spring, ?, final exam, pr. 8

CMU, 2004 fall, Carlos Guestrin, HW2, pr. 2.1

We want to derive an EM algorithm for estimating the mixing parameter for a mixture of arbitrary probability densities f_1 and f_2 .

For *example*, $f_1(x)$ could be a standard normal distribution centered at 0, and $f_2(x)$ could be the uniform distribution between $[0, 1]$. You can think about such mixtures in the following way: First, you flip a coin. With probability λ (i.e., the coin comes up *heads*), you will sample x from density f_1 , and with probability $(1 - \lambda)$ you sample from density f_2 .

More formally, let $f_\lambda(x) = \lambda f_1(x) + (1 - \lambda) f_2(x)$, where f_1 and f_2 are arbitrary probability density functions on \mathbb{R} , and $\lambda \in [0, 1]$ is an unknown mixture parameter.

a.

Given a data point x , and a value for the mixture parameter λ , compute the probability that x was generated from density f_1 .

b.

Now, suppose you are given a data set $\{x_1, \dots, x_n\}$ drawn i.i.d. from the mixture density, and a set of coin flips $\{z_1, z_2, \dots, z_n\}$, such that $z_i = 1$ means that x_i is a sample from f_1 , and $z_i = 0$ means that x_i was generated from density f_2 .

For a fixed parameter λ , compute the complete log-likelihood of the data, i.e., $\ln P(x_1, z_1, x_2, z_2, \dots, x_n, z_n | \lambda)$.

c.

Now, suppose you are given only a sample $\{x_1, \dots, x_n\}$ drawn i.i.d. from the mixture density, without the knowledge about which component the samples were drawn from (i.e., the z_i are unknown).

Using your derivations from part *a* and *b*, derive the E- and M-steps for an EM algorithm to compute the maximum likelihood estimate (MLE) of the mixture parameter λ .

Solution

$$\text{a. } P(Y = 1|X = x) = \frac{P(X = x|Y = 1) \cdot P(Y = 1)}{P(X = x)} = \frac{\lambda f_1(x)}{f_\lambda(x)}.$$

$$\text{b. } P(x_1, z_1, x_2, z_2, \dots, x_n, z_n|\lambda) = \prod_{i=1}^n P(x_i, z_i|\lambda) = \prod_{i=1}^n P(x_i|z_i, \lambda) \cdot P(z_i|\lambda).$$

$$P(x_i|z_i, \lambda) \cdot P(z_i|\lambda) = \begin{cases} \lambda f_1(x_i) & \text{if } z_i = 1 \\ (1 - \lambda) f_2(x_i) & \text{if } z_i = 0 \end{cases} = f_1(x_i)^{z_i} (1 - \lambda)^{1-z_i} f_2(x_i)^{1-z_i}$$

Therefore,

$$\begin{aligned} \ln P(x_1, z_1, x_2, z_2, \dots, x_n, z_n|\lambda) &= \sum_{i=1}^n \ln P(x_i, z_i|\lambda) \\ &= \sum_{i=1}^n \ln \left(\lambda^{z_i} f_1(x_i)^{z_i} (1 - \lambda)^{1-z_i} f_2(x_i)^{1-z_i} \right) \\ &= \sum_{i=1}^n z_i (\ln \lambda + \ln f_1(x_i)) + (1 - z_i) (\ln(1 - \lambda) + \ln f_2(x_i)) \end{aligned}$$

c.

E-step: $q(z_i) \stackrel{not.}{=} P(z_i = 1|x_i, \lambda^{(t)}) \stackrel{B.Th.}{=} \frac{\lambda^{(t)} f_1(x_i)}{f_{\lambda^{(t)}}(x_i)}$

M-step:

$$\begin{aligned} \lambda^{(t+1)} &= \operatorname{argmax}_{\lambda} E_q \left[\ln \prod_{i=1}^n P(x_1, z_1, x_2, z_2, \dots, x_n, z_n | \lambda) \right] \\ &= \operatorname{argmax}_{\lambda} \left(\ln \lambda \cdot \underbrace{\sum_{i=1}^n q(z_i)}_c + \ln(1 - \lambda) \cdot \underbrace{\sum_{i=1}^n (1 - q(z_i))}_{n-c} \right) \end{aligned}$$

$$\begin{aligned} \frac{\partial}{\partial \lambda} (c \ln \lambda + (n - c) \ln(1 - \lambda)) &= 0 \Leftrightarrow \frac{c}{\lambda} = \frac{n - c}{1 - \lambda} \\ \Leftrightarrow c(1 - \lambda) &= \lambda(n - c) \Leftrightarrow c = n\lambda \Leftrightarrow \lambda = \frac{c}{n} = \frac{\sum_{i=1}^n q(z_i)}{n} \\ \Rightarrow \lambda^{(t+1)} &= \frac{\lambda^{(t)}}{n} \sum_{i=1}^n \frac{f_1(x_i)}{f_{\lambda^{(t)}}(x_i)} \end{aligned}$$